


# Microservice Design and Architecture



---

# The following course materials are **copyright protected** materials.

They may not be reproduced or distributed and may only be used by students attending this Google Cloud Partner Learning Services program.

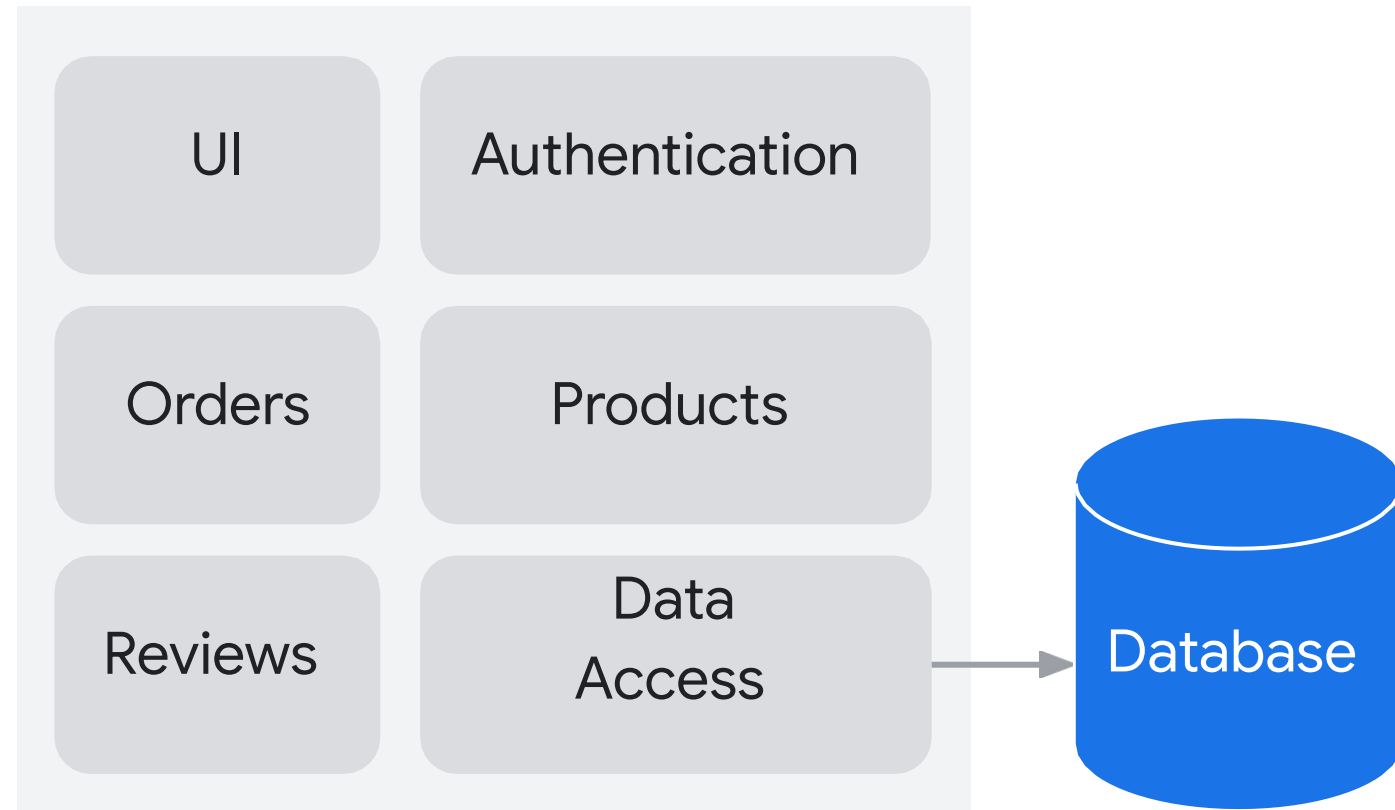




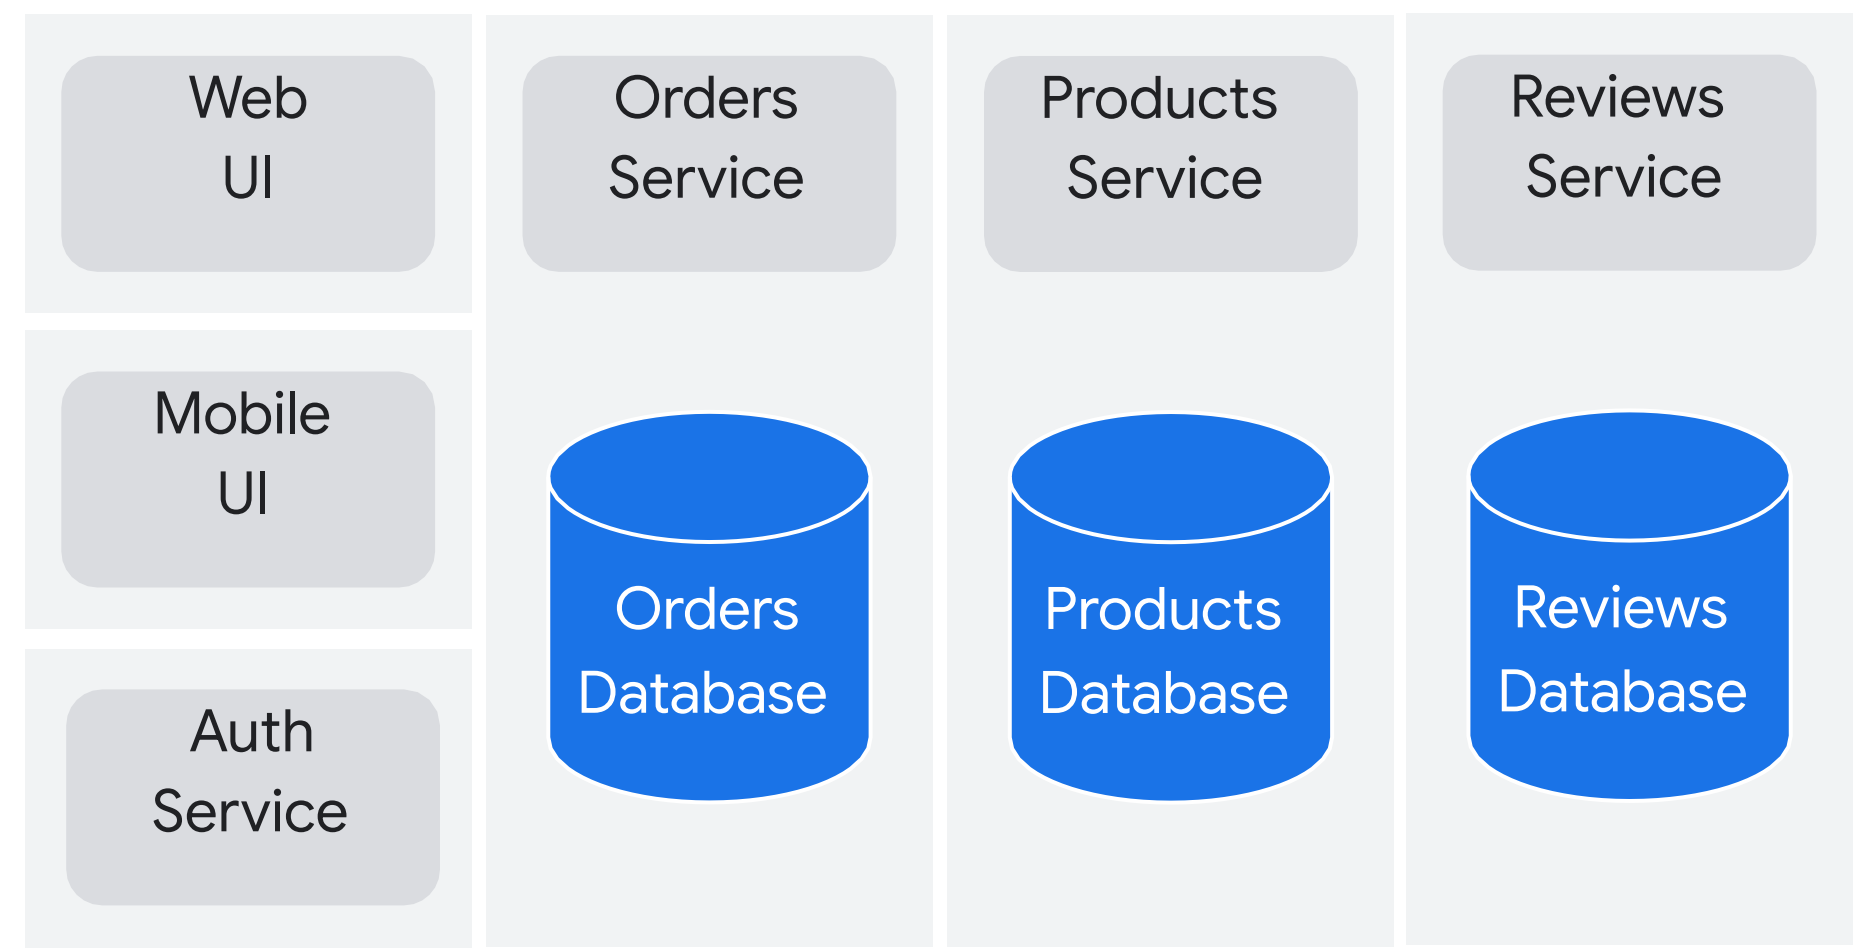
# Microservices

# Microservices divide a large program into multiple smaller, independent services

Monolithic applications implement all features in a single code base with a database for all data.



Microservices have multiple code bases, and each service manages its own data.



# Pros and cons of microservice architectures

## Pros



- Easier to develop and maintain.
- Reduced risk when deploying new versions.
- Services scale independently to optimize use of infrastructure.
- Faster to innovate and add new features.
- Can use different languages and frameworks for different services.
- Choose the runtime appropriate to each service.

## Con



- Increased complexity when communicating between services.
- Increased latency across service boundaries.
- Services scale independently to optimize use of infrastructure.
- Concerns about securing inter-service traffic.
- Multiple deployments.
- Need to ensure that you don't break clients as versions change.
- Must maintain backward compatibility with clients as the microservice evolves.

# The key to architecting microservice applications is recognizing service boundaries

01

Decompose applications by feature to minimize dependencies

- Reviews service
- Orders service
- Products service
- Etc.

02


Organize services by architectural layer

- Web, Android, and iOS user interfaces
- Data access services

03

Isolate services that provide shared functionality

- Authentication service
- Reporting service
- Etc.



# Microservice Best Practices

# The 12-factor app is a set of best practices for building web or software-as-a-service applications

- Maximize portability
- Deploy to the cloud
- Enable continuous deployment
- Scale easily



## THE TWELVE-FACTOR APP



# The 12 factors

01

## Codebase

One codebase tracked in revision control, many deploys

- Use a version control system like Git.
- Each app has one code repo and vice versa.

02

## Dependencies

Explicitly declare and isolate dependencies

- Use a package manager like Maven, Pip, NPM to install dependencies.
- Declare dependencies in your code base.

03

## Config

Store config in the environment

- Don't put secrets, connection strings, endpoints, etc., in source code.
- Store those as environment variables.

04

## Backing Services

Treat backing services as attached resources

- Databases, caches, queues, and other services are accessed via URLs.
- Should be easy to swap one implementation for another.

# The 12 factors (continued)

05

## Build, release, run

Strictly separate build and run stages

- Build creates a deployment package from the source code.
- Release combines the deployment with configuration in the runtime environment.
- Run executes the application.

06

## Processes

Execute the app as one or more stateless processes

- Apps run in one or more processes.
- Each instance of the app gets its data from a separate database service.

07

## Port binding

Export services via port binding

- Apps are self-contained and expose a port and protocol internally.
- Apps are not injected into a separate server like Apache.

08

## Concurrency

Scale out via the process model

- Because apps are self-contained and run in separate process, they scale easily by adding instances.

# The 12 factors (continued)

09

## Disposability

Maximize robustness with fast startup and graceful shutdown

- App instances should scale quickly when needed.
- If an instance is not needed, you should be able to turn it off with no side effects.

10

## Dev/prod parity

Keep development, staging, and production as similar as possible

- Container systems like Docker makes this easier.
- Leverage infrastructure as code to make environments easy to create.

11

## Logs

Treat logs as event streams

- Write log messages to standard output and aggregate all logs to a single source.

12

## Admin processes

Run admin/management tasks as one-off processes

- Admin tasks should be repeatable processes, not one-off manual tasks.
- Admin tasks shouldn't be a part of the application.

03



REST

# A good microservice design is loosely coupled

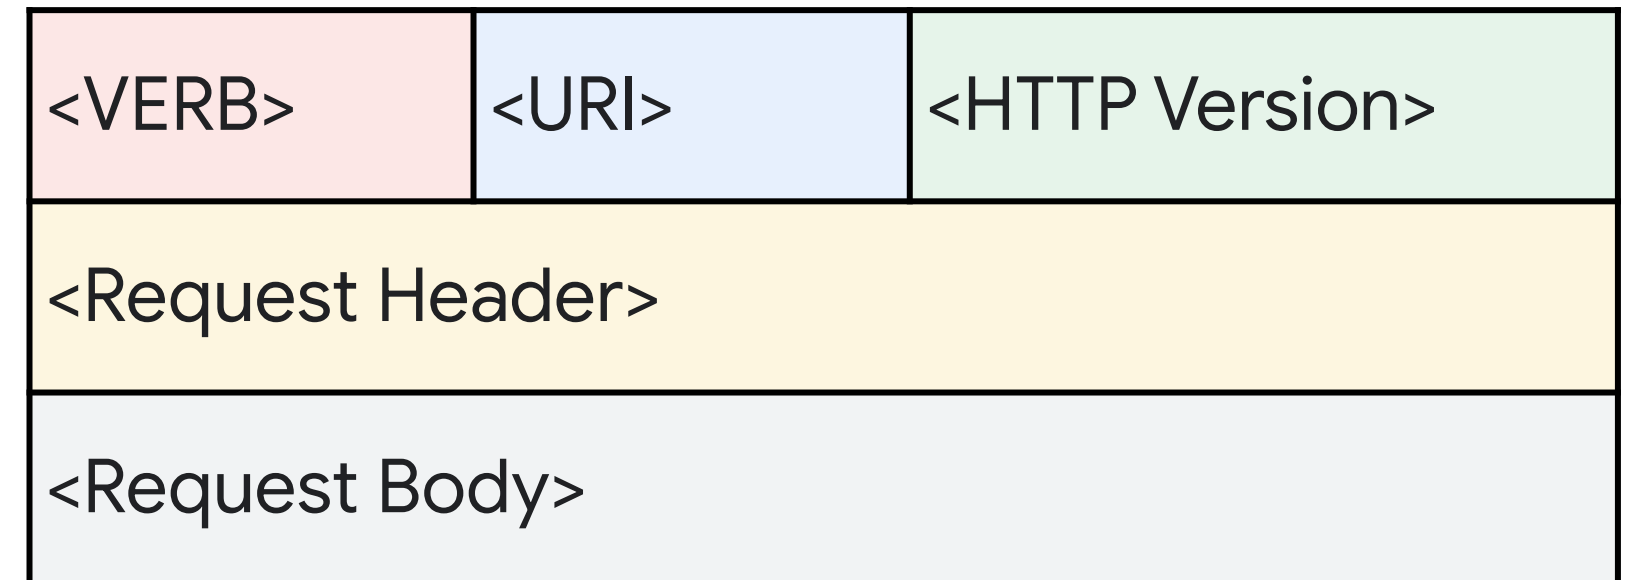
- Clients shouldn't need to know too many details of services they use
- Services communicate via HTTPS using text-based payloads
  - Client makes GET, POST, PUT, or DELETE request
  - Body of the request is formatted as JSON or XML
  - Results returned as JSON, XML, or HTML
- Services should add functionality without breaking existing clients
  - Add, but don't remove, items from responses



*If microservices aren't loosely coupled, you'll end up with a really complicated monolith.*

# Clients access services using HTTP requests

- VERB: GET, PUT, POST, DELETE
- URI: Uniform Resource Identifier (endpoint)
- Request Header: metadata about the message
  - Preferred representation formats (e.g., JSON, XML)
- Request Body: (Optional) Request state
  - Representation (JSON, XML) of resource

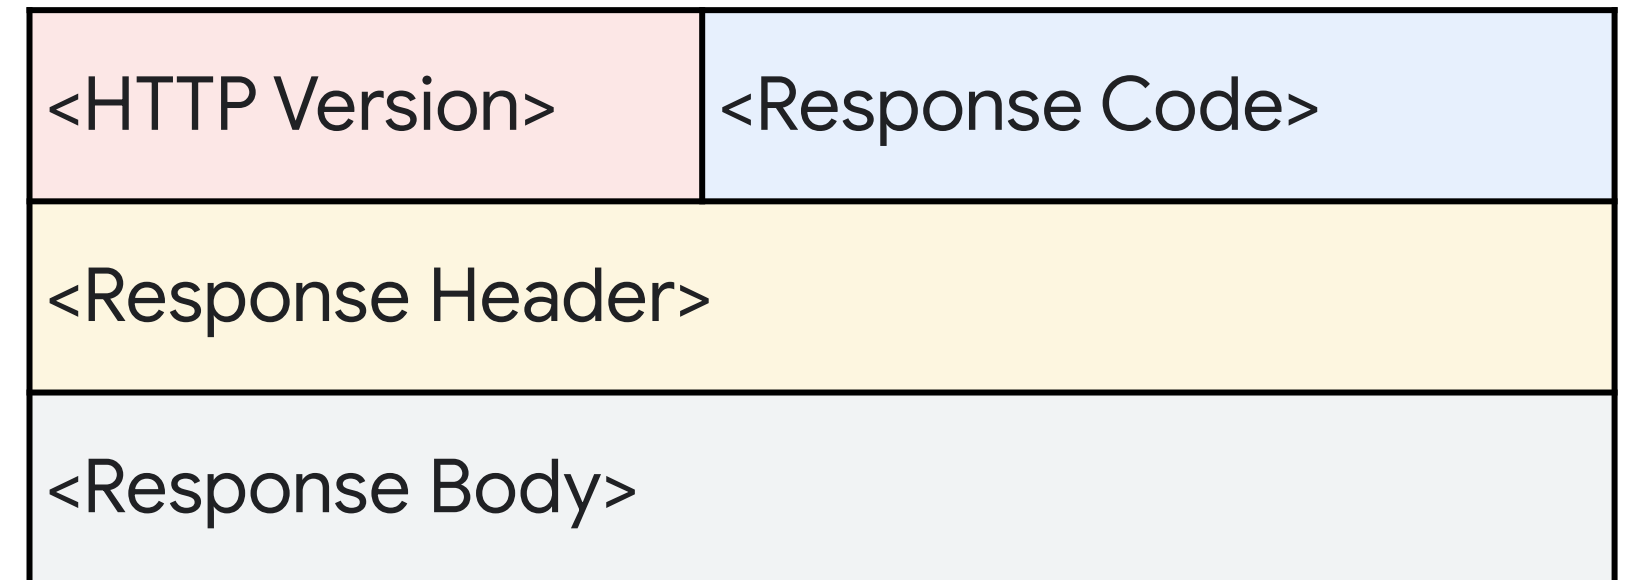


# The HTTP verb tells the server what to do

- **GET** is used to retrieve data
- **POST** is used to create data
  - Generates entity ID and returns it to the client
- **PUT** is used to create data or alter existing data
  - Entity ID must be known
  - *PUT should be idempotent, which means that whether the request is made once or multiple times, the effects on the data are exactly the same*
- **DELETE** is used to remove data

# Services return HTTP responses

- Response Code: 3-digit HTTP status code
  - 200 codes for success
  - 400 codes for client errors
  - 500 codes for server errors
- Response Body: contains resource representation
  - JSON, XML, HTML, etc.







APIs

# It's important to design consistent APIs for services


- Each Google Cloud service exposes a REST API
  - Functions are in the form: `service.collection.verb`
  - Parameters are passed either in the URL or in the request body in JSON format
- For example, the Compute Engine API has...
  - A service endpoint at: <https://compute.googleapis.com>
  - Collections include instances, instanceGroups, instanceTemplates, etc.
  - Verbs include insert, list, get, etc.
- So, to see all your instances, make a GET request to:  
<https://compute.googleapis.com/compute/v1/projects/{project}/zones/{zone}/instances>

# Google Cloud provides two tools, Cloud Endpoints and Apigee, for managing APIs


Both provide tools for:

- User authentication
- Monitoring
- Securing APIs
- Etc.

Both support OpenAPI and gRPC



Cloud  
Endpoints



Apigee API  
Platform

# More resources

## API Design Guide

<https://cloud.google.com/apis/design/>


## Authenticating service-to-service calls with Google Cloud Endpoints

<https://youtu.be/4PgX3yBJEyw>

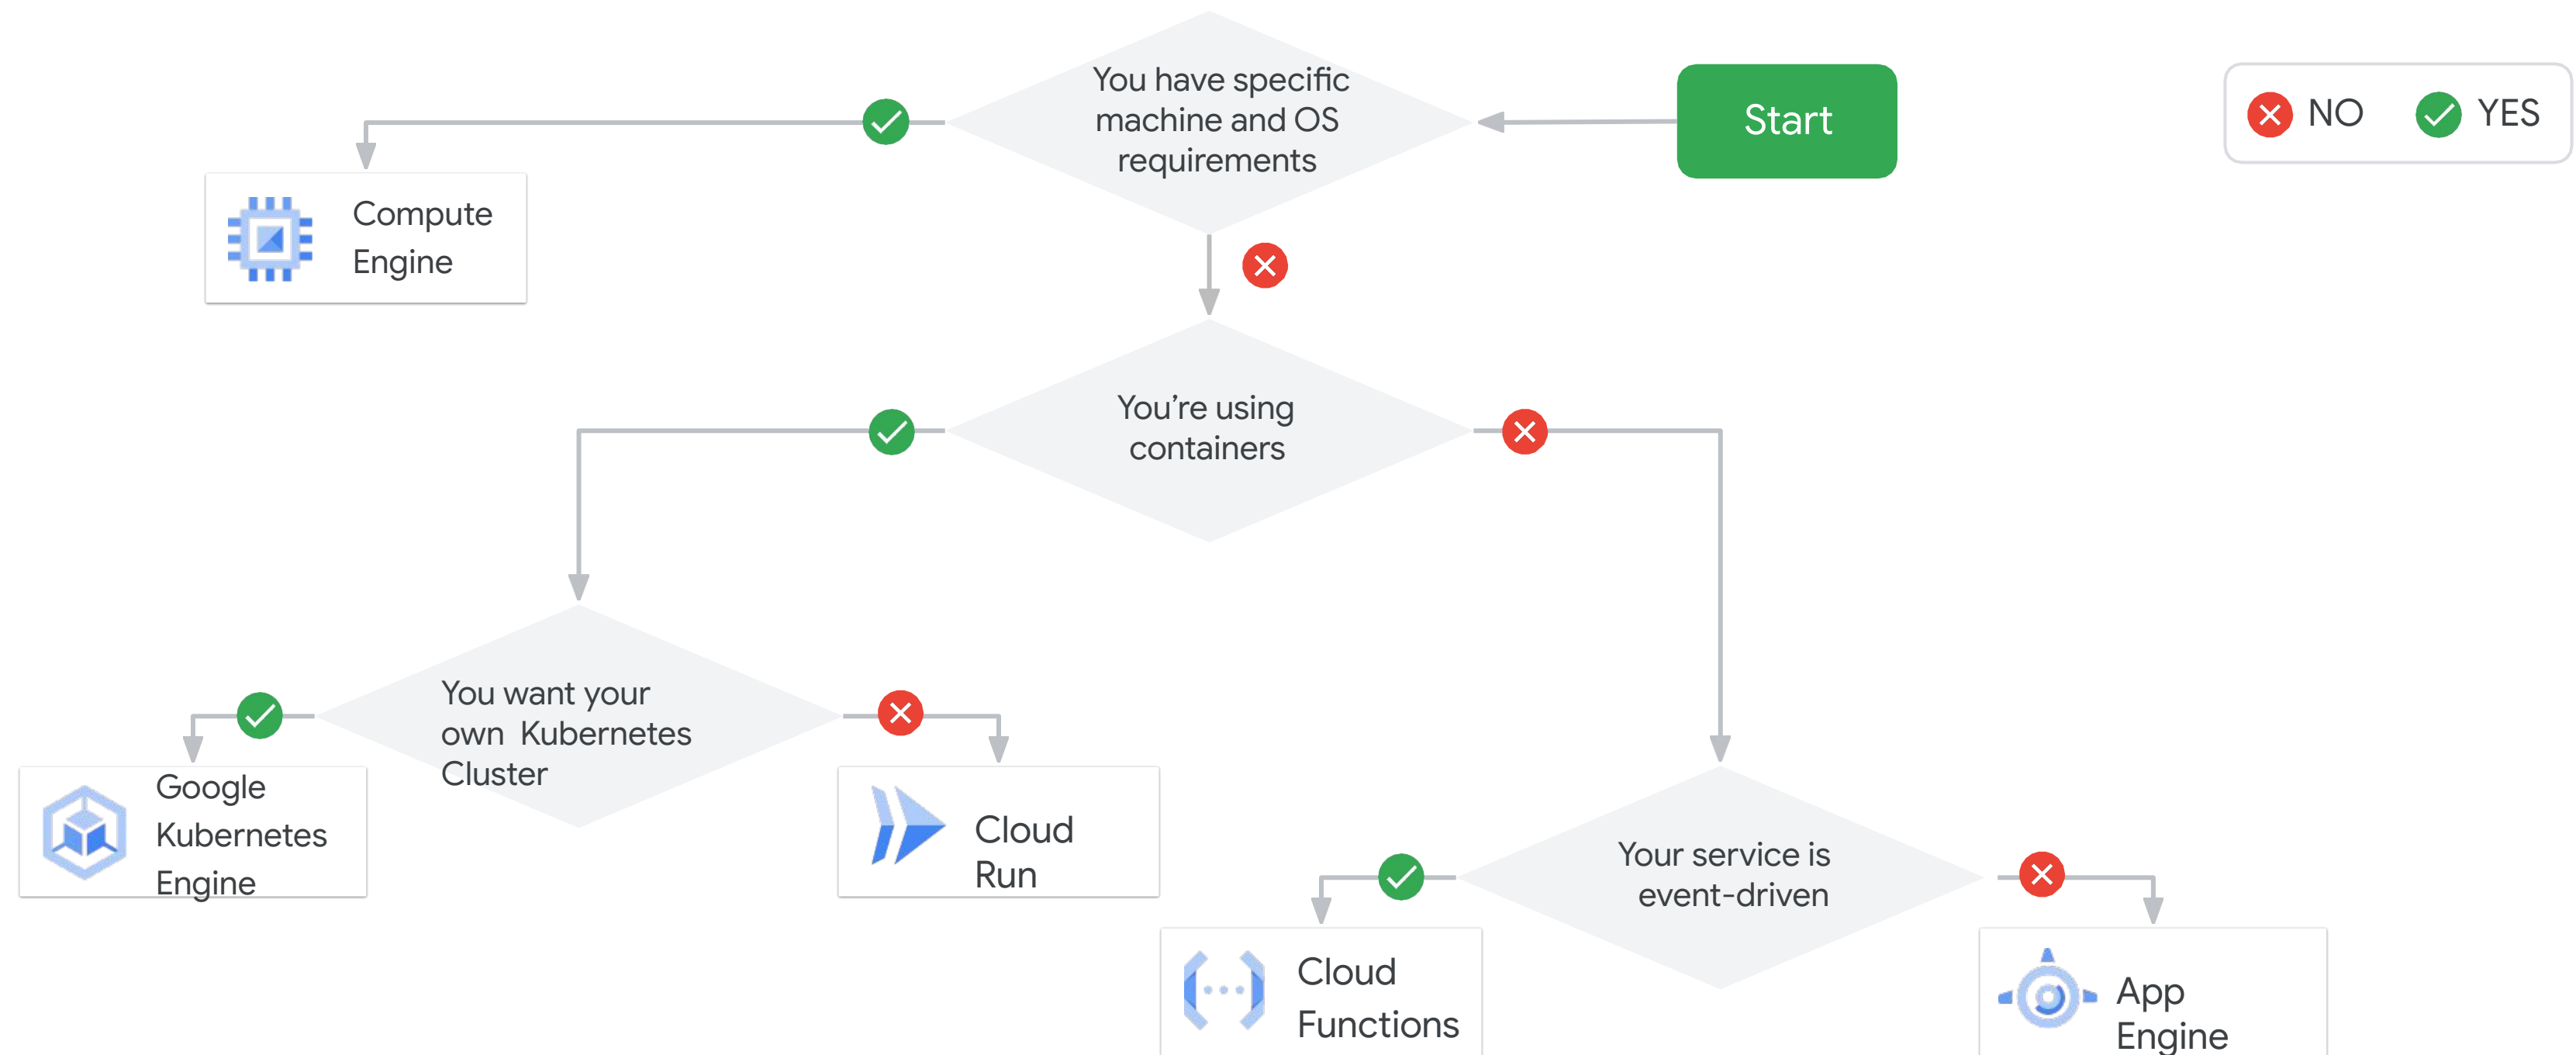




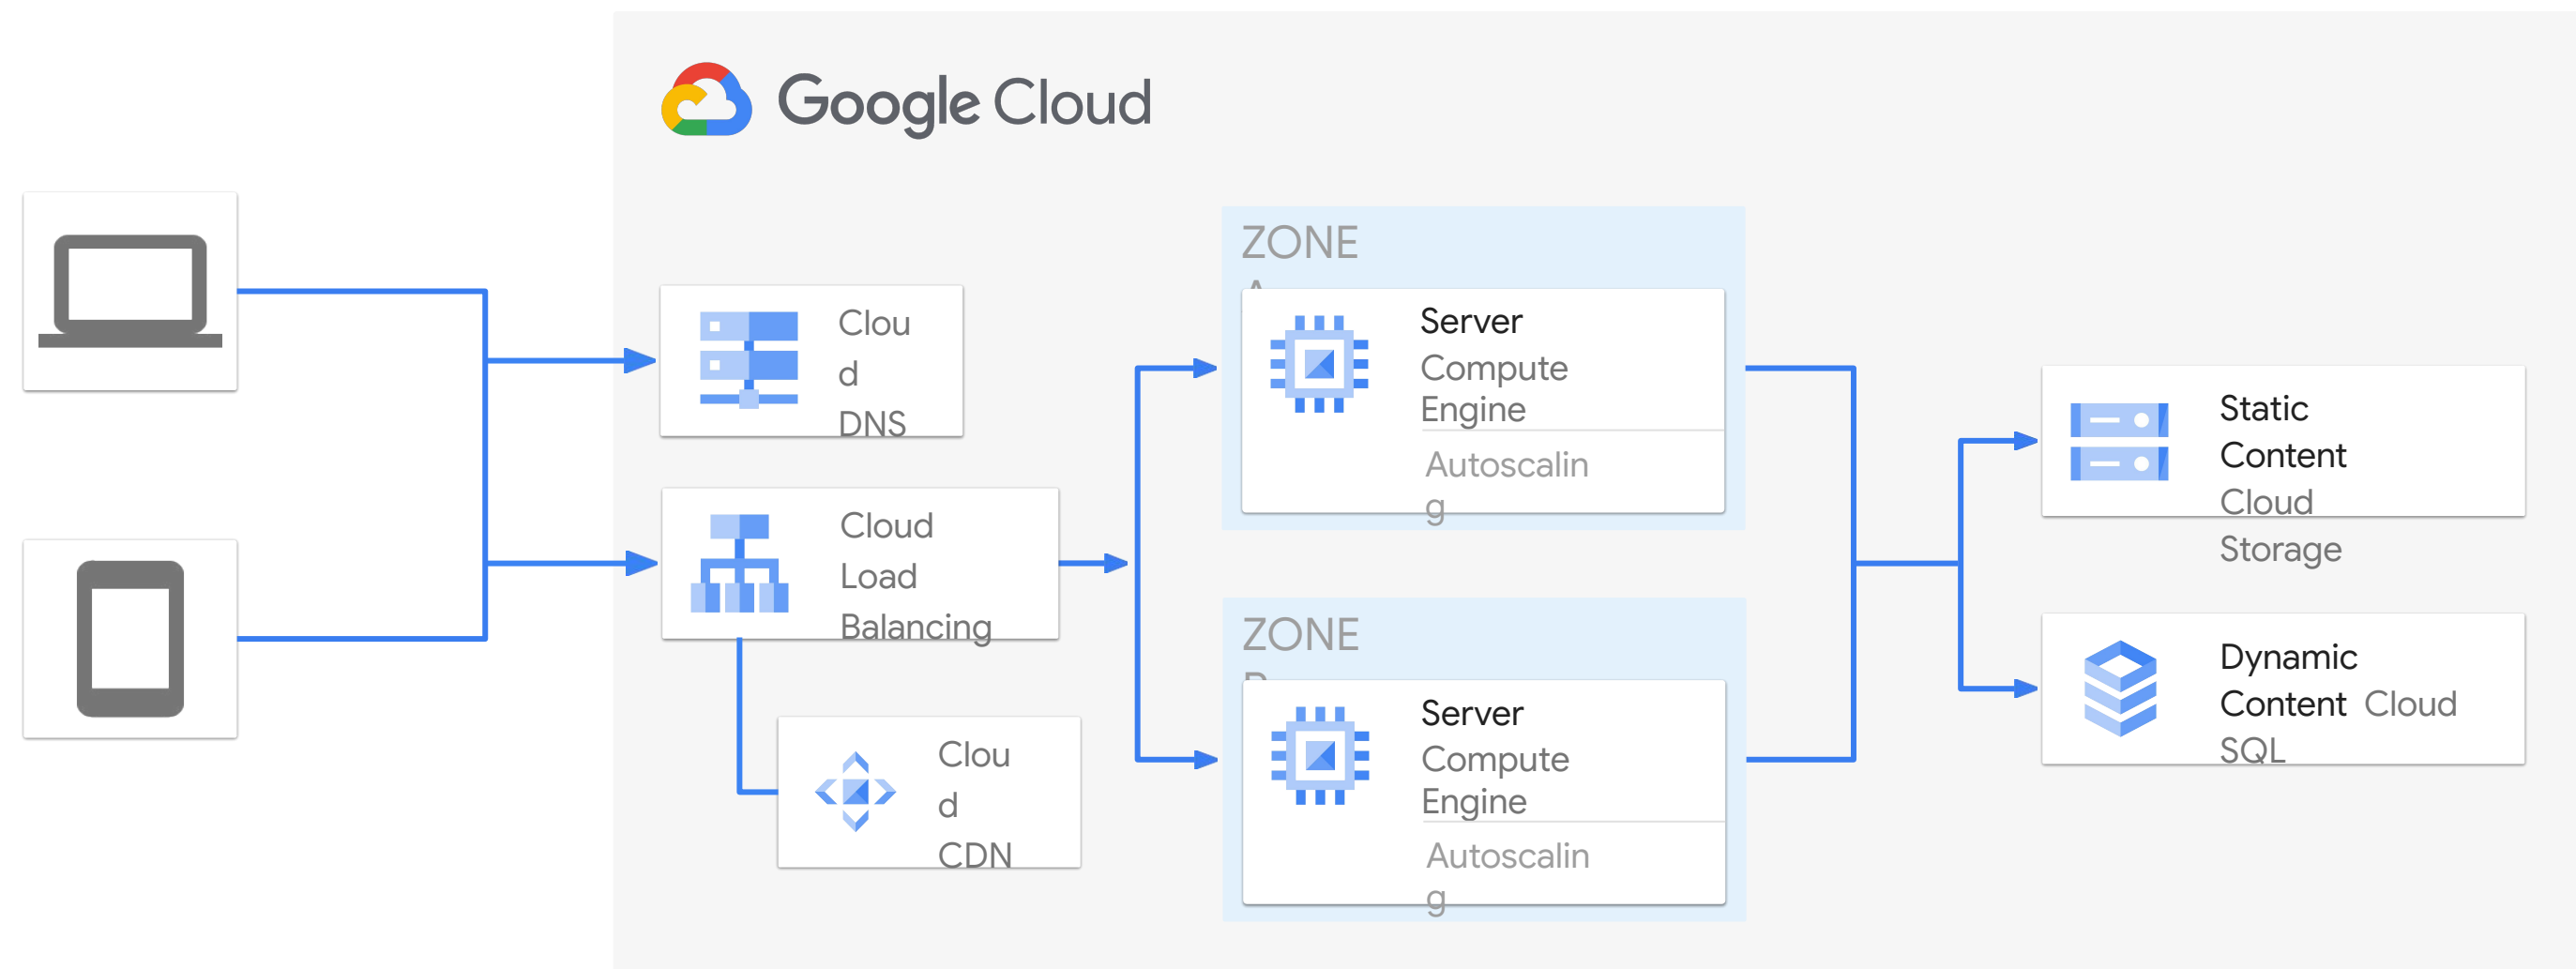
# Deploying Applications to Google Cloud



# Choosing a Google Cloud deployment platform



# Use Compute Engine when you need complete control over operating systems, for apps that are not containerized or self-hosted databases



# Managed instance groups create VMs based on instance templates

- Instance templates define the VMs:
  - image, machine type, etc.
  - Test to find the smallest machine type that will run your program.
- Use a Startup Script to install your program from a Git repo.
- Instance group manager creates the machines.
- Set up auto scaling to optimize cost and meet varying user workloads.
- Add a health check to enable auto healing.
- Use multiple zones for high availability.




# Use one or more instance groups as the backend for load balancers

- Use a global load balancer if you have instance groups in multiple regions.
- Enable the CDN to cache static content.
- For external services, set up SSL.
- For internal services, don't provide a public IP address.

### Create backend service

**Name** ?  
Name is permanent


⌵ **Description**


Protocol: HTTP   Named port: http   Timeout: 30 seconds 

**Backend type**

☒ Instance groups  
☐ Network endpoint groups

**Backends**  
Regions: europe-north1, us-central1


instance-group-europe (Zone: europe-north1-a, Port: 80) 

instance-group-us (Zone: us-central1-a, Port: 80) 

+ Add backend

**Cloud CDN** ?

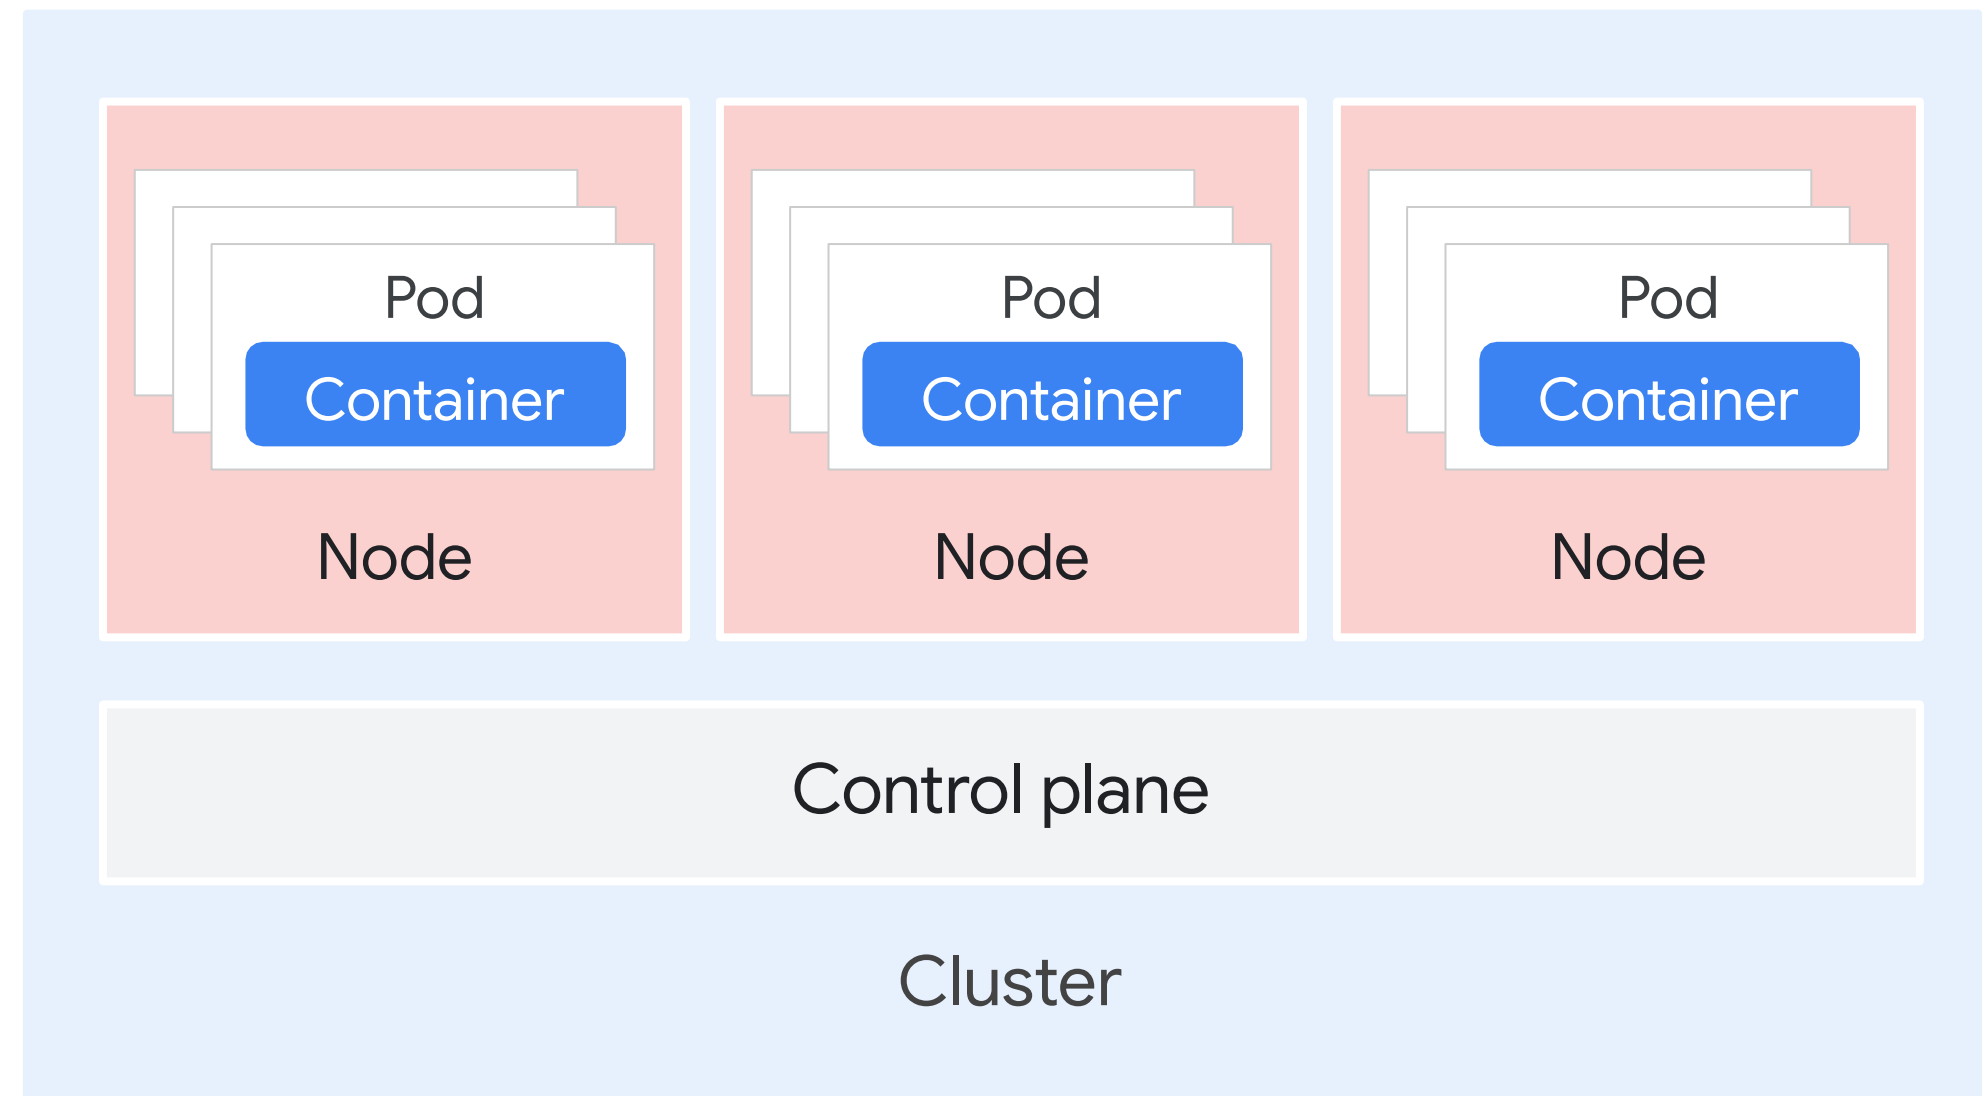
☒ Enable Cloud CDN



# Google Cloud Deployment Platforms

# Google Kubernetes Engine (GKE) automates the creation and management of compute infrastructure

- Kubernetes clusters have a collection of nodes.
- In GKE, nodes are Compute Engine VMs.
- Services are deployed into pods.
- Optimize resource utilization by deploying multiple services to the same cluster.
- You pay for the VMs.



# Cloud Run allows you to deploy containers to Google-managed Kubernetes clusters

- Cloud Run allows you to use Kubernetes without the cluster management or configuration code.
- Apps must be stateless.
- Need to deploy apps using Docker images in Container Registry.
- Can also use Cloud Run to automate deployment to your own GKE cluster.

## Container

Container image URL \*

gcr.io/doug-rehnstrom/pets-app@sha256:bc43dbb6adf9e3ff4083971ff48

SELECT

E.g. gcr.io/cloudrun/hello

Must be stateless and listen for HTTP requests on \$PORT. [How to build a container?](#)

## Deployment platform ?

☒ Cloud Run (fully managed)

Location \*

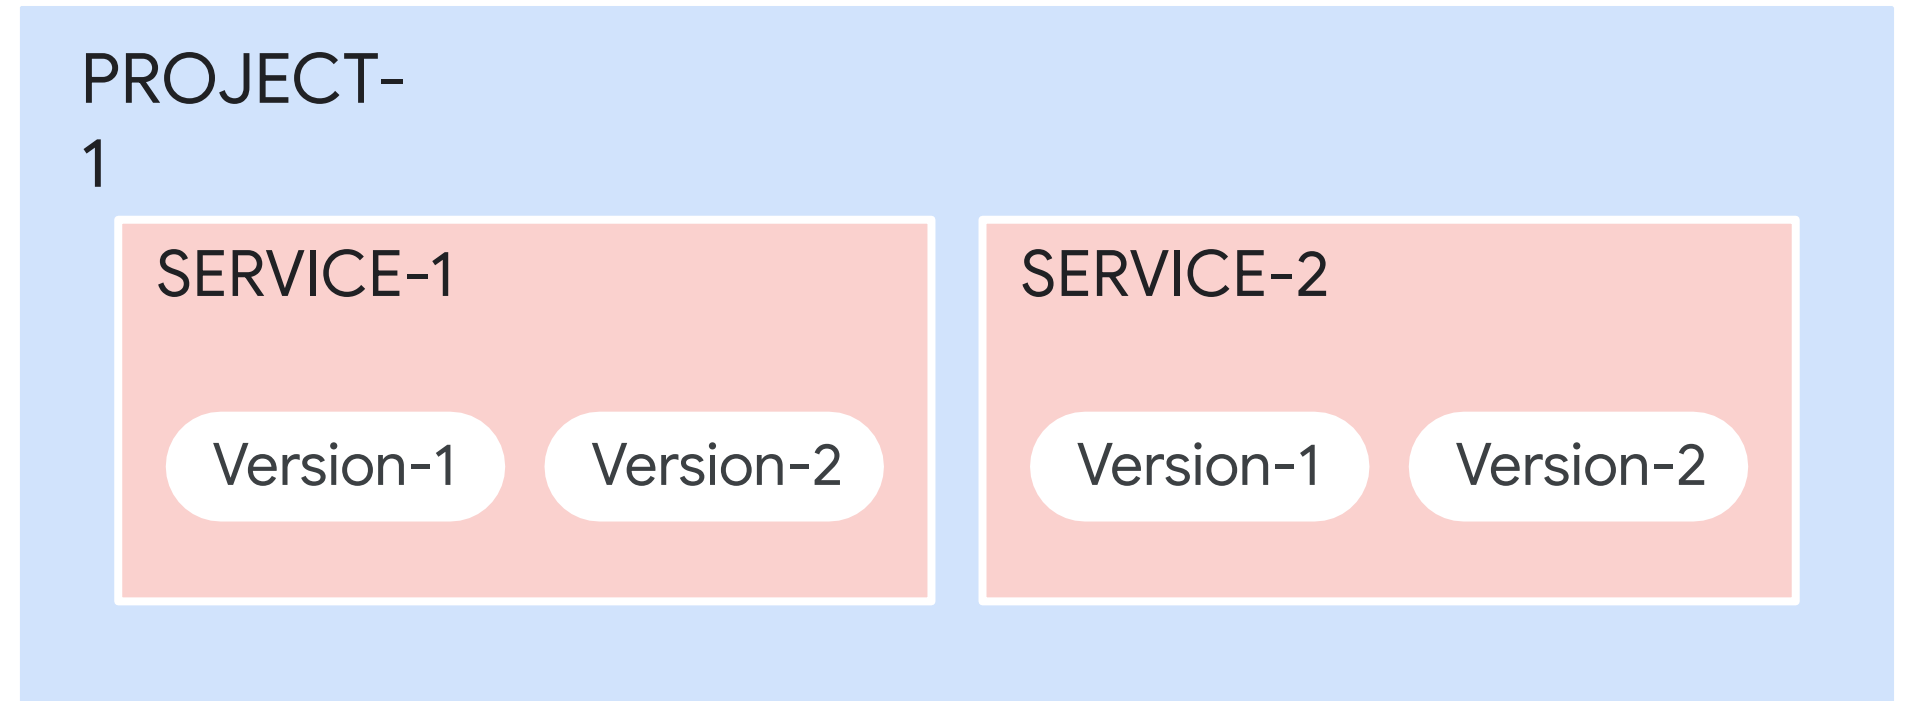
us-central1

Region for this Service can't be changed later. [How to pick a region?](#)

☐ Cloud Run for Anthos

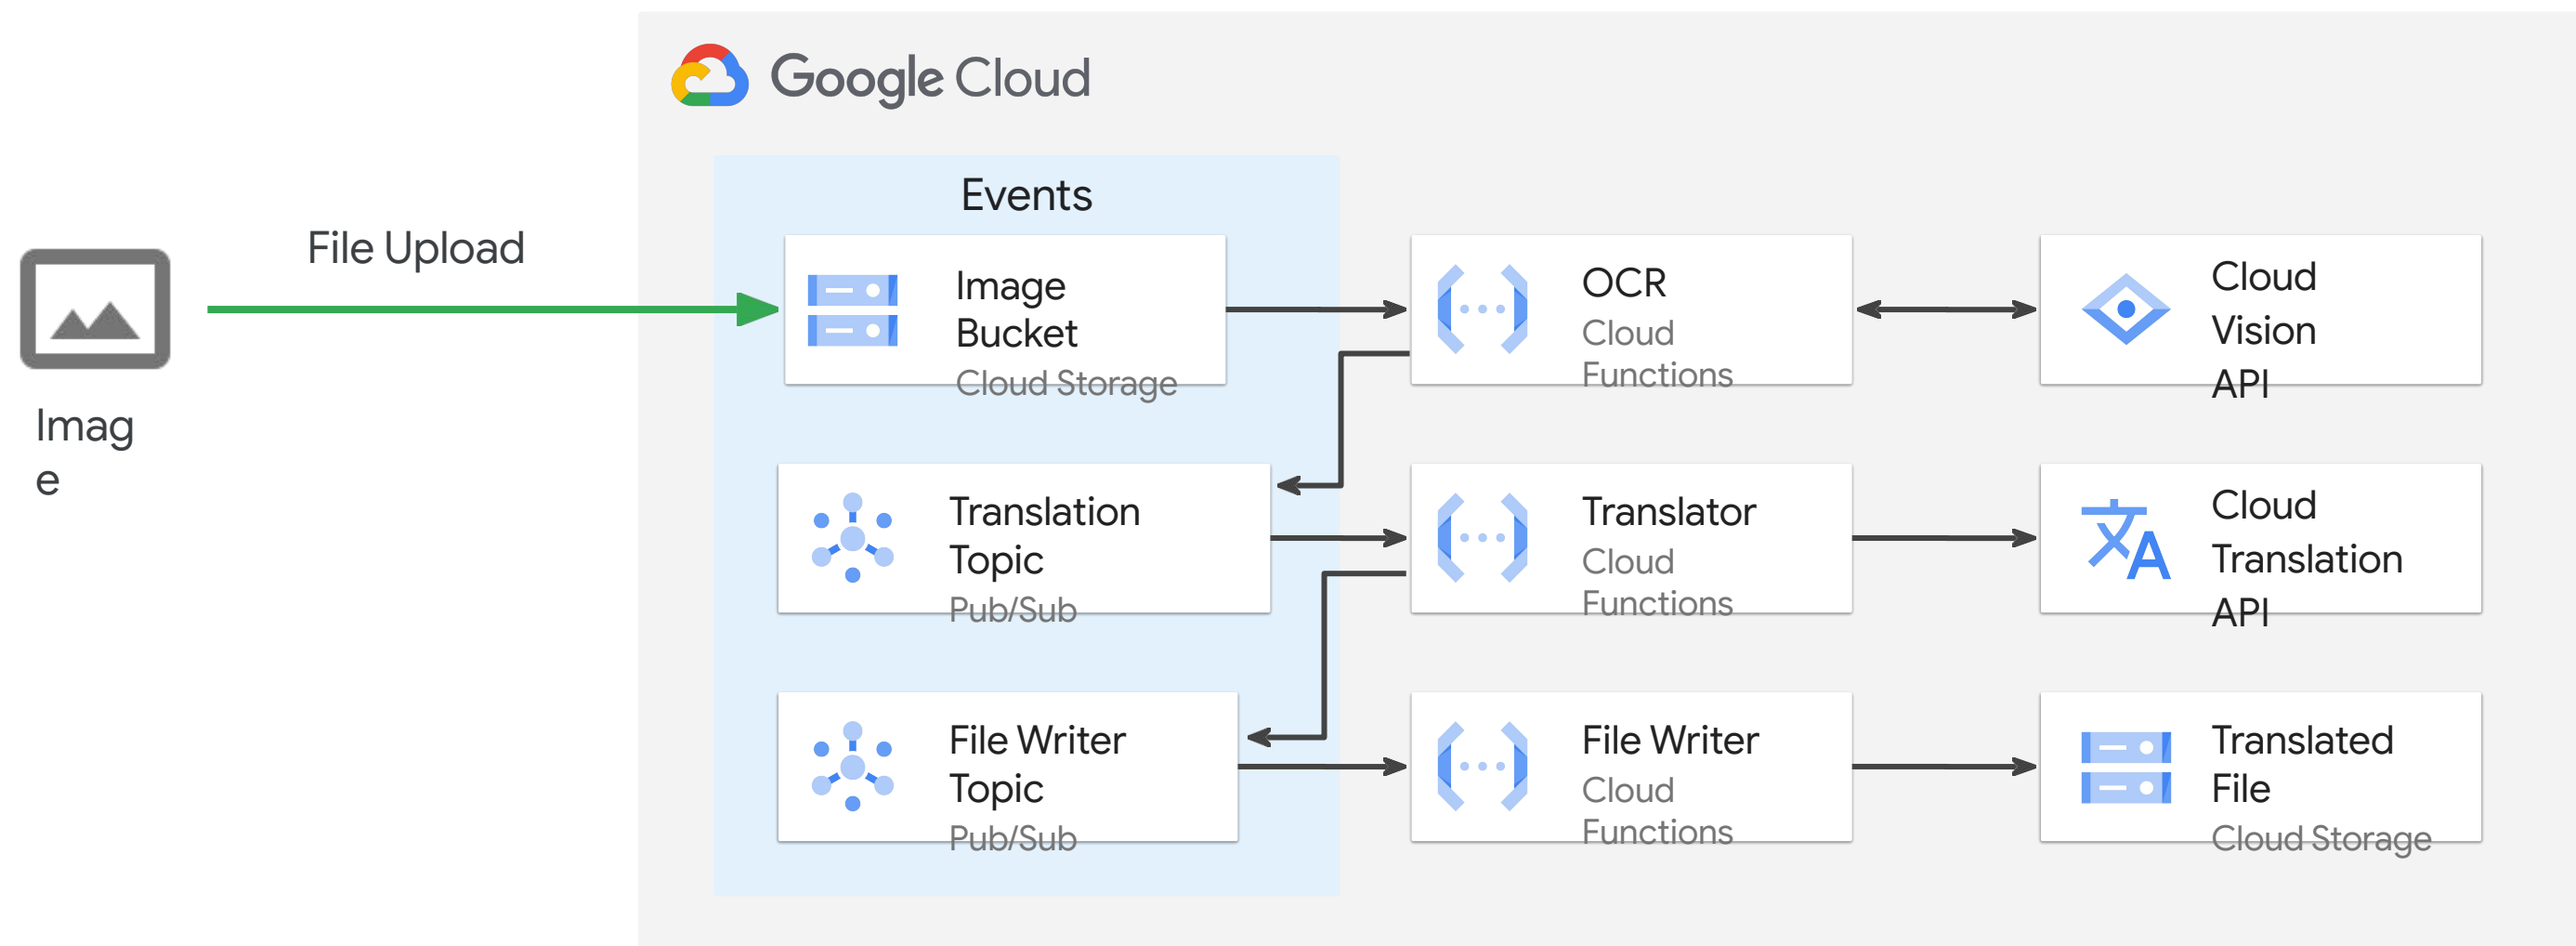
# App Engine was designed for microservices

- Each Google Cloud project can contain 1 App Engine application.
- An application has 1 or more services.
- Each service has 1 or more versions.
- Versions have 1 or more instances.
- Automatic traffic splitting for switching versions.



# Cloud Functions is great way to create loosely coupled, event-driven microservices

- Can be triggered by changes in a storage bucket, Pub/Sub messages, or web requests
- Completely managed, scalable, and inexpensive



# More resources

Migration to Google Cloud:  
Deploying your workloads

<https://cloud.google.com/solutions/migration-to-gcp-deploying-your-workloads>

Compute Engine

<https://cloud.google.com/compute/>

GKE

<https://cloud.google.com/kubernetes-engine/>

App Engine

<https://cloud.google.com/appengine/>



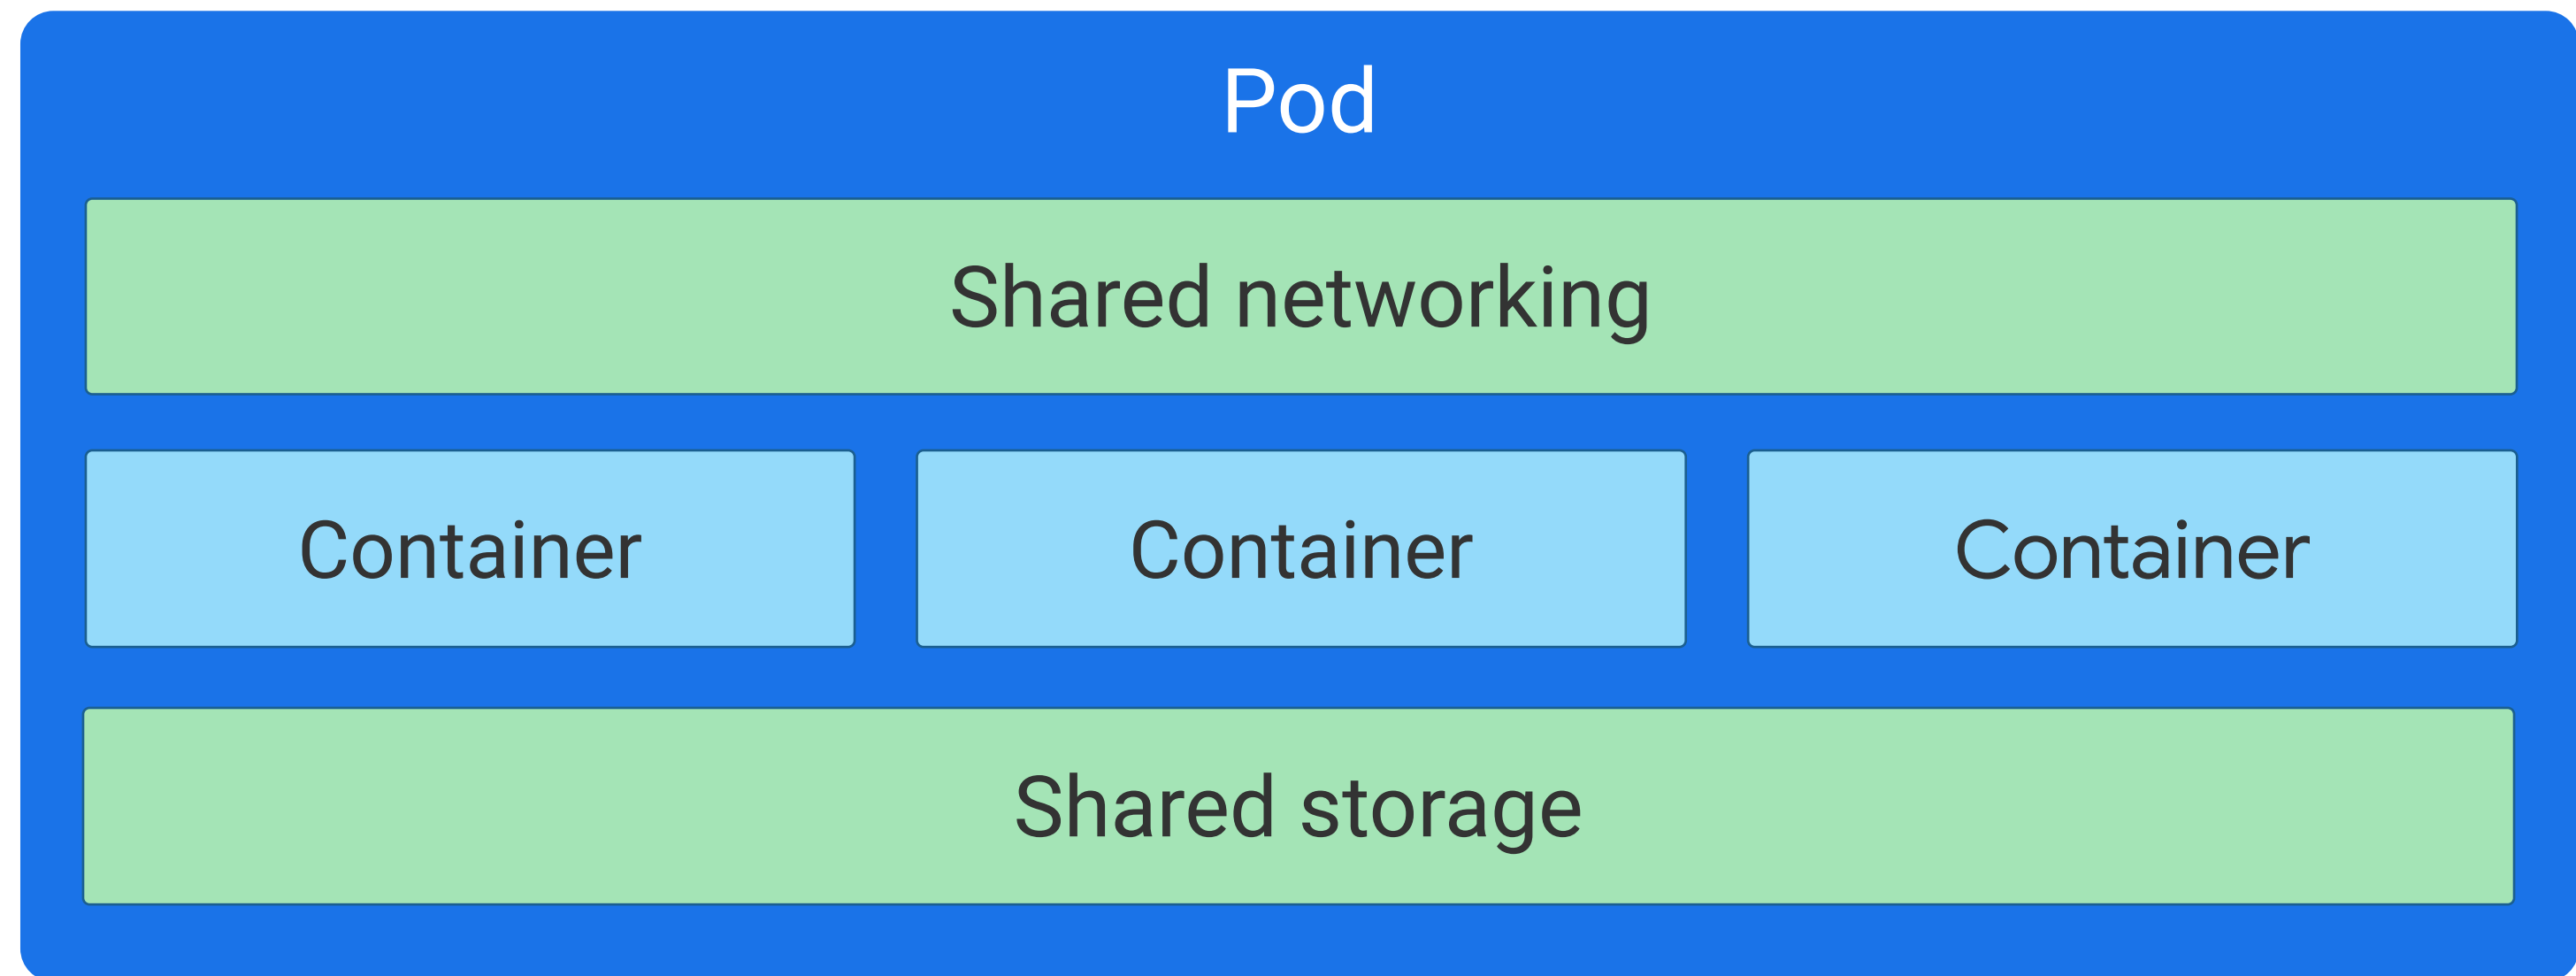


# Kubernetes Architecture

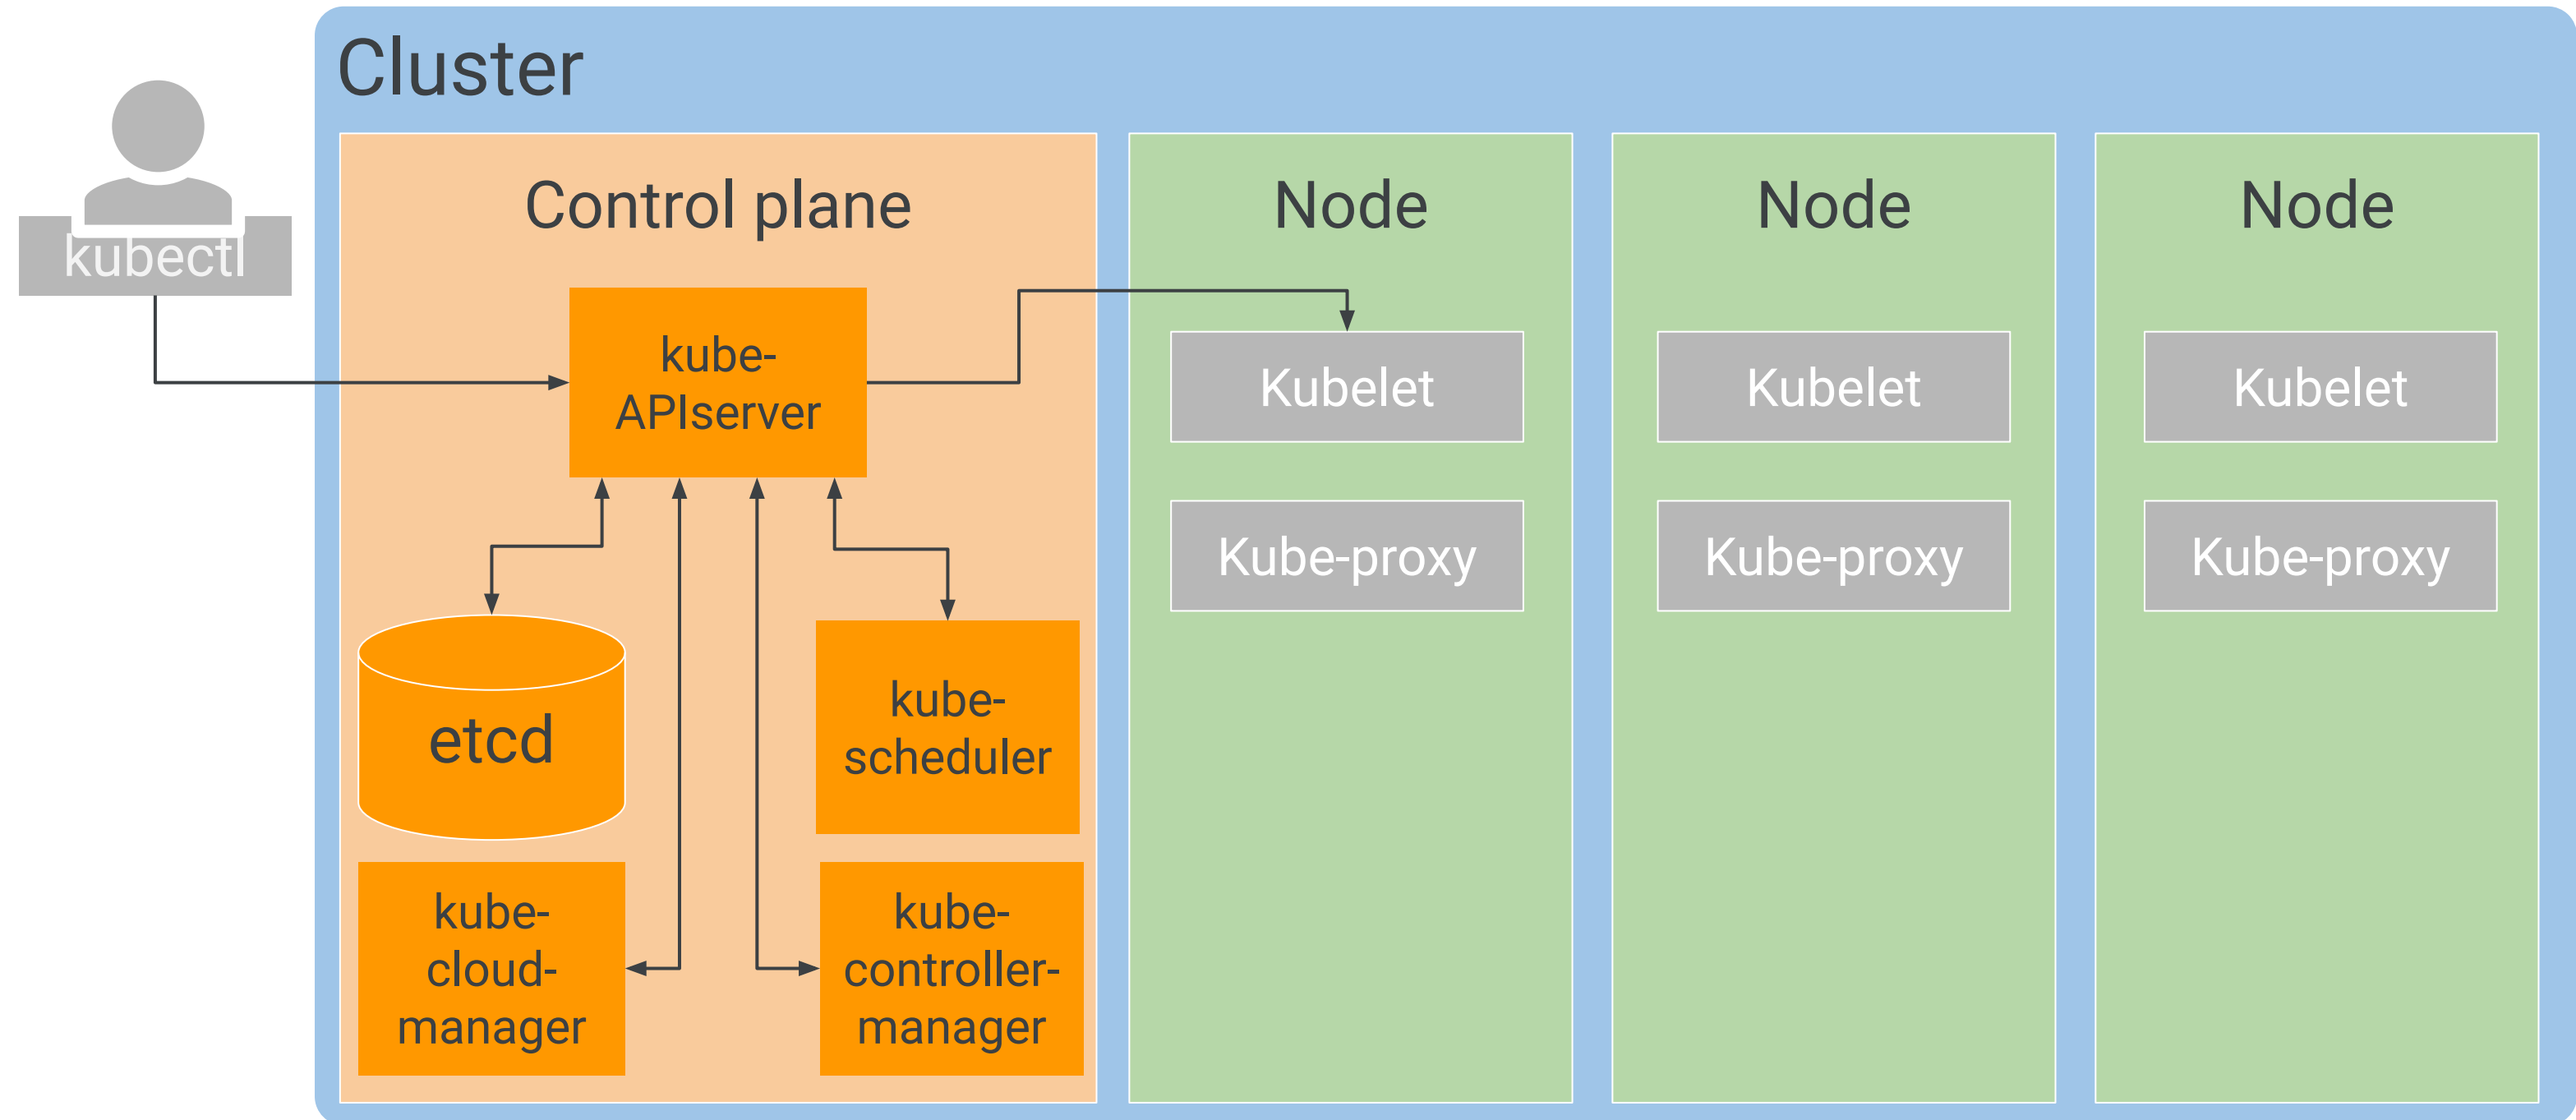




# Containers in a Pod share resources



# Cooperating processes make a Kubernetes cluster work



# GKE: More about nodes

### Machine configuration

**Machine family**

GENERAL-PURPOSE COMPUTE-OPTIMIZED MEMORY-OPTIMIZED GPU

Machine types for common workloads, optimized for cost and flexibility

Series

E2

CPU platform selection based on availability

Machine type

e2-medium (2 vCPU, 4 GB memory)

**vCPU**

1 shared core

**Memory**

4 GB

CPU platform

Automatic

### Machine configuration

**Machine family**

GENERAL-PURPOSE COMPUTE-OPTIMIZED MEMORY-OPTIMIZED GPU

High-performance machine types for compute-intensive workloads

**Series**

C2

Powered by Intel Cascade Lake CPU platform

Machine type

c2-standard-4

4 vCPU, 16 GB memory

c2-standard-8

8 vCPU, 32 GB memory

c2-standard-16

16 vCPU, 64 GB memory

c2-standard-30

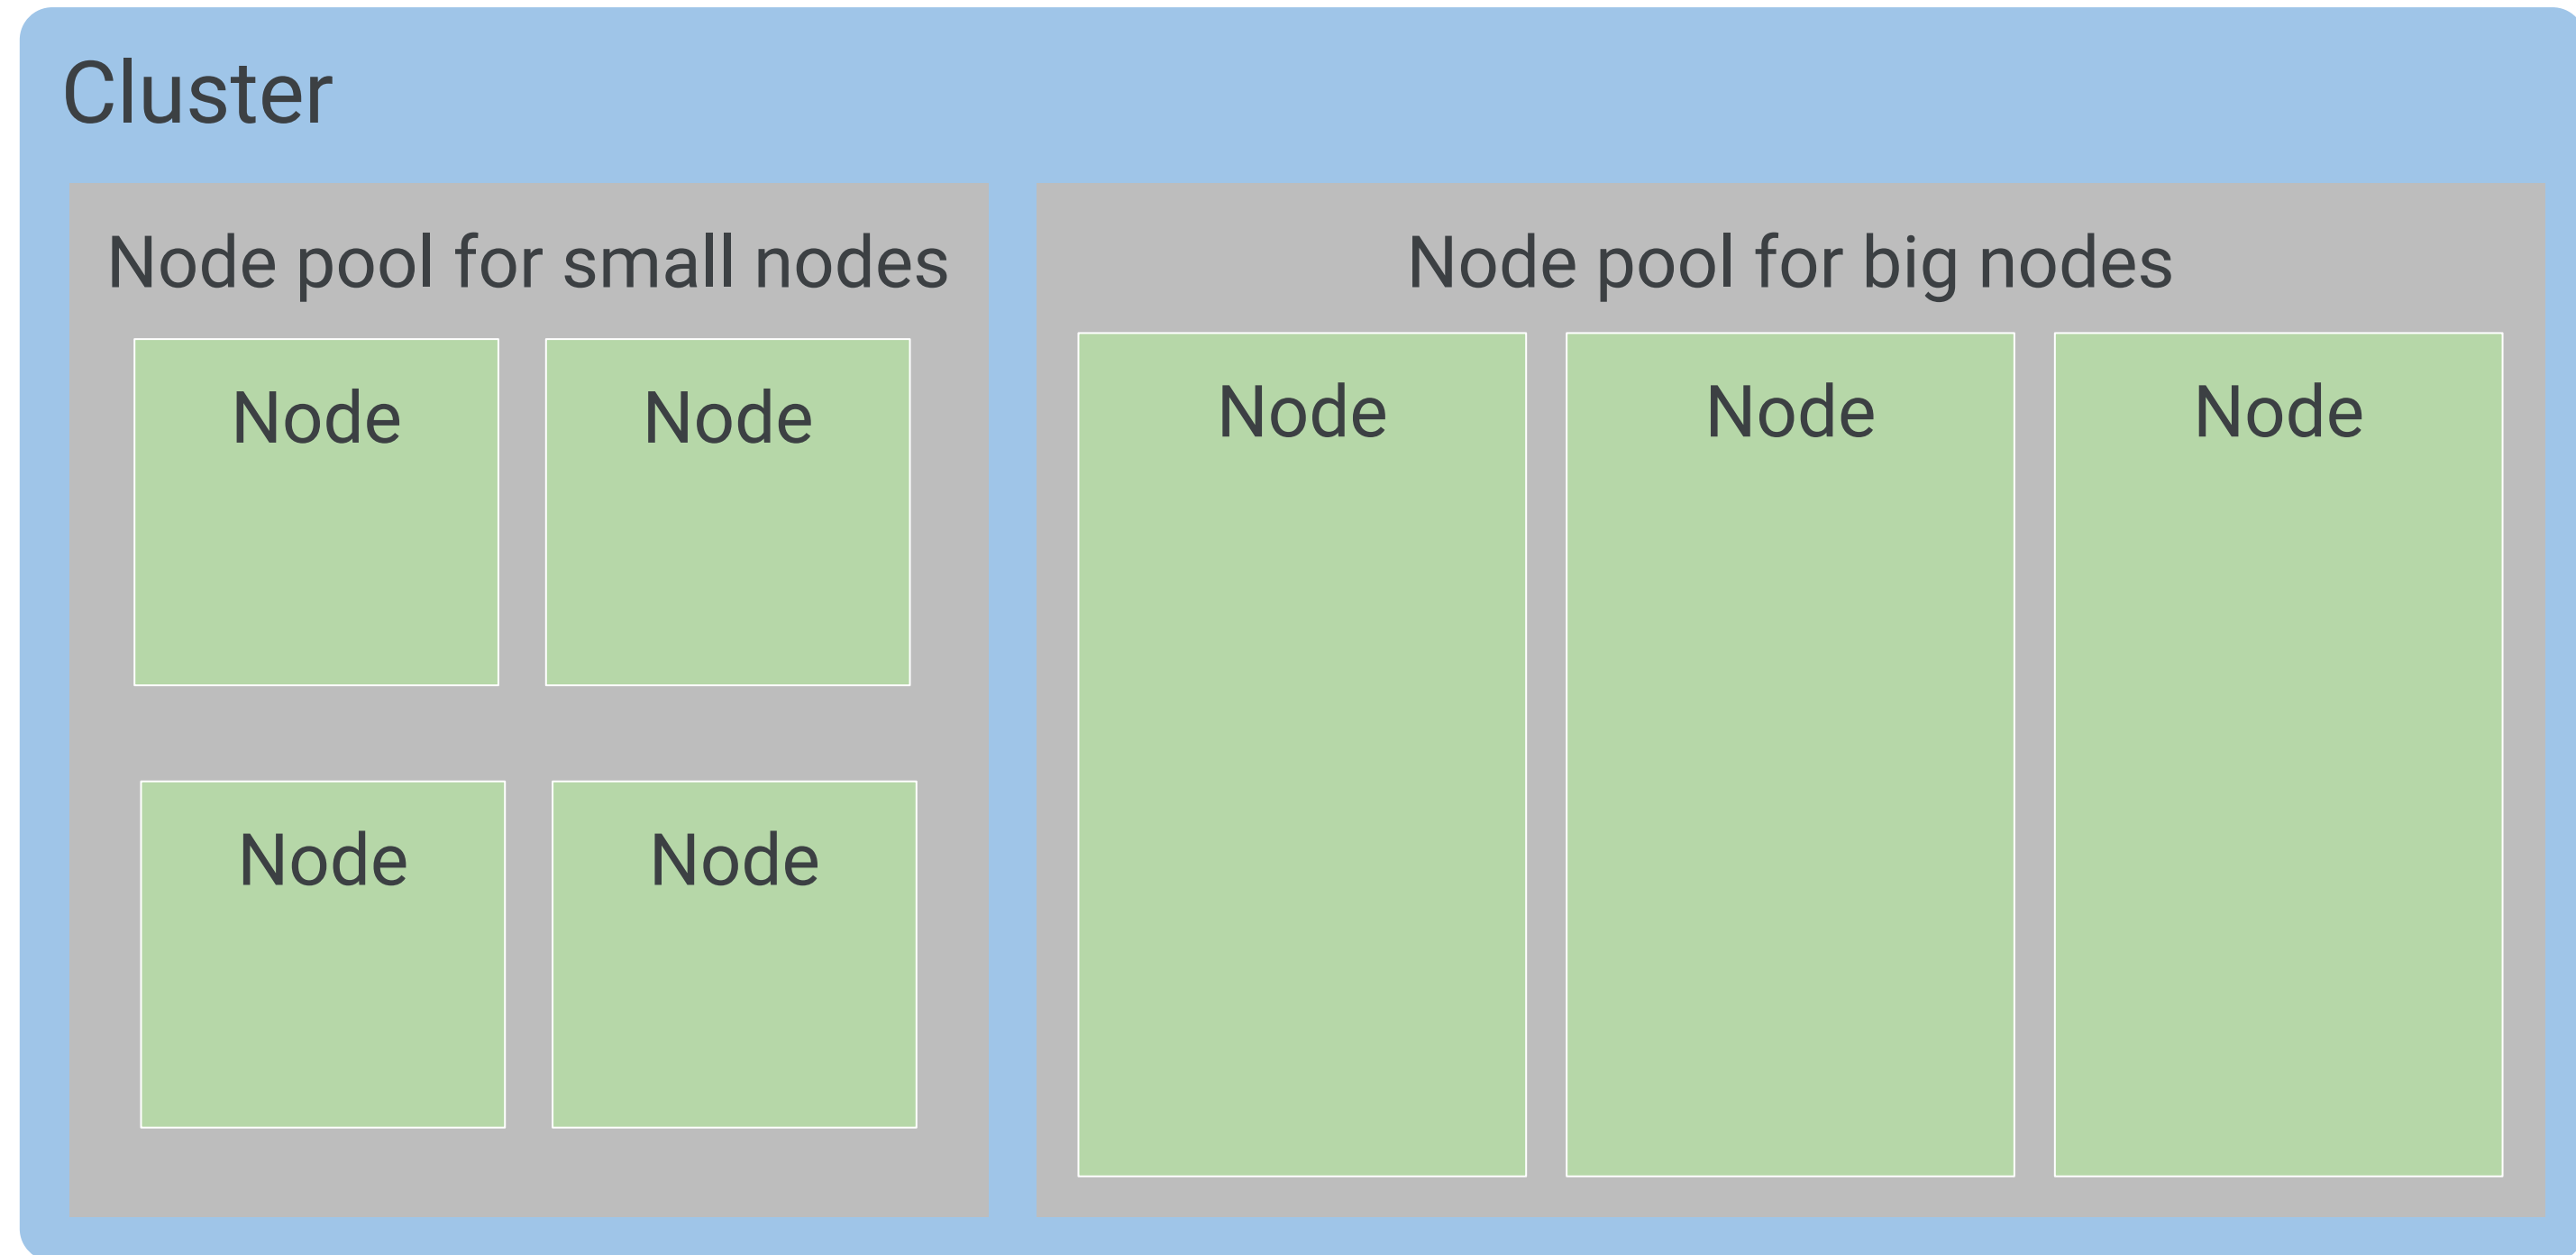
30 vCPU, 120 GB memory

c2-standard-60

60 vCPU, 240 GB memory

+ ADD GPU

# Use node pools to manage different kinds of nodes



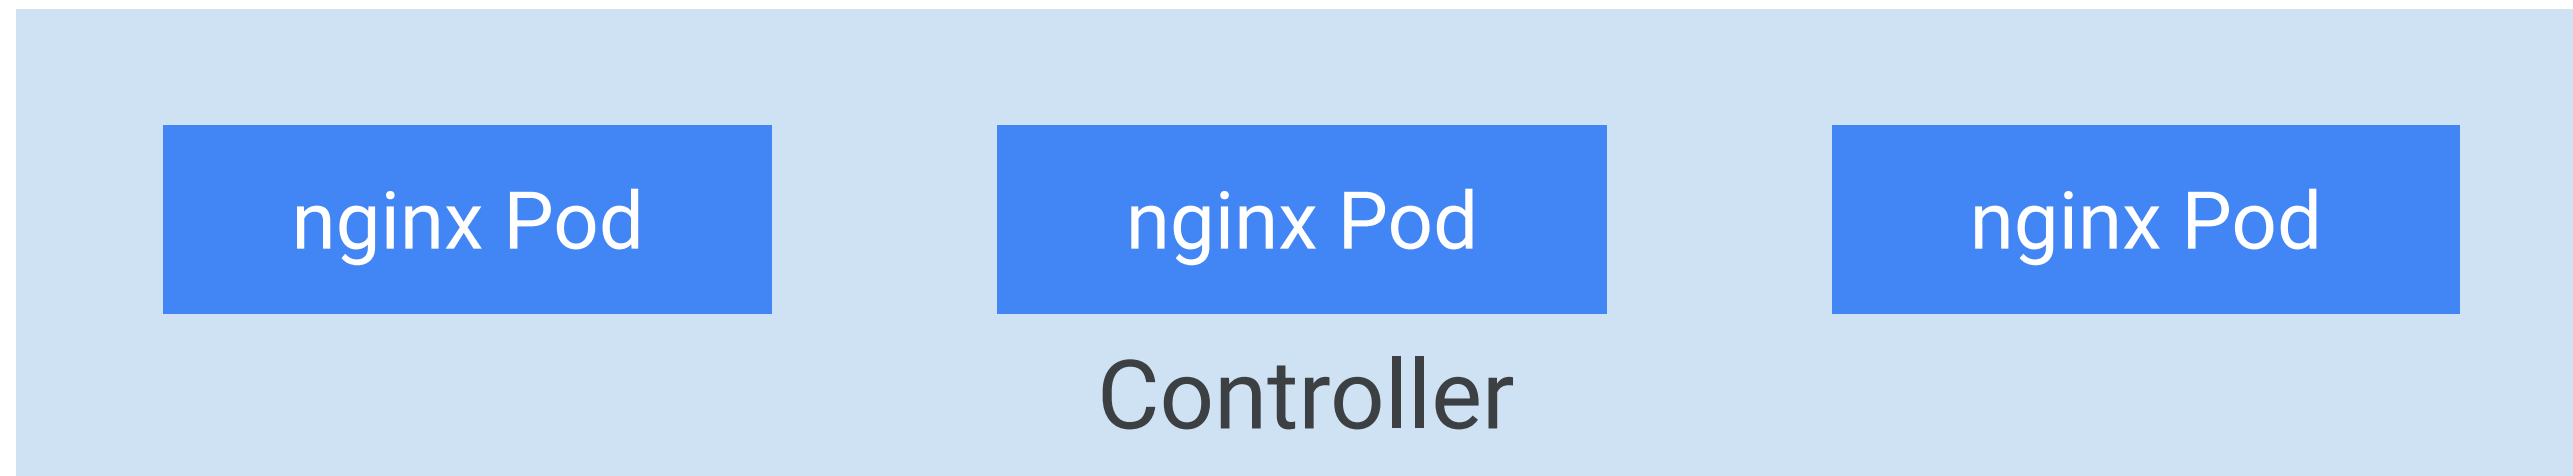
---

# Objects are defined in a YAML file

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:latest
```

---

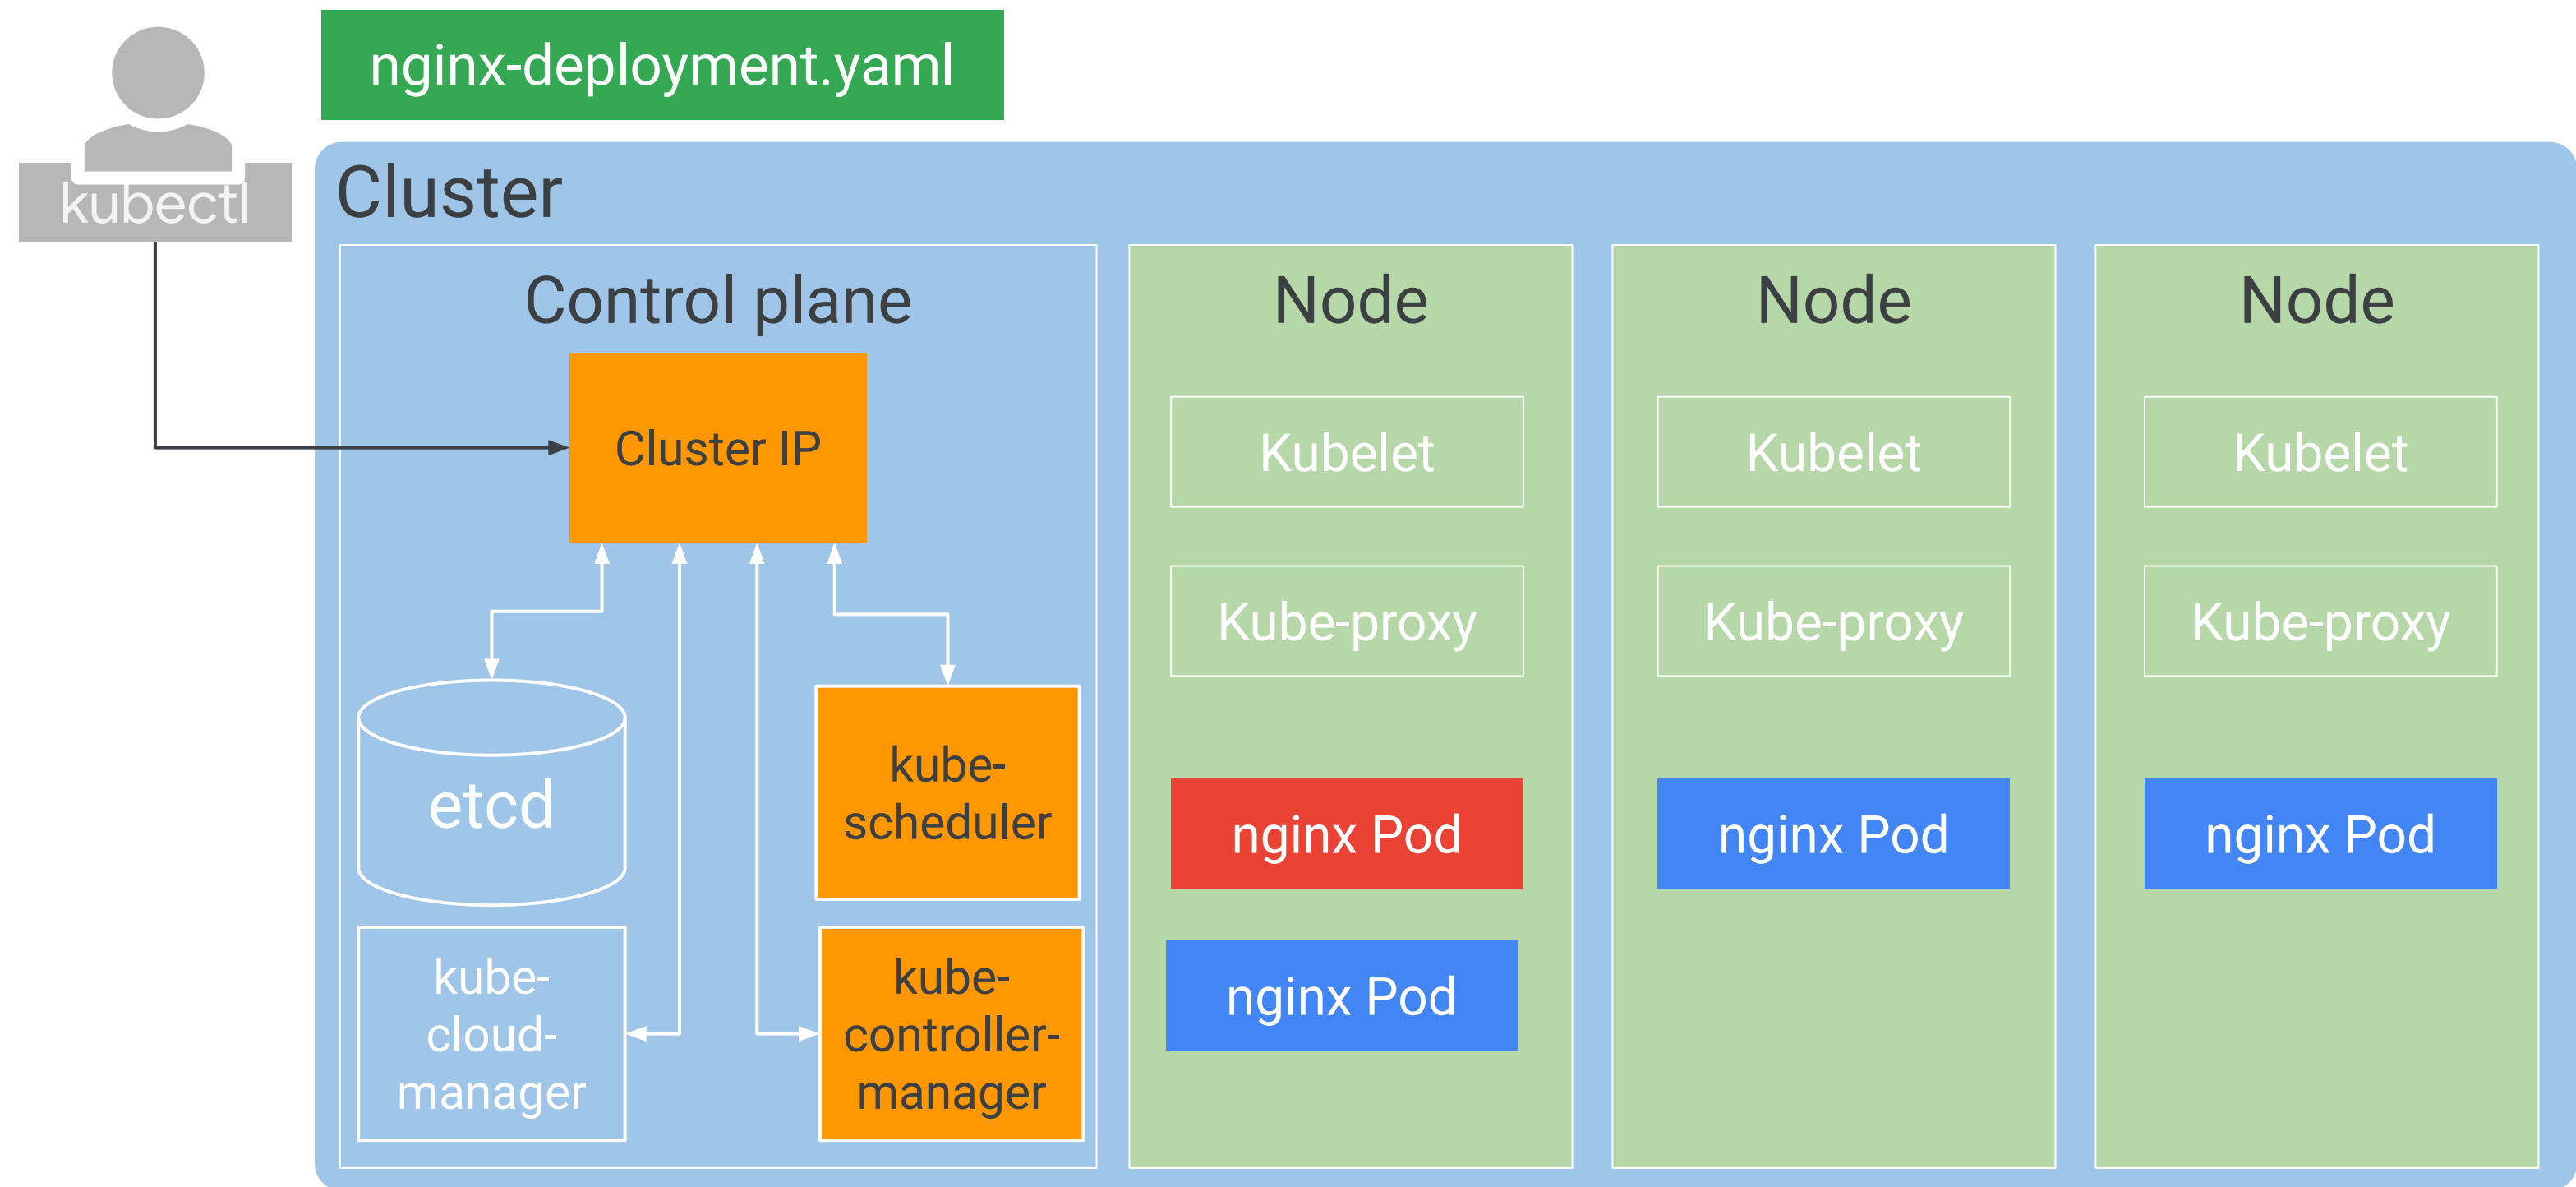
# Pods and Controller Objects



Controller object  
types

- Deployment
- StatefulSet
- DaemonSet
- Job

# A Deployment maintains the desired state



---

# Deployments ensure that sets of Pods are running

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
```






---

## Introducing Anthos

# Hybrid Cloud Overview

- Companies would like to adopt the cloud in their own pace, alongside their on premise.
- The hybrid environment empowers companies to pick and choose the best of both worlds, and create bespoke infrastructure



# Hybrid environment wishlist

**Write once,  
deploy in any  
cloud**

**Accelerate  
developer  
velocity**

**Consistency  
across  
environments**

**Interoperability  
with legacy  
workloads**

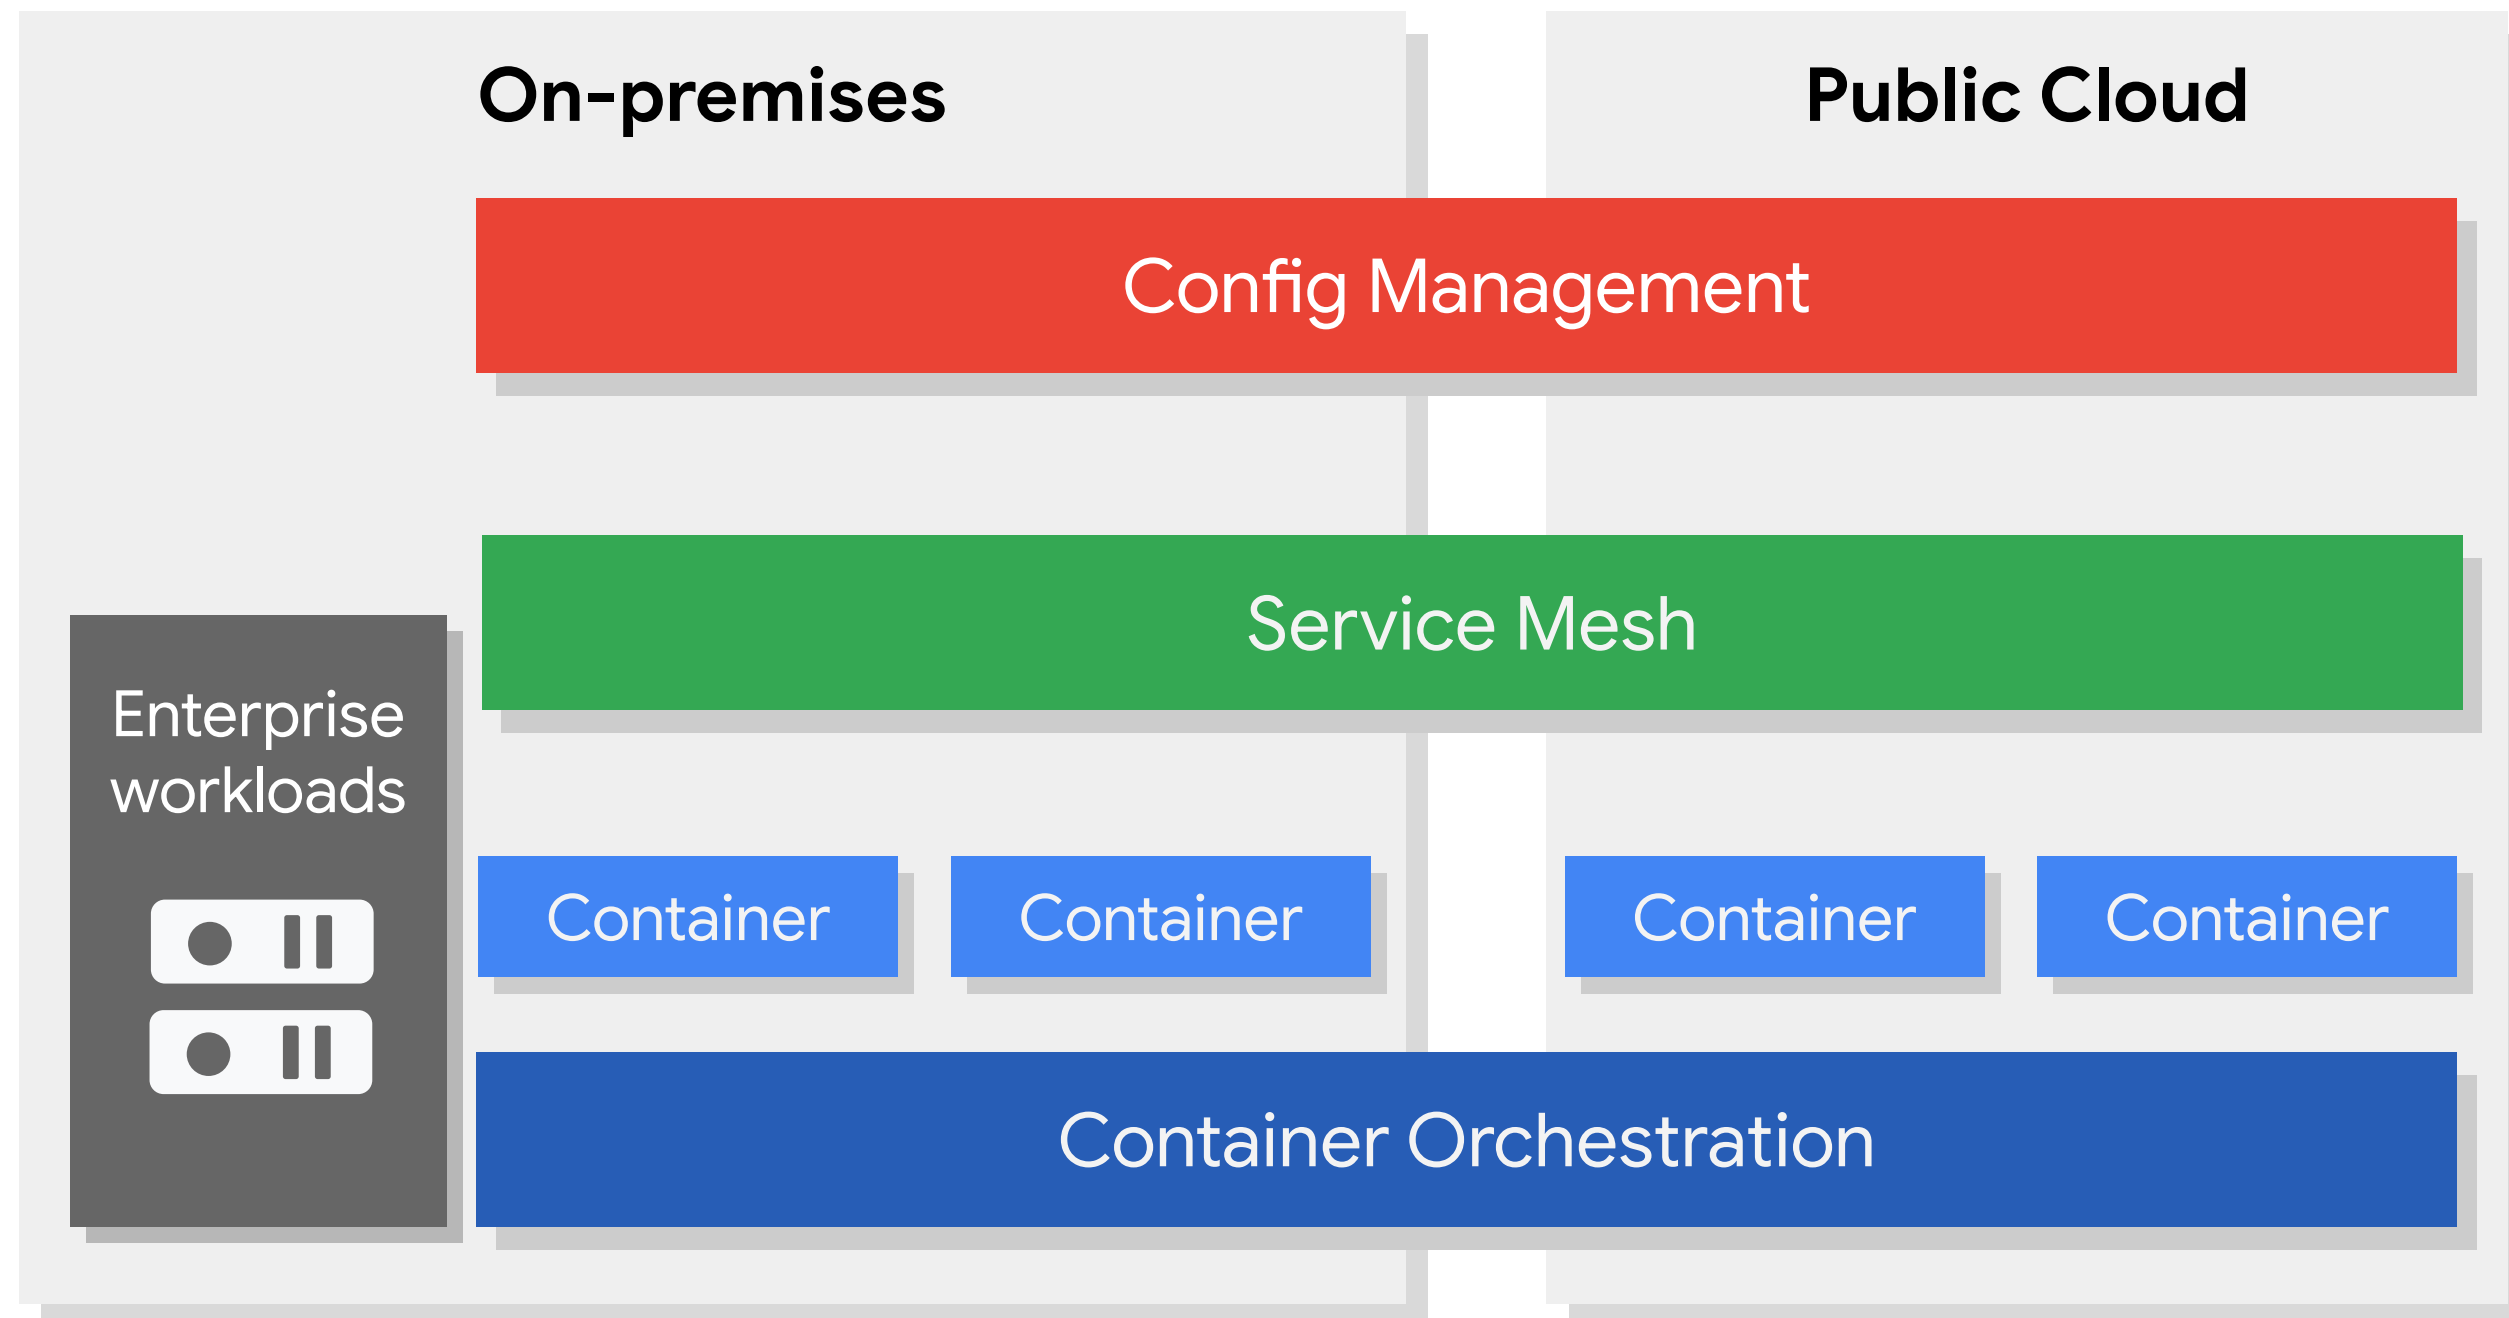
**Increased  
observability  
and SLO**

**Decoupling  
across critical  
components**

**Increased  
workload  
mobility**

**Avoid vendor  
lock in**

# Bringing it together



# Container Orchestration



## Google Kubernetes Engine


- Operates within GCP
- Managed, production-ready environment for deploying containerized applications
- Operate Seamlessly with High Availability and SLA
- Runs Certified Kubernetes ensuring portability across clouds and on-premises.
- Auto node repair, auto upgrade, auto scaling
- Regional clusters for high availability with multiple masters, node storage replication across 3 zones



## Anthos clusters

- Operates On-Premises
- Turn-key, production-grade, conformant Kubernetes with best-practice configuration
- Easy upgrade path to the latest Kubernetes releases that have been validated and tested by Google
- Access to Container services on GCP such as Cloud Build, Container Registry, Audit Logging, and more.
- Integration with Istio, Knative, Marketplace Solutions
- Consistent Kubernetes versions and experience across environments

# Service Mesh




**Anthos Service Mesh**

---



- Based on Istio
- Automated/managed installation
- Automatic upgrades
- Managed control plane



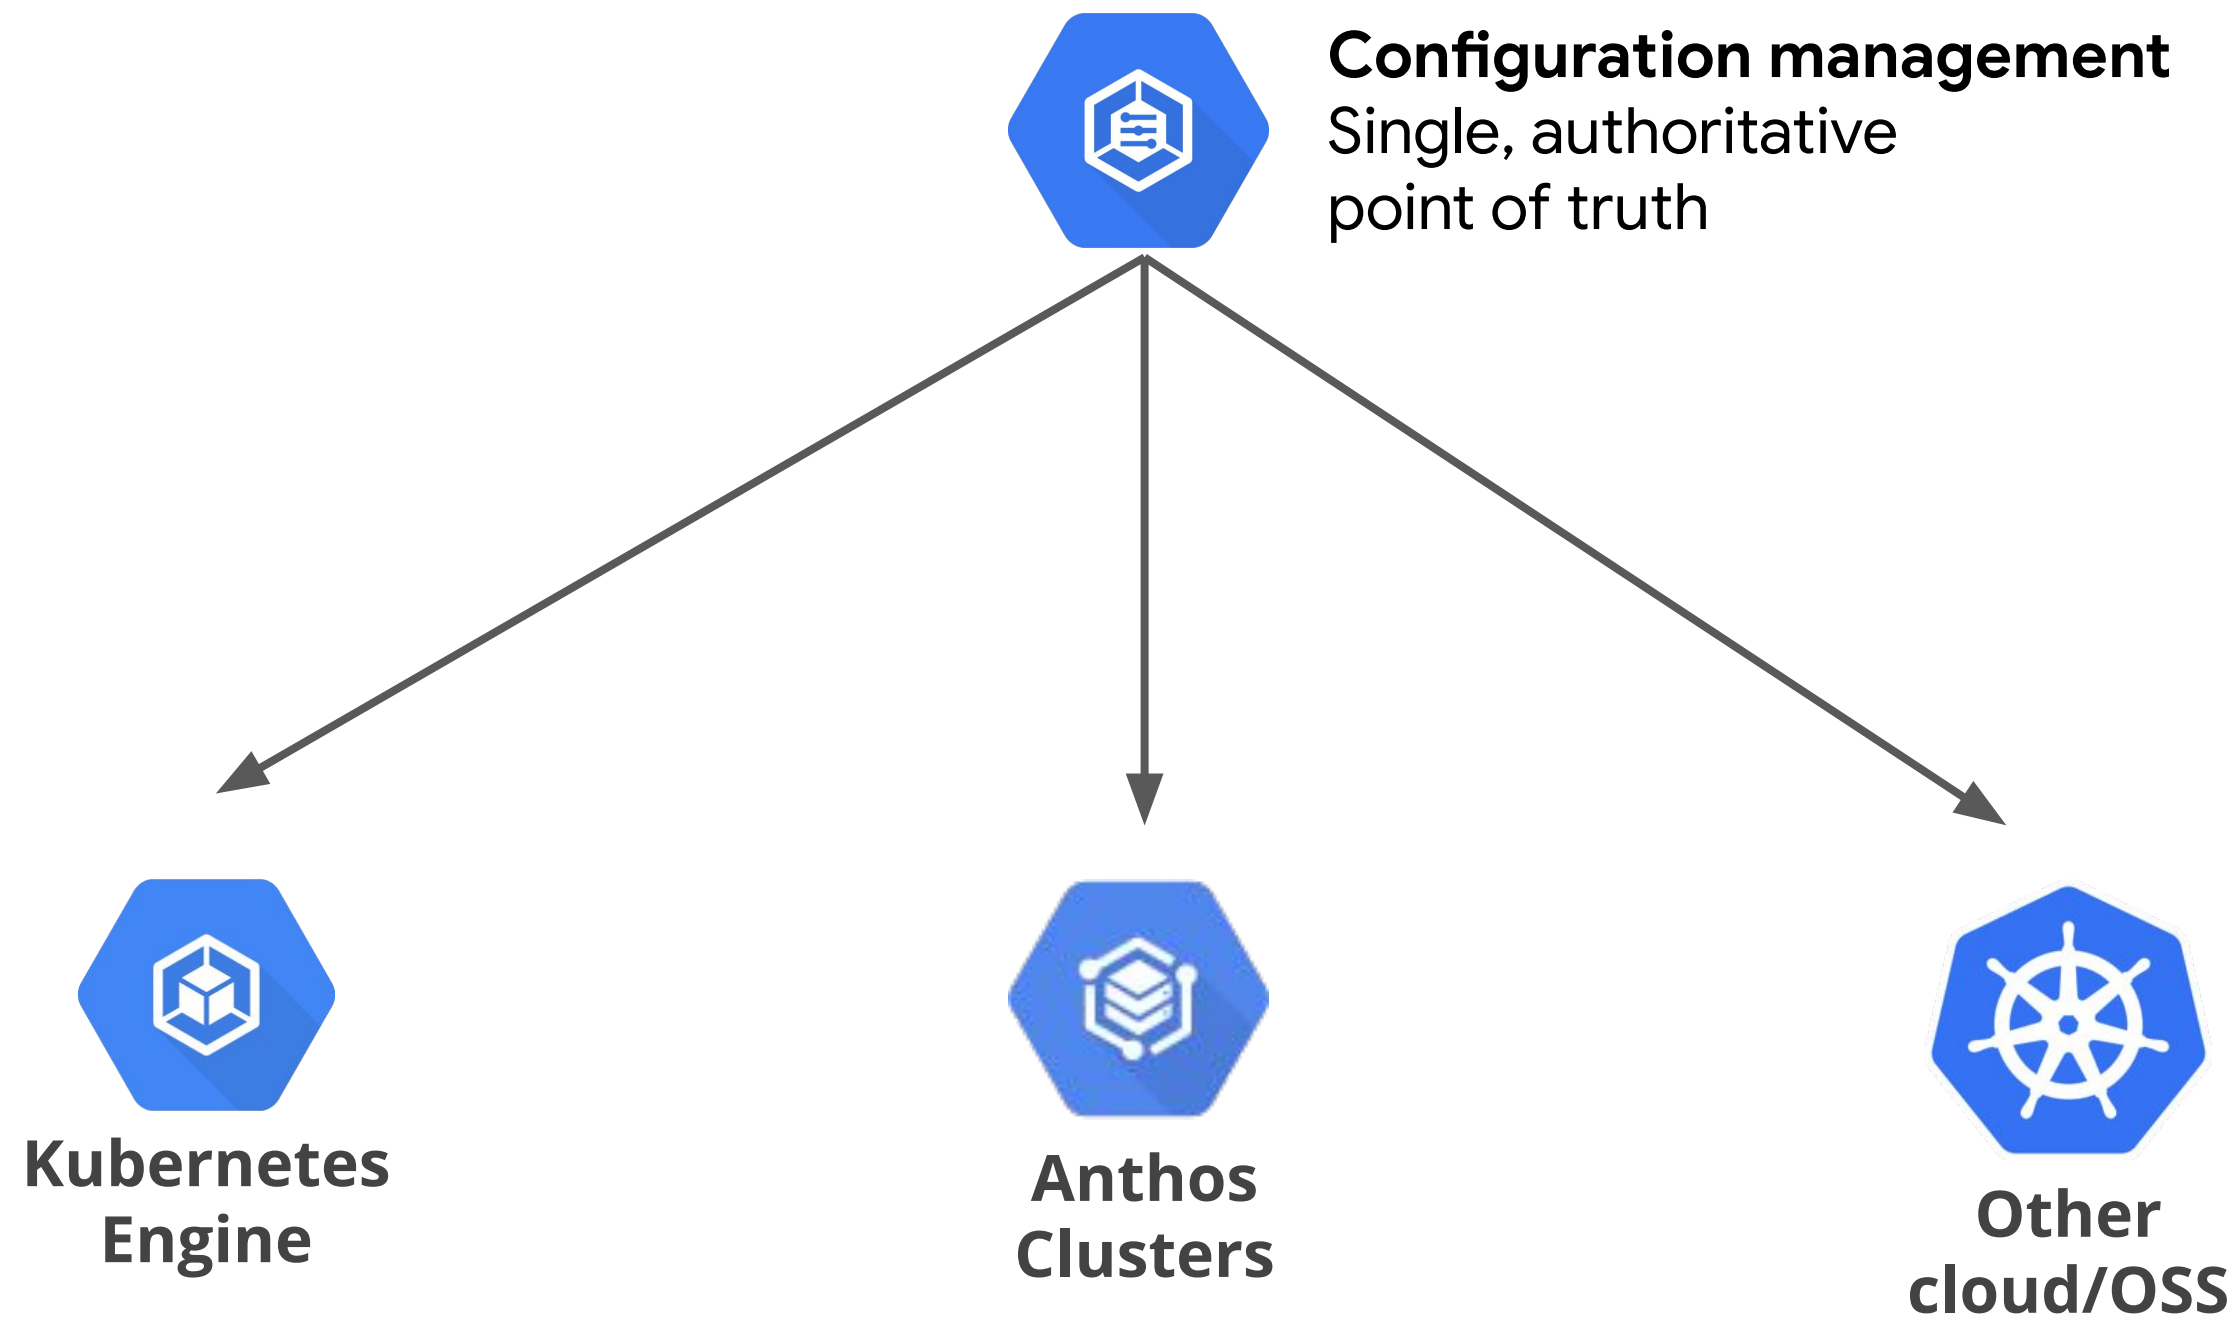
**Istio OSS**

---




- Self managed open source version
- Fully integrates with Anthos Service Mesh
- Creates a seamless interoperability between the services hosted in different environments

# Multicloud Management



# Observability







# Anthos

Anthos is a modern application management platform that provides a unified model for computing, networking, and even service management across clouds and data centers.




Kubernetes  
Engine



Anthos GKE



Anthos Config  
Management




Anthos Service  
Mesh



Cloud Logging  
& Monitoring



Ingress for  
Anthos



Migrate  
for Anthos



Cloud Run




Cloud Build



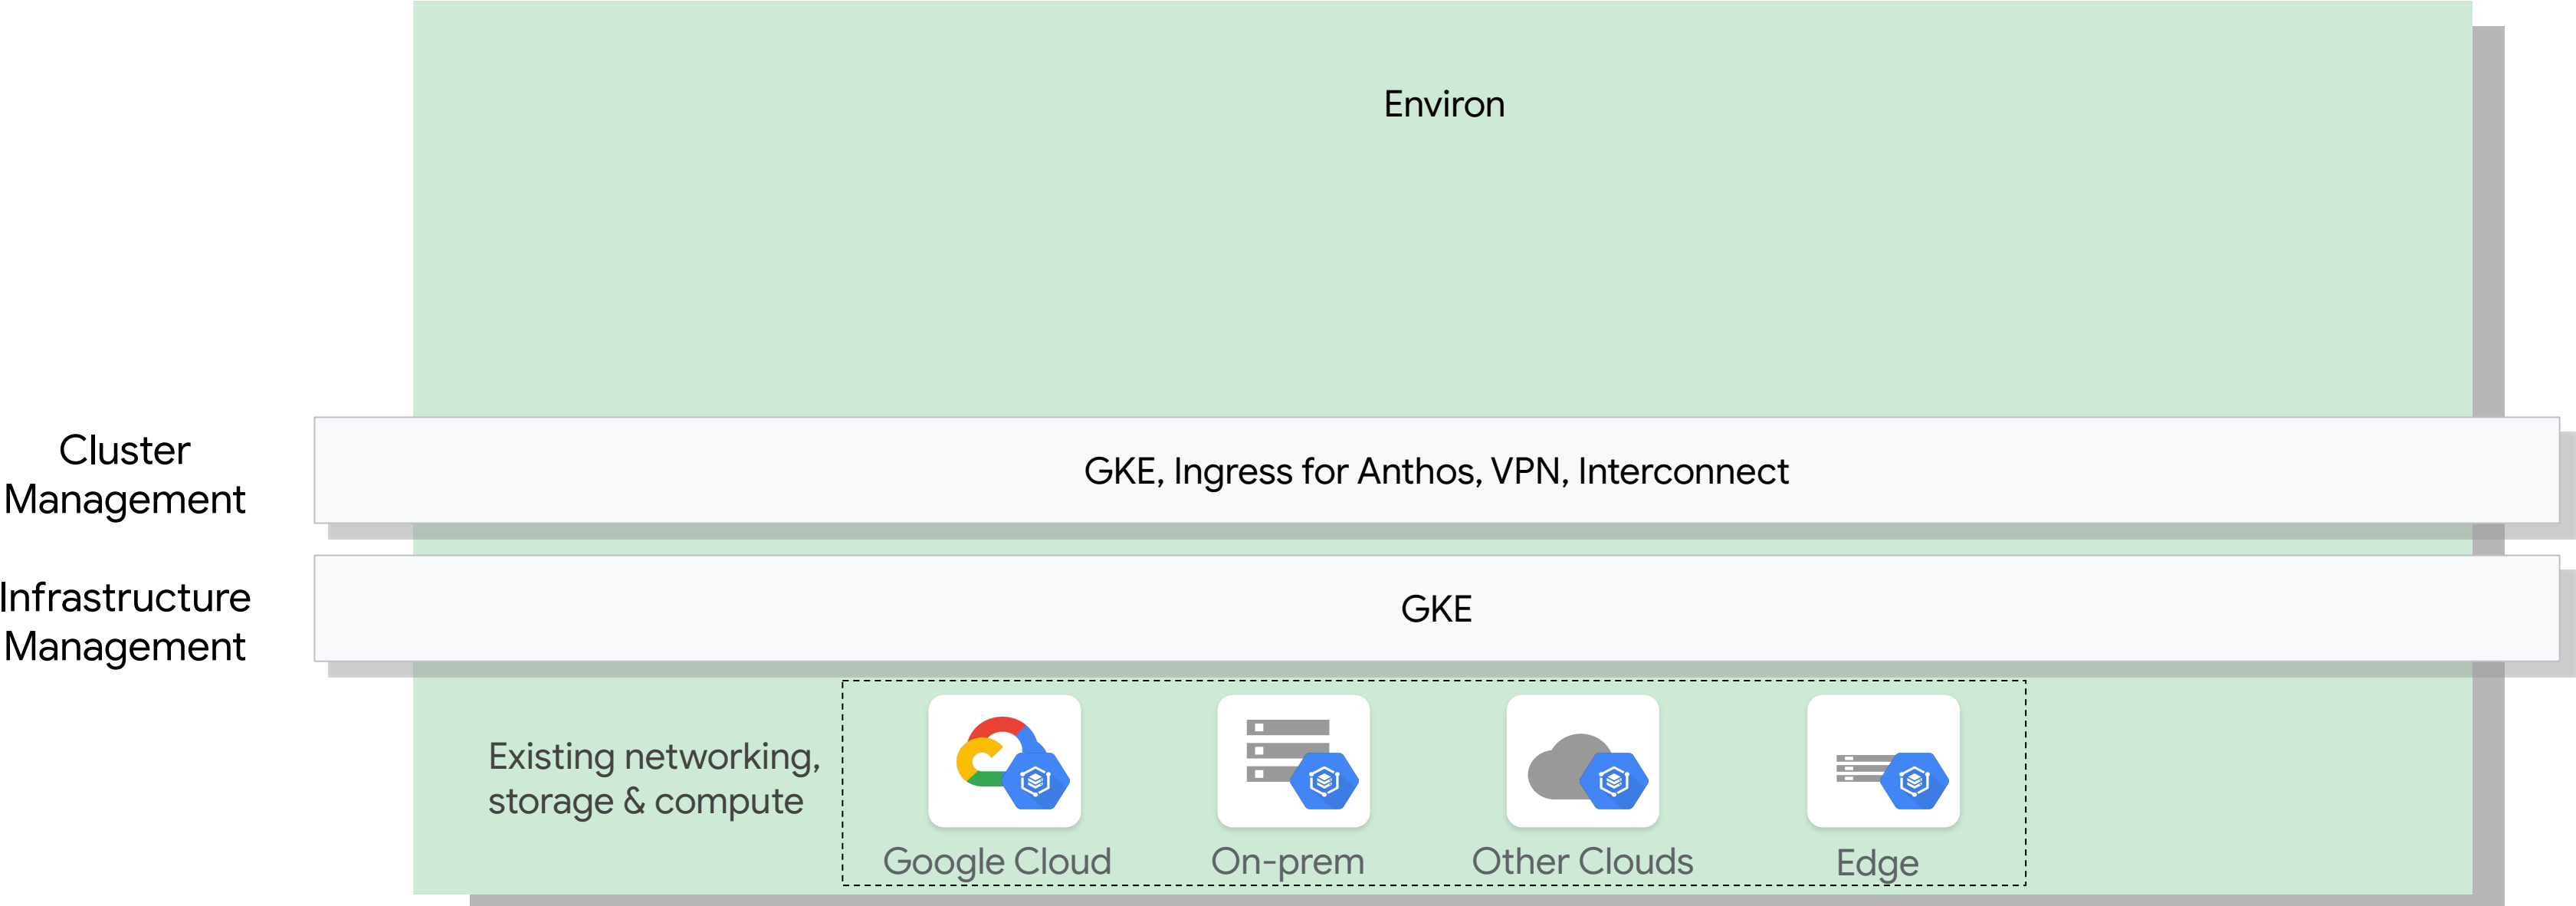
GCP Marketplace  
for Anthos

The technology stack is built on consistent set of APIs based on open-source technologies which empowers developers and operators with a single methodology that applies to on premise, GCP and other cloud providers

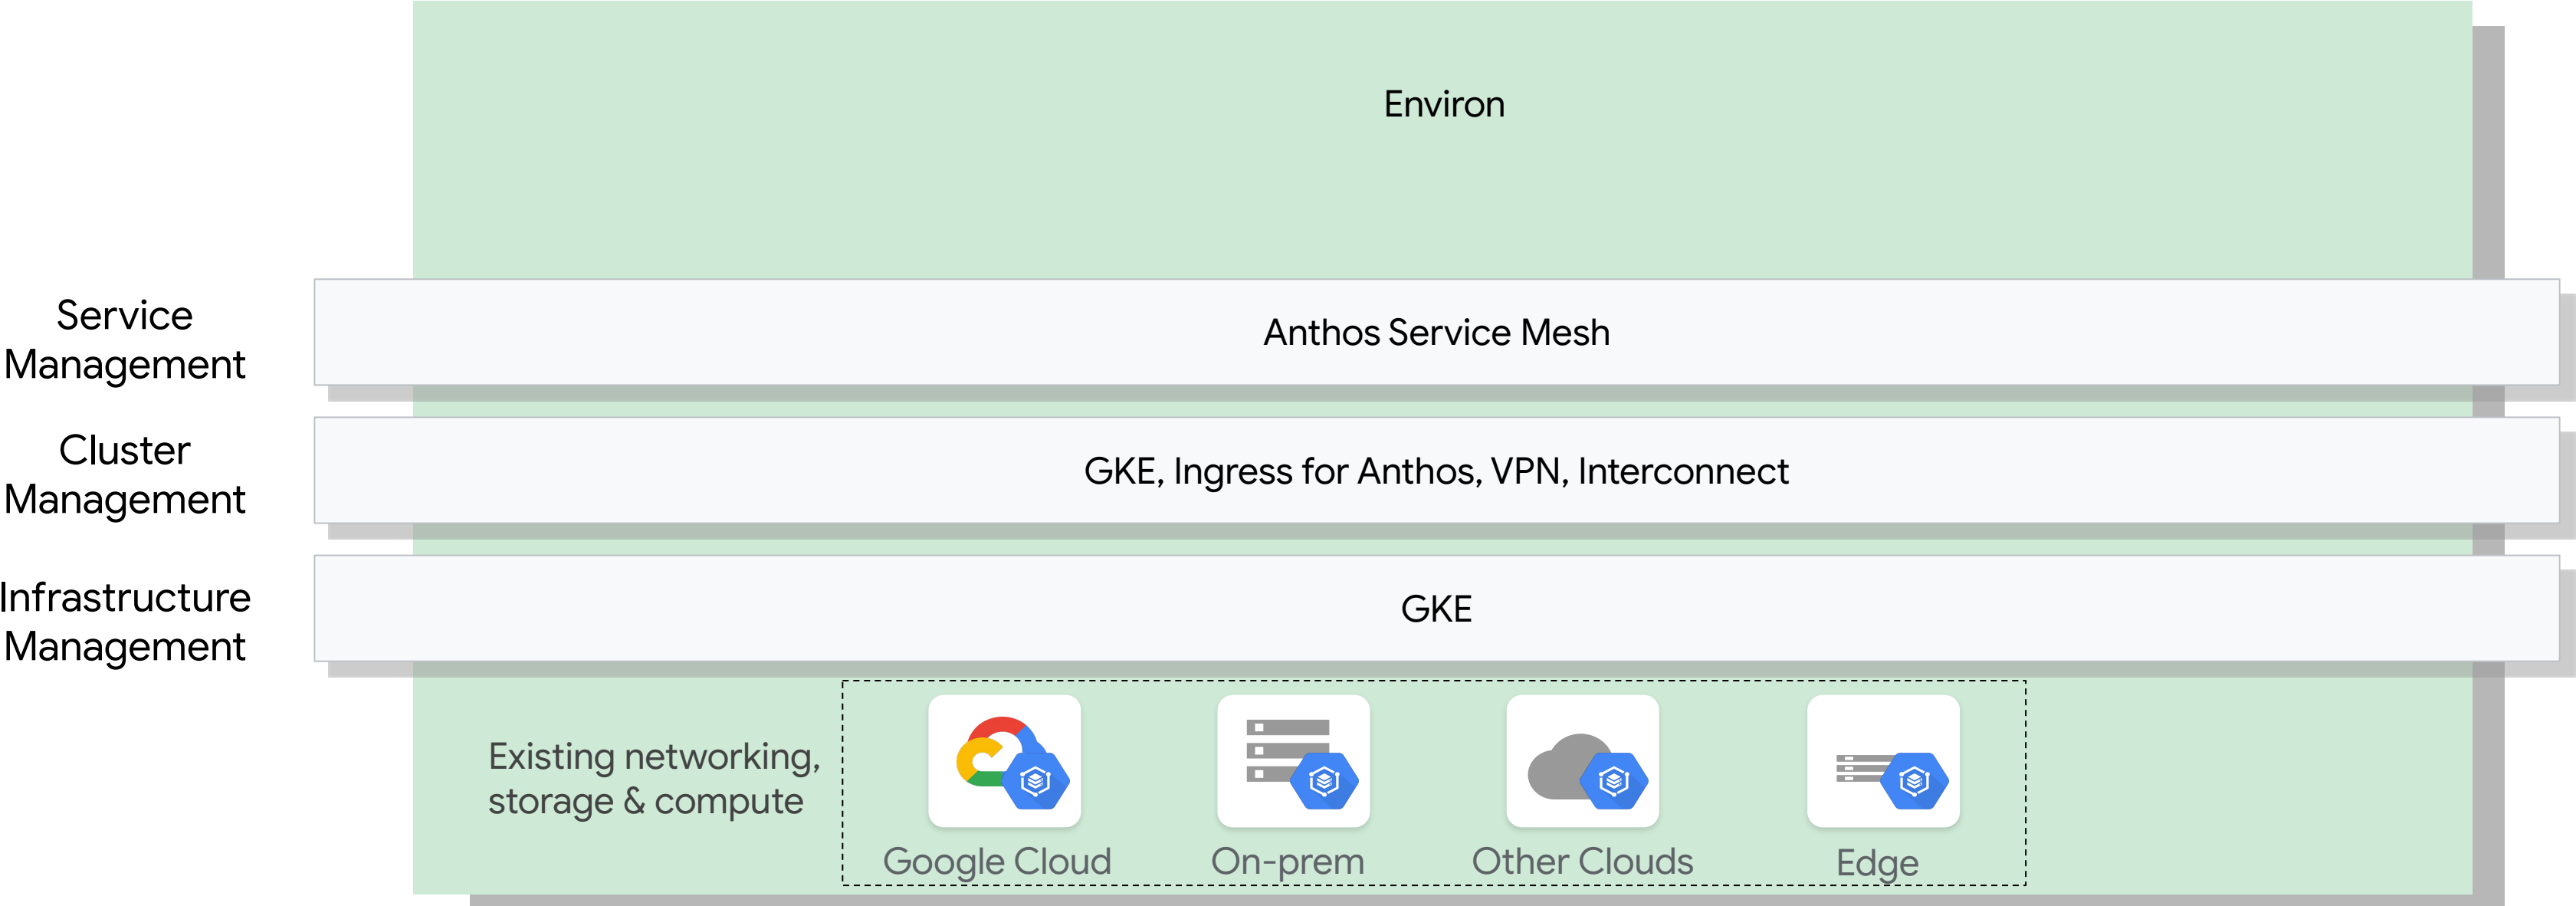
# Anthos technical overview



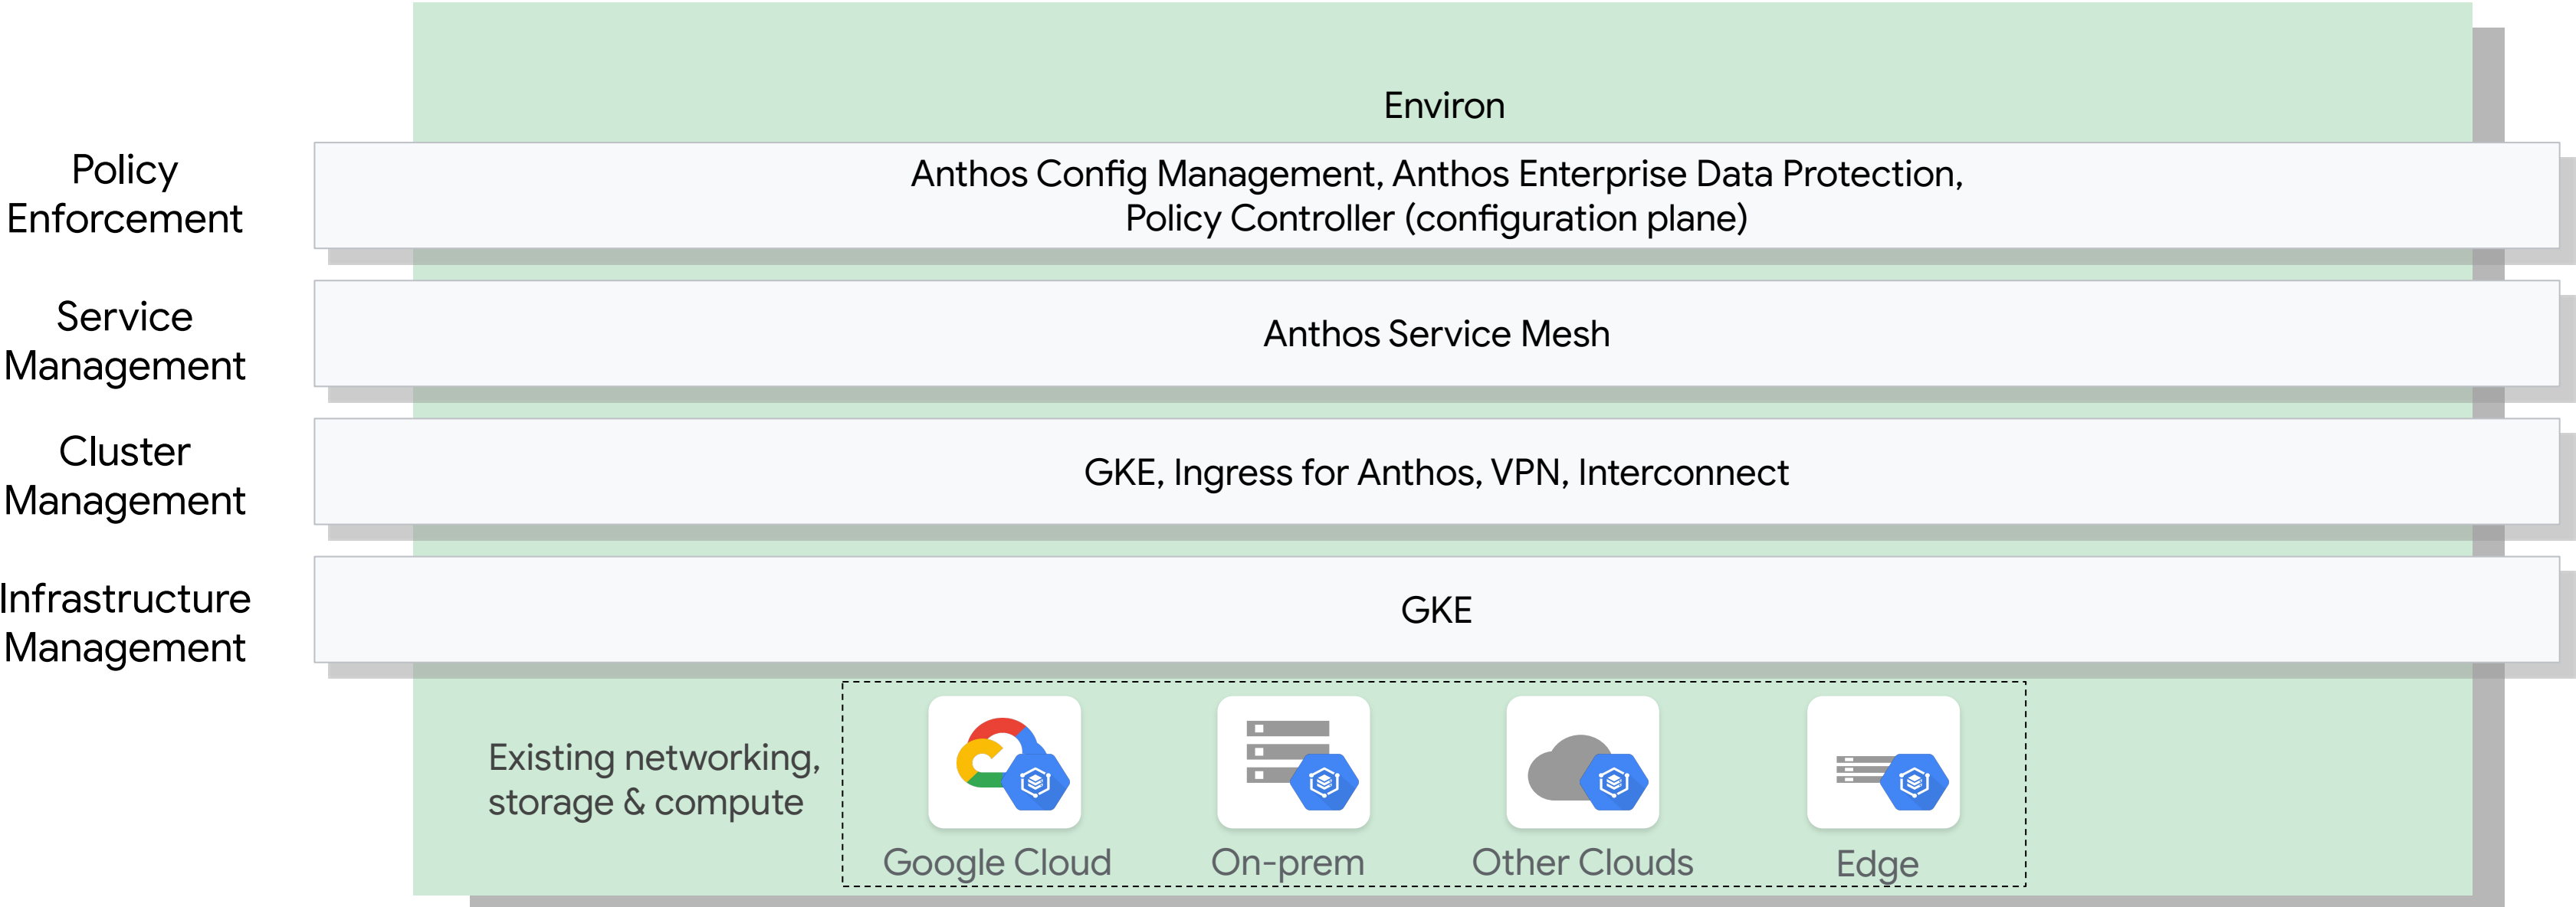
# Anthos technical overview



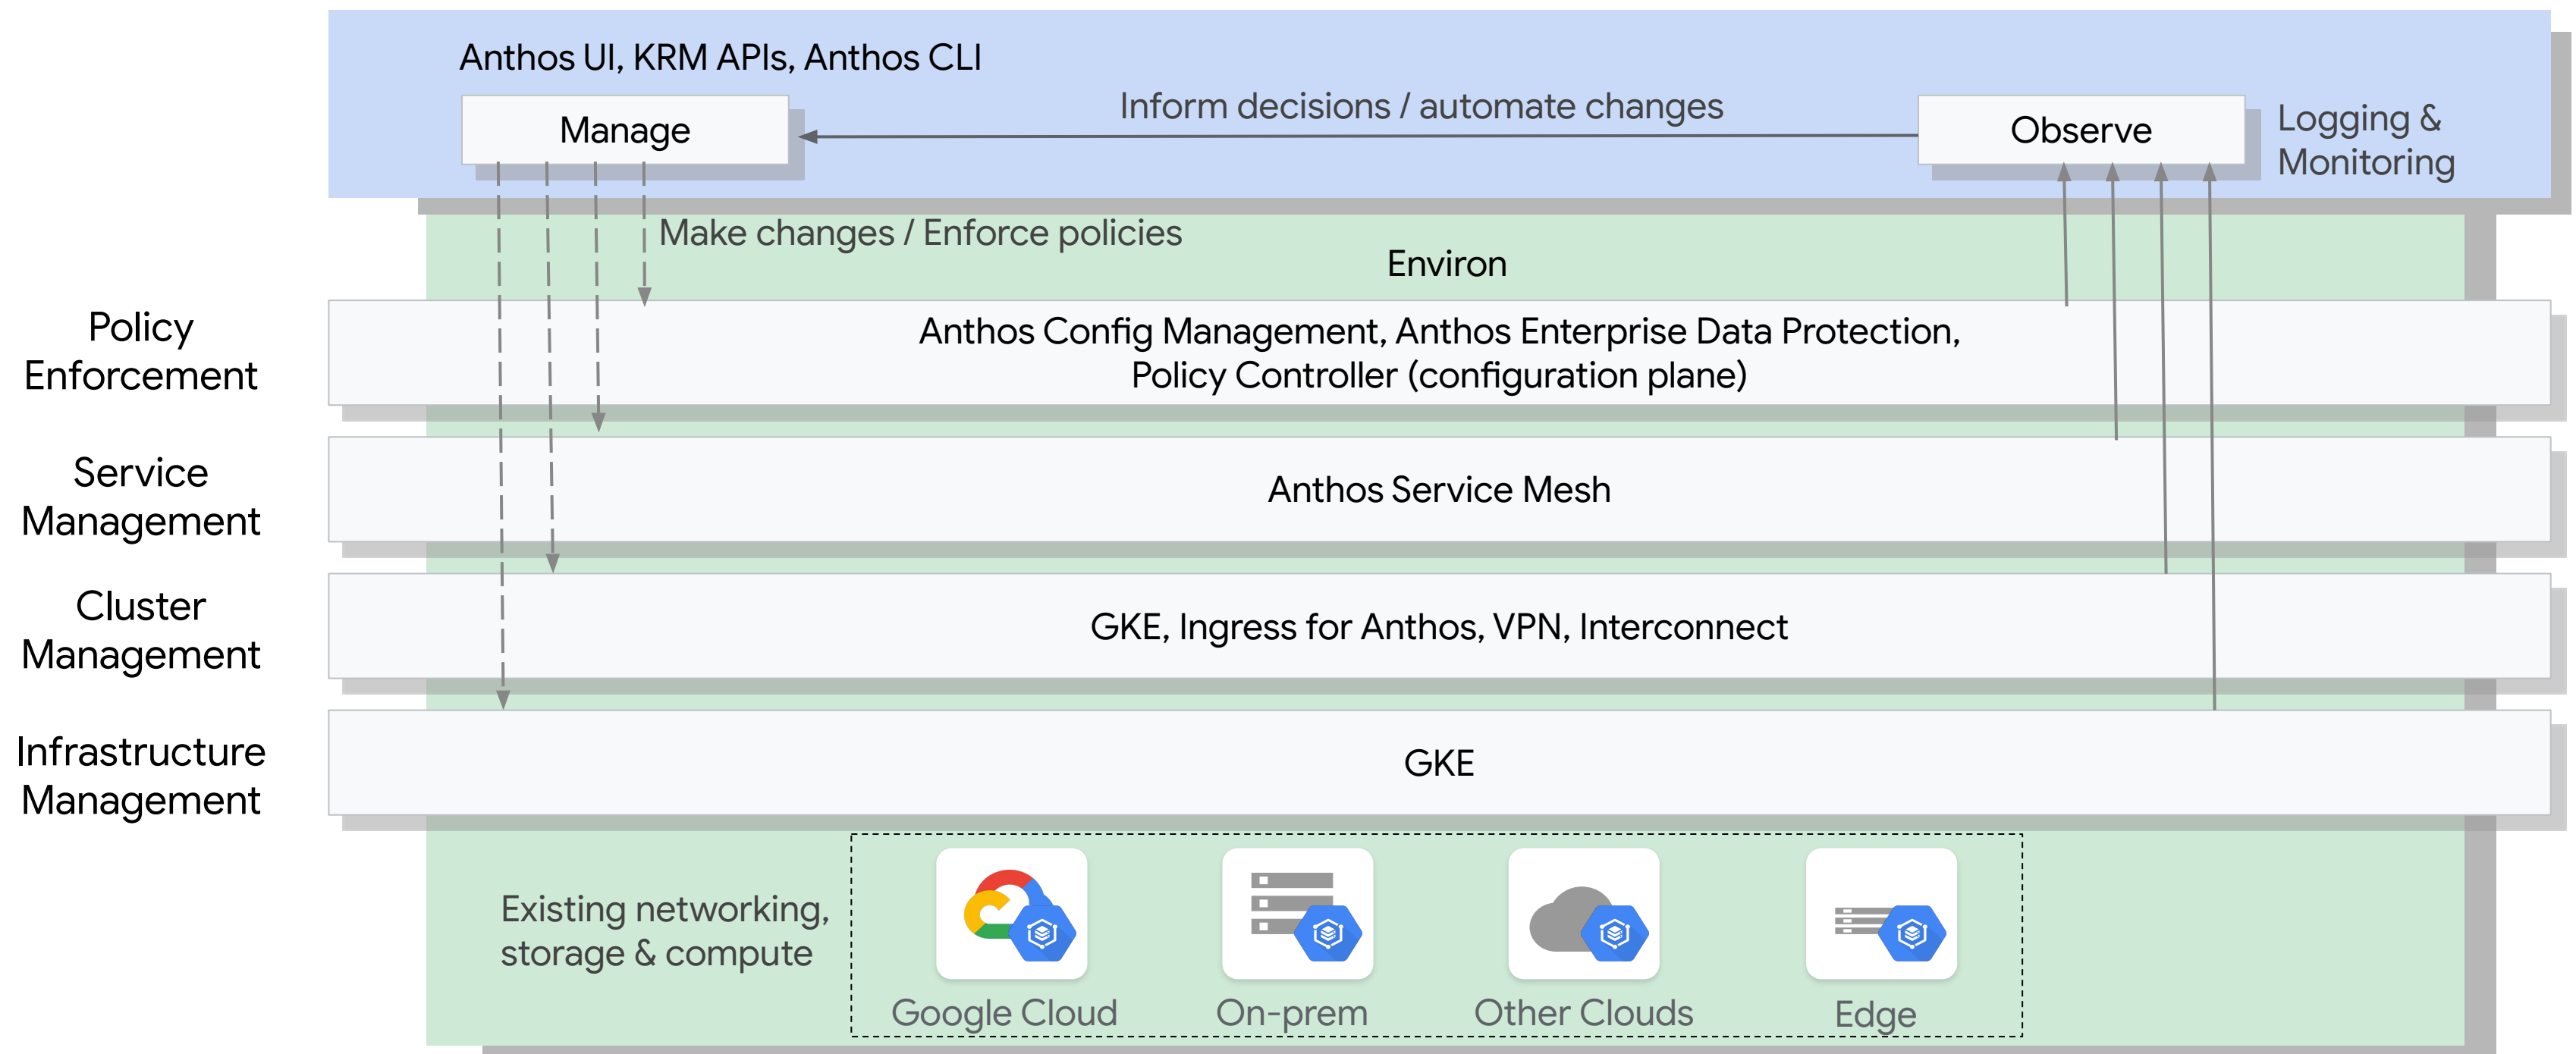
# Anthos technical overview



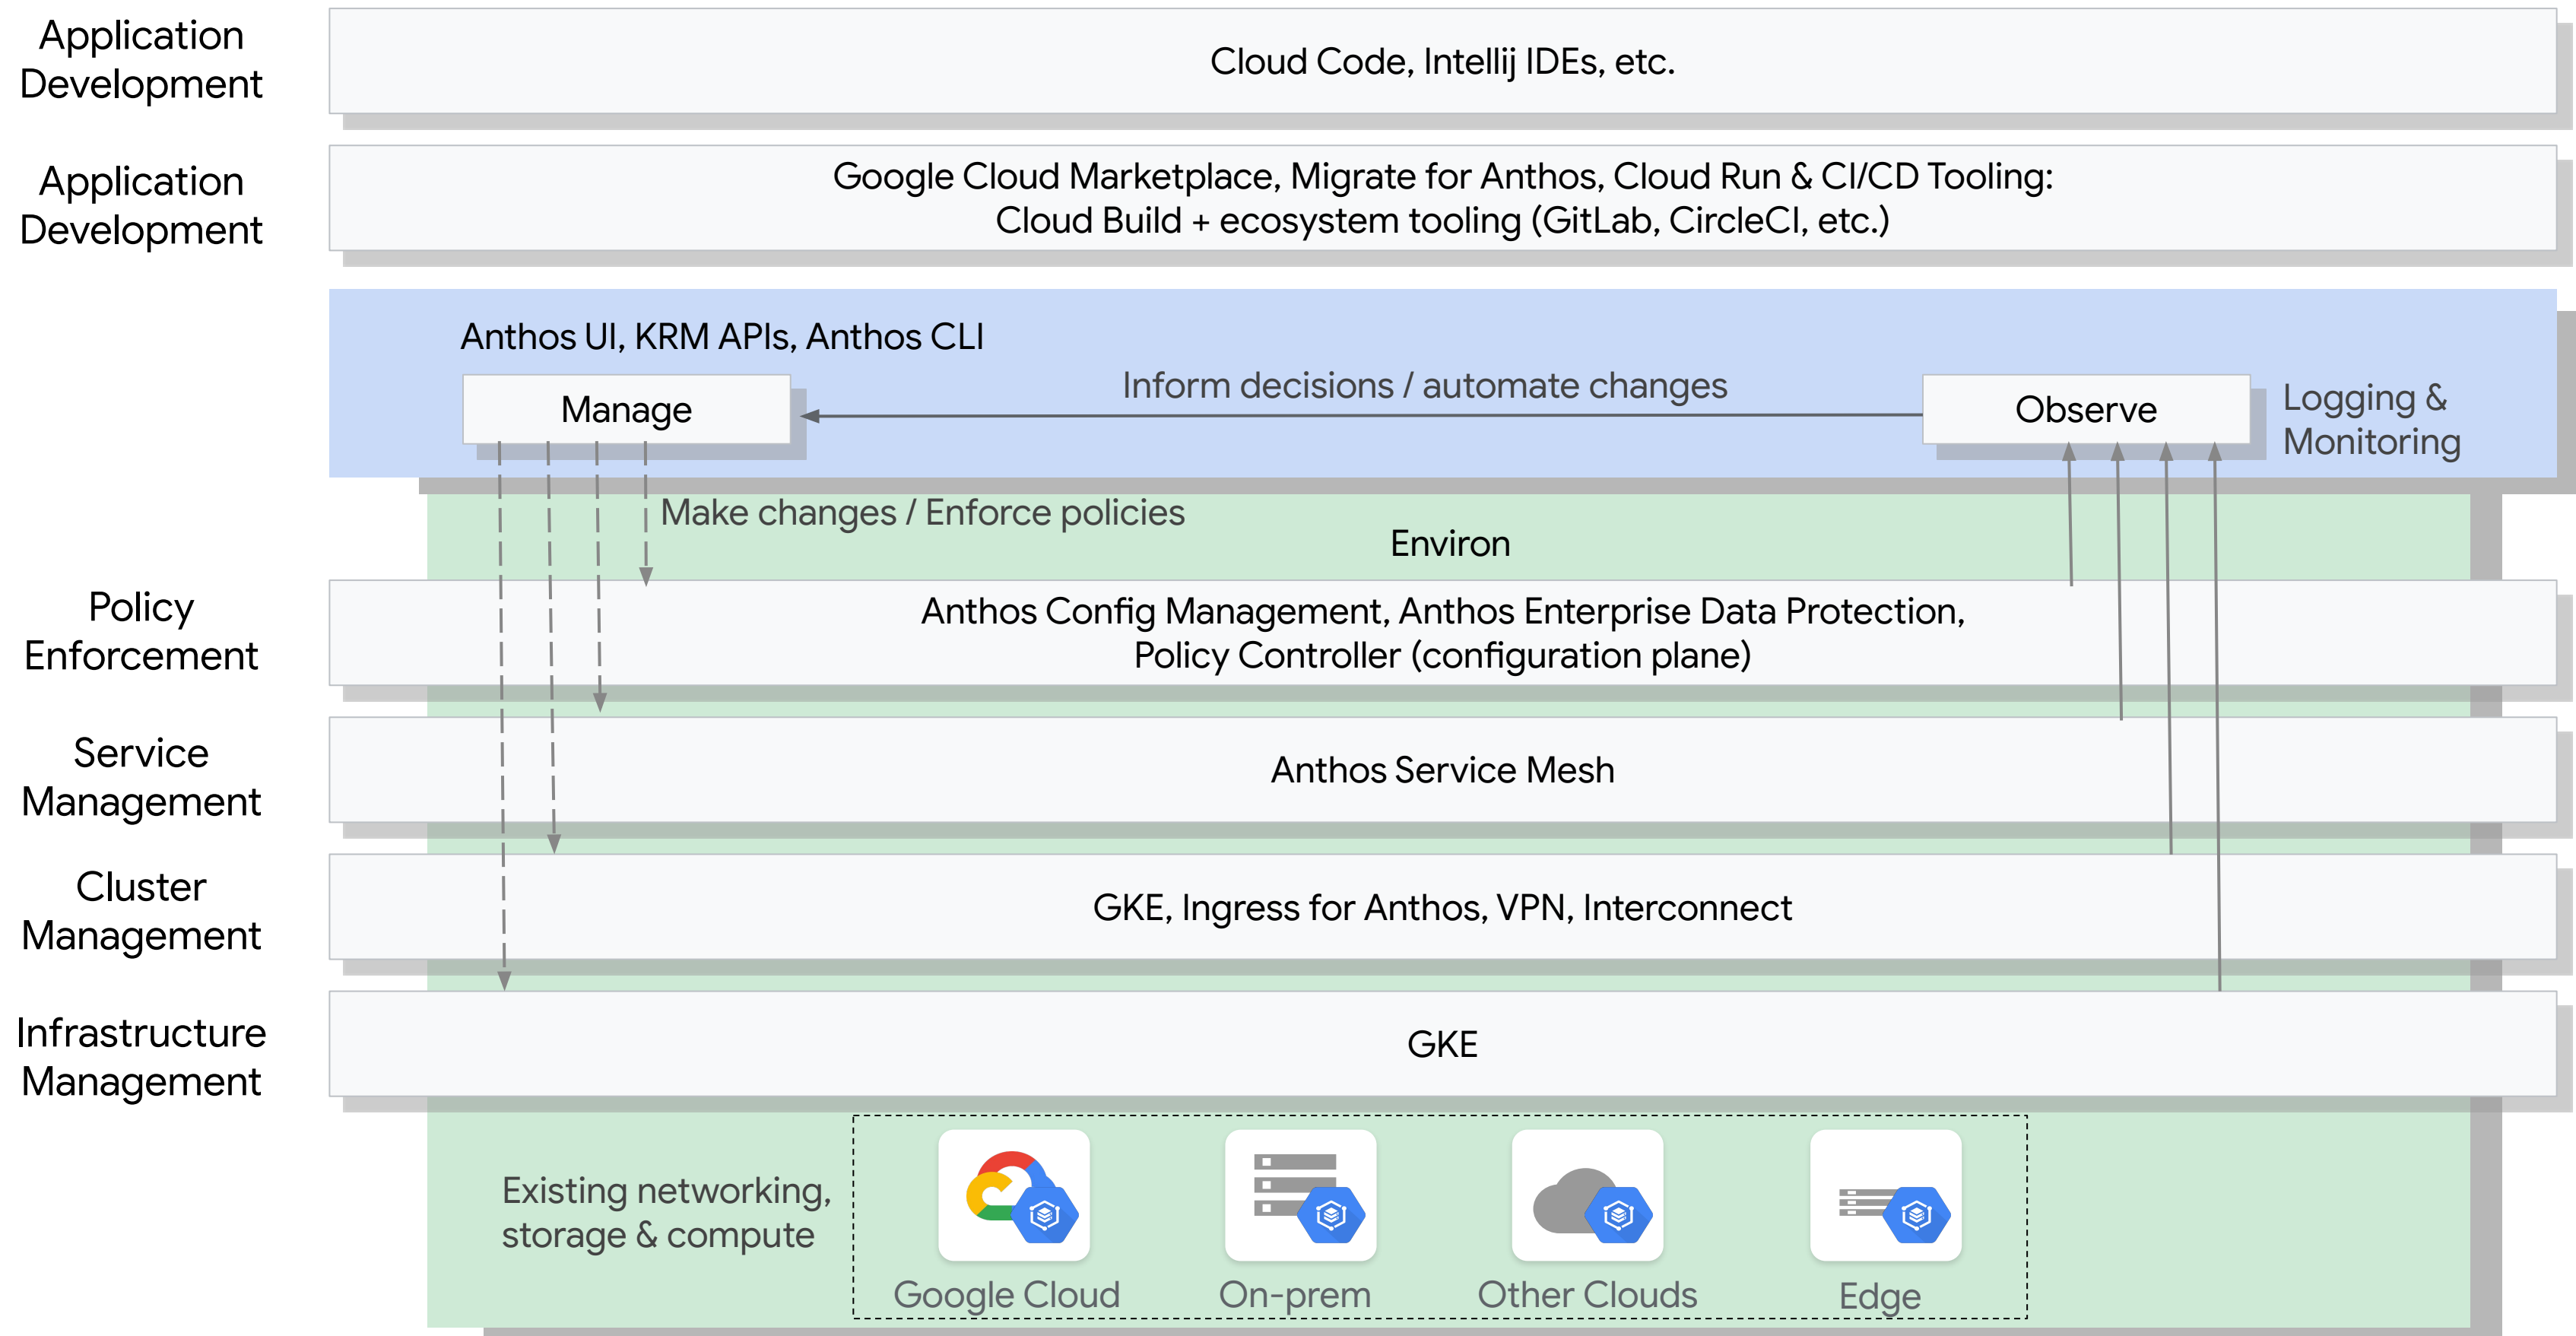
# Anthos technical overview



# Anthos technical overview



# Anthos technical overview



# Anthos Benefits

- A consistent platform for all your application deployments, both legacy as well as cloud native, while offering a service-centric view of all your environments.
- Build enterprise-grade containerized applications faster with managed Kubernetes on cloud and on-premises environments. Create a fast, scalable software delivery pipeline with cloud-native tooling and guidance.
- Leverage a programmatic, outcome-focused approach to managing policies for apps across environments, and enable greater awareness and control with a unified view of your services' health and performance.



