



# Configuring Google Cloud Services for Observability



**The following course materials are **copyright protected** materials.**

They may not be reproduced or distributed and may only be used by students attending this Google Cloud Partner Learning Services program.



---

# Agenda

Working with Agents  
Monitoring

Logging

Images and Agent Policies

Non-VM Resources

Exposing Custom Metrics

---

# Agenda

Working with Agents  
Monitoring

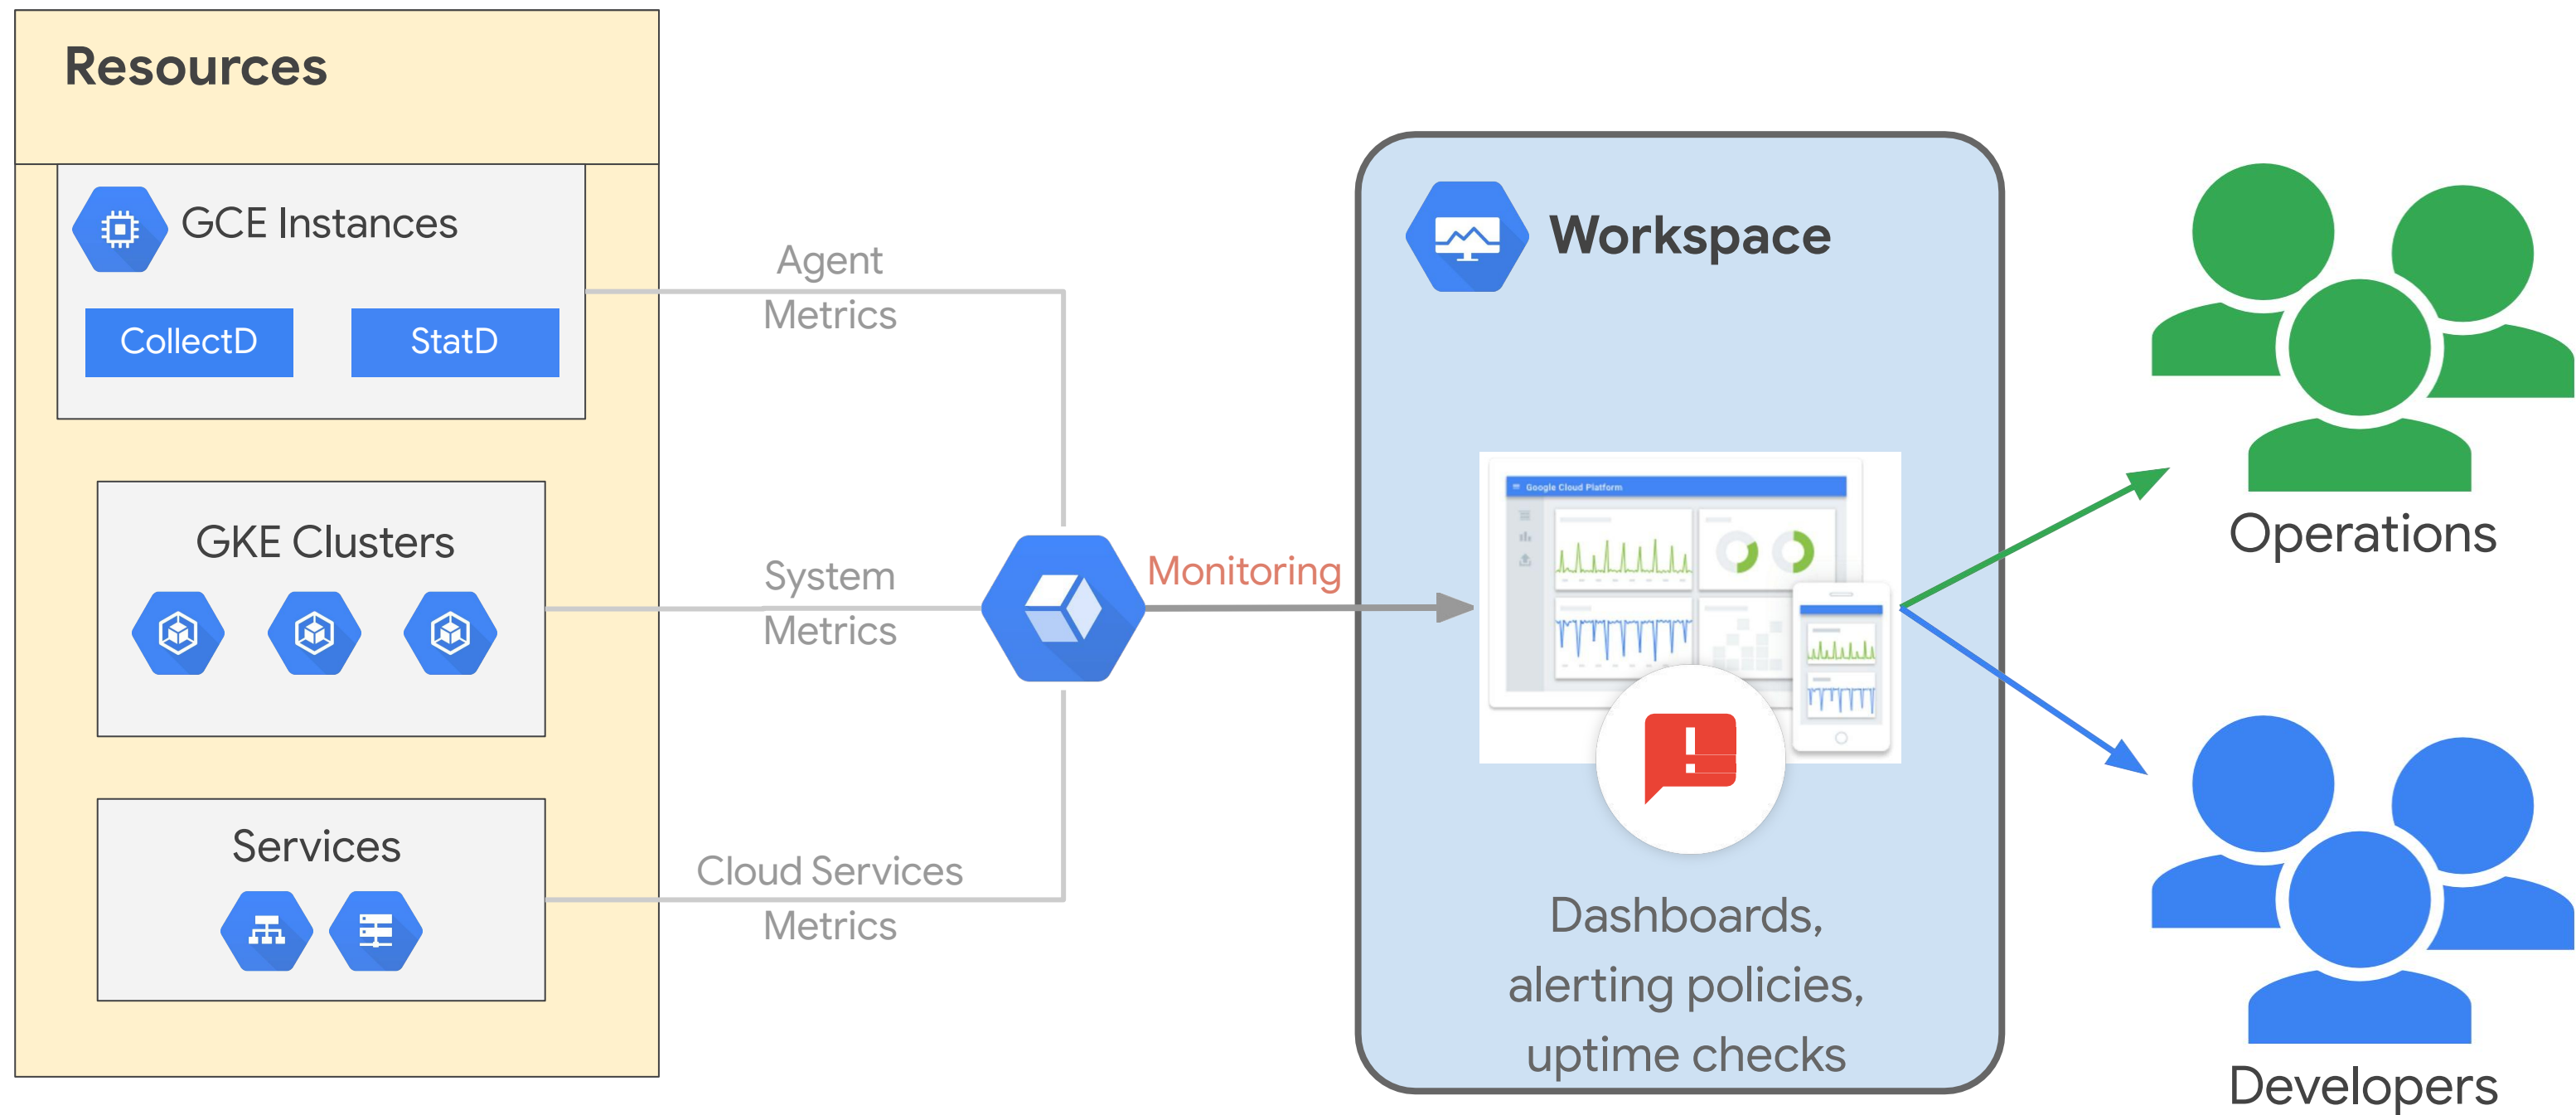
Logging

Images and Agent Policies

Non-VM Resources

Exposing Custom Metrics

# Monitoring Workspace



---

# OS Monitoring

## agent

Gathers system and application metrics from VM instances and sends them to Monitoring

- Based on the open-source collectd
- Gathers additional system resources and application metrics
- Optional, but recommended
- Supports third-party applications, such as:
  - Apache/Nginx/MySQL
- Additional support offered through BindPlane from Blue Medora
- Supports major operating systems:
  - CentOS, Debian, Red Hat Enterprise Linux
  - Ubuntu LTS, SUSE Linux Enterprise Server
  - Windows server



---

# Services with “other” Monitoring support

Don't try to manually install or configure the agent

- App Engine standard has monitoring built-in
- App Engine flex has agent pre-installed and configured
- GKE nodes has monitoring configurable and enabled by default
- Anthos GKE On-Prem agent collects system but not application metrics
- Cloud Run provides integrated monitoring support
- Cloud Function supports integrated monitoring

---

# Installing the Monitoring agent

```
curl -sSO  
https://dl.google.com/cloudagents/add-monitoring-agent-repo.sh sudo  
bash add-monitoring-agent-repo.sh  
sudo zypper install  
stackdriver-agent sudo service  
stackdriver-agent start
```

```
curl -sSO  
https://dl.google.com/cloudagents/add-monitoring-agent-repo.sh sudo  
bash add-monitoring-agent-repo.sh  
sudo apt-get update  
sudo apt-get install  
stackdriver-agent sudo service  
stackdriver-agent start
```



---

# Installing the Monitoring agent

```
CentOS 8
curl -sSO
https://dl.google.com/cloudagents/add-monitoring-agent-repo.sh sudo
bash add-monitoring-agent-repo.sh
sudo yum install -y
stackdriver-agent sudo service
stackdriver-agent start
```

## Other

- All other Linux distros, see [here](#)
- Windows, see [here](#)

---

# Verifying Monitoring agent authorization

Execute on the

```
curl --silent --connect-timeout 1 -f -H "Metadata-Flavor: Google" \
https://169.254.169.254/computeMetadata/v1/instance/service-accounts/default/scopes
```

Check for one or more of the

following

<https://www.googleapis.com/auth/monitoring.admin>

<https://www.googleapis.com/auth/monitoring.write>

<https://www.googleapis.com/auth/cloud-platform>

---

# Agenda

## Working with Agents

Monitoring

Logging

Images and Agent Policies

Non-VM Resources

Exposing Custom Metrics

---

# OS Logging

## agent

Streams logs from common third-party applications and system software to Google Cloud Logging

- Supports third-party applications, such as:
  - Apache/Tomcat/Nginx
  - Chef/Jenkins/Puppet
  - Cassandra/Mongodb/MySQL
- Based on fluentd log data collector—can add own fluentd configuration files
- Supports major operating systems:
  - CentOS
  - Debian
  - Red Hat Enterprise Linux
  - Ubuntu LTS
  - SUSE
  - Windows Server



---

# Services with “other” Logging support

Don't try to manually install or configure the agent

- App Engine flex and standard have built-in support for logging
- GKE nodes can enable GKE logging
- Anthos GKE On-Prem agent collects system but not app logs
- Cloud Run has built-in logging support
- Cloud Functions have built-in logging support

---

# Installing the Logging agent

```
curl -sSO  
https://dl.google.com/cloudagents/install-logging-agent.sh sudo  
bash install-logging-agent.sh
```

## Windows (PowerShell terminal)

```
cd $env:UserProfile;  
(New-Object Net.WebClient).DownloadFile(  
    "https://dl.google.com/cloudagents/windows/StackdriverLogging-v1-10.exe",  
    ".\StackdriverLogging-v1-10.exe")  
.\StackdriverLogging-v1-10.exe
```

# Installing the Ops agent

## Linux

```
curl -sSO  
https://dl.google.com/cloudagents/add-google-cloud-ops-agent-repo.sh sudo  
bash add-google-cloud-ops-agent-repo.sh --also-install
```

## Windows (PowerShell terminal)

```
(New-Object  
Net.WebClient).DownloadFile("https://dl.google.com/cloudagents/add-google-cloud-ops-agent-repo.ps1",  
"${env:UserProfile}\add-google-cloud-ops-agent-repo.ps1") Invoke-Expression  
"${env:UserProfile}\add-google-cloud-ops-agent-repo.ps1 -AlsoInstall"
```

The Ops Agent is the primary agent for collecting telemetry from your Compute Engine instances. Combining logging and metrics into a single agent, the Ops Agent uses Fluent Bit for logs, which supports high-throughput logging, and the OpenTelemetry Collector for metrics.

---

# Verifying Logging agent authorization

Execute on the

```
curl --silent --connect-timeout 1 -f -H "Metadata-Flavor: Google" \
https://169.254.169.254/computeMetadata/v1/instance/service-accounts/default/scopes
```

Check for one or more of the

following

```
https://www.googleapis.com/auth/logging.write  
https://www.googleapis.com/auth/logging.admin  
https://www.googleapis.com/auth/cloud-platform
```



---

# Agenda

## Working with Agents

Monitoring

Logging

## Images and Agent Policies

Non-VM Resources

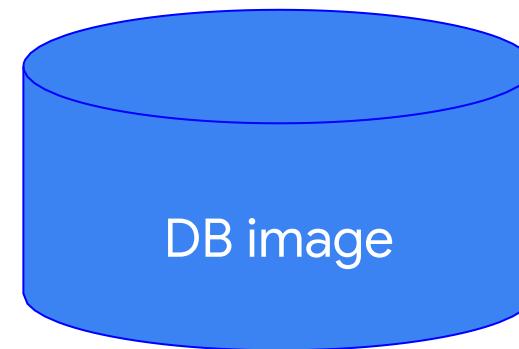
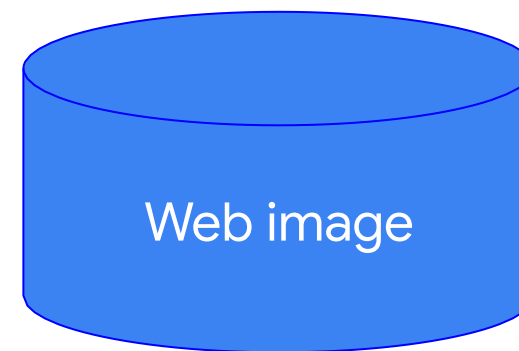
Exposing Custom Metrics

# Organizational maturity

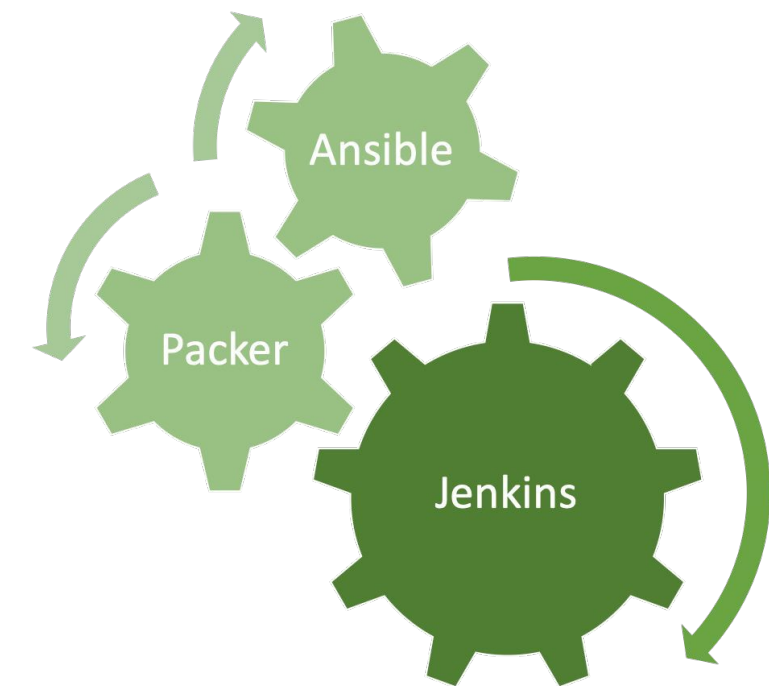
What's an image?



Core image library



Robust image factory

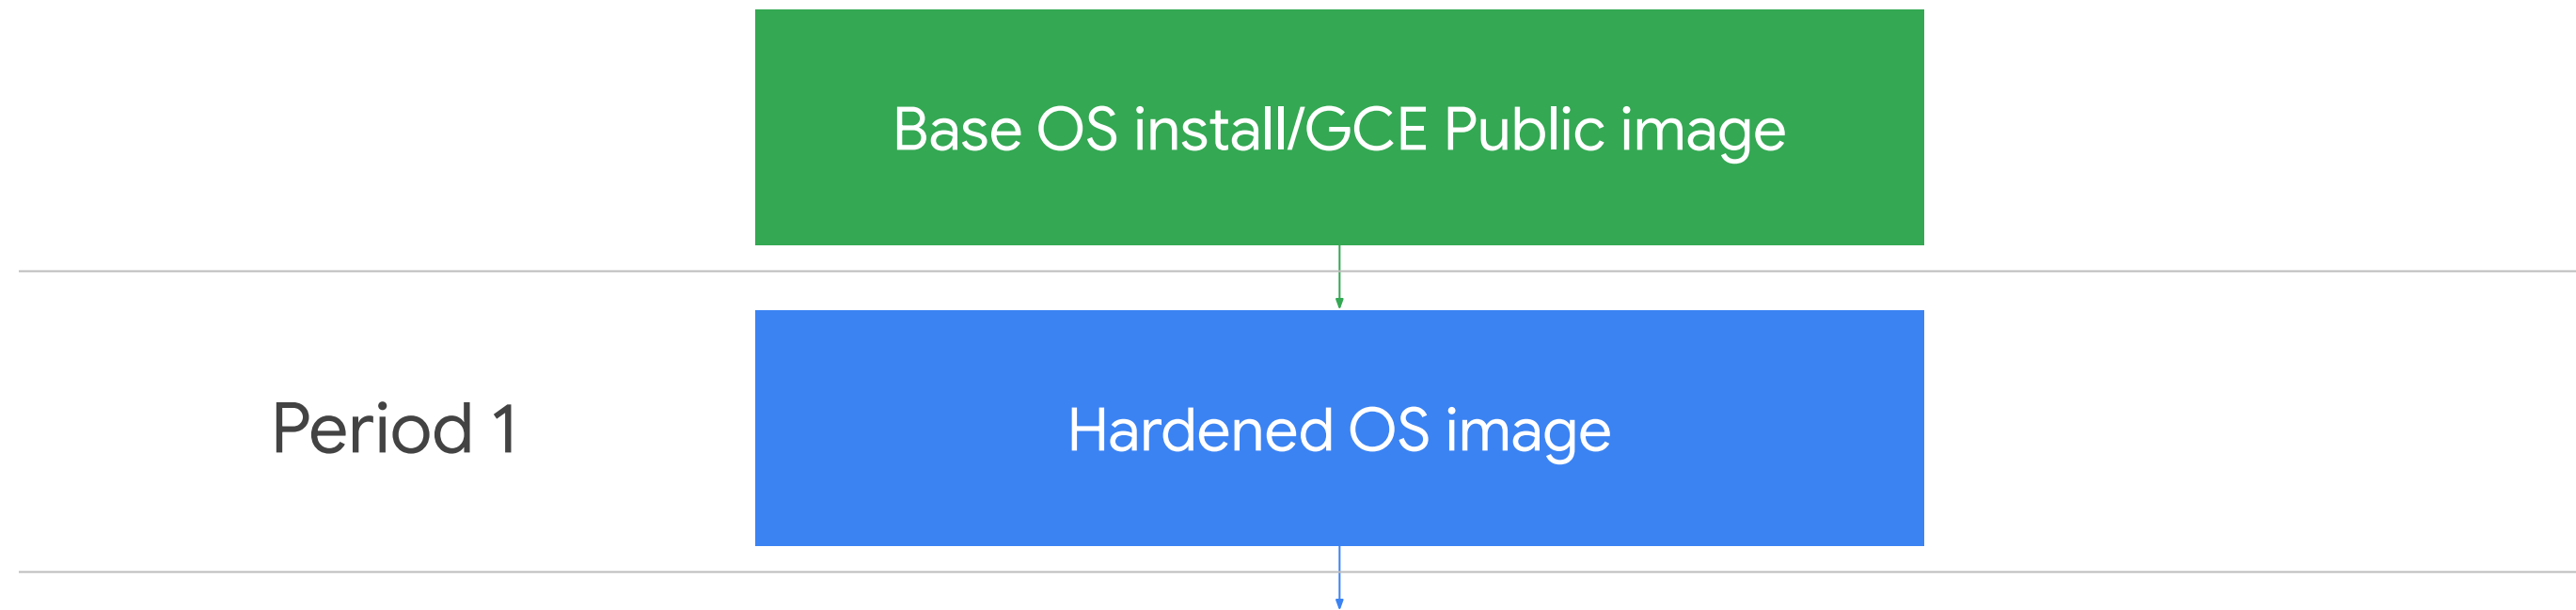


# Basic image management scheme

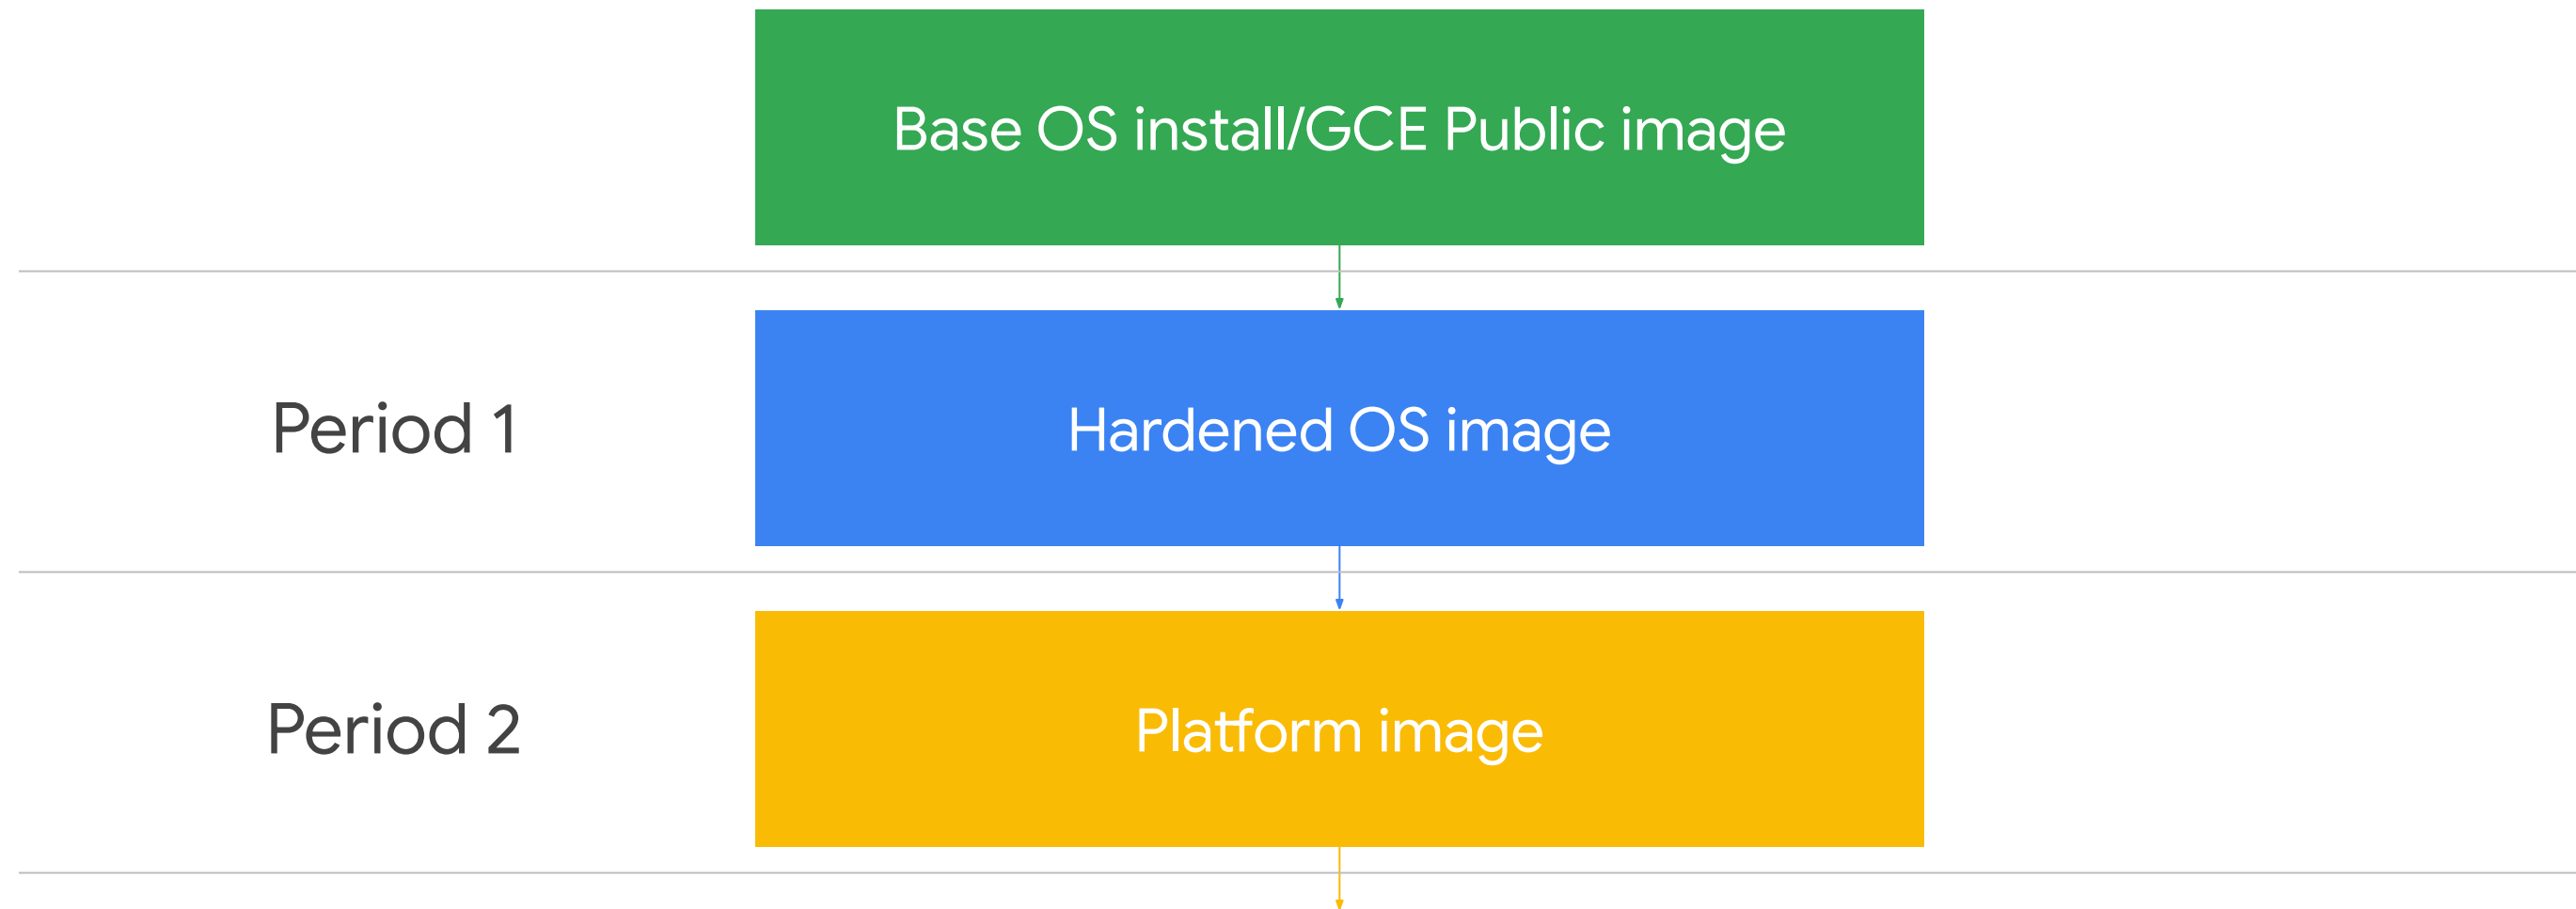
Base OS install/GCE Public image

The diagram consists of a green rectangular box containing the text 'Base OS install/GCE Public image'. A thin green arrow points downwards from the bottom center of this box to a horizontal grey line that spans the width of the diagram area.

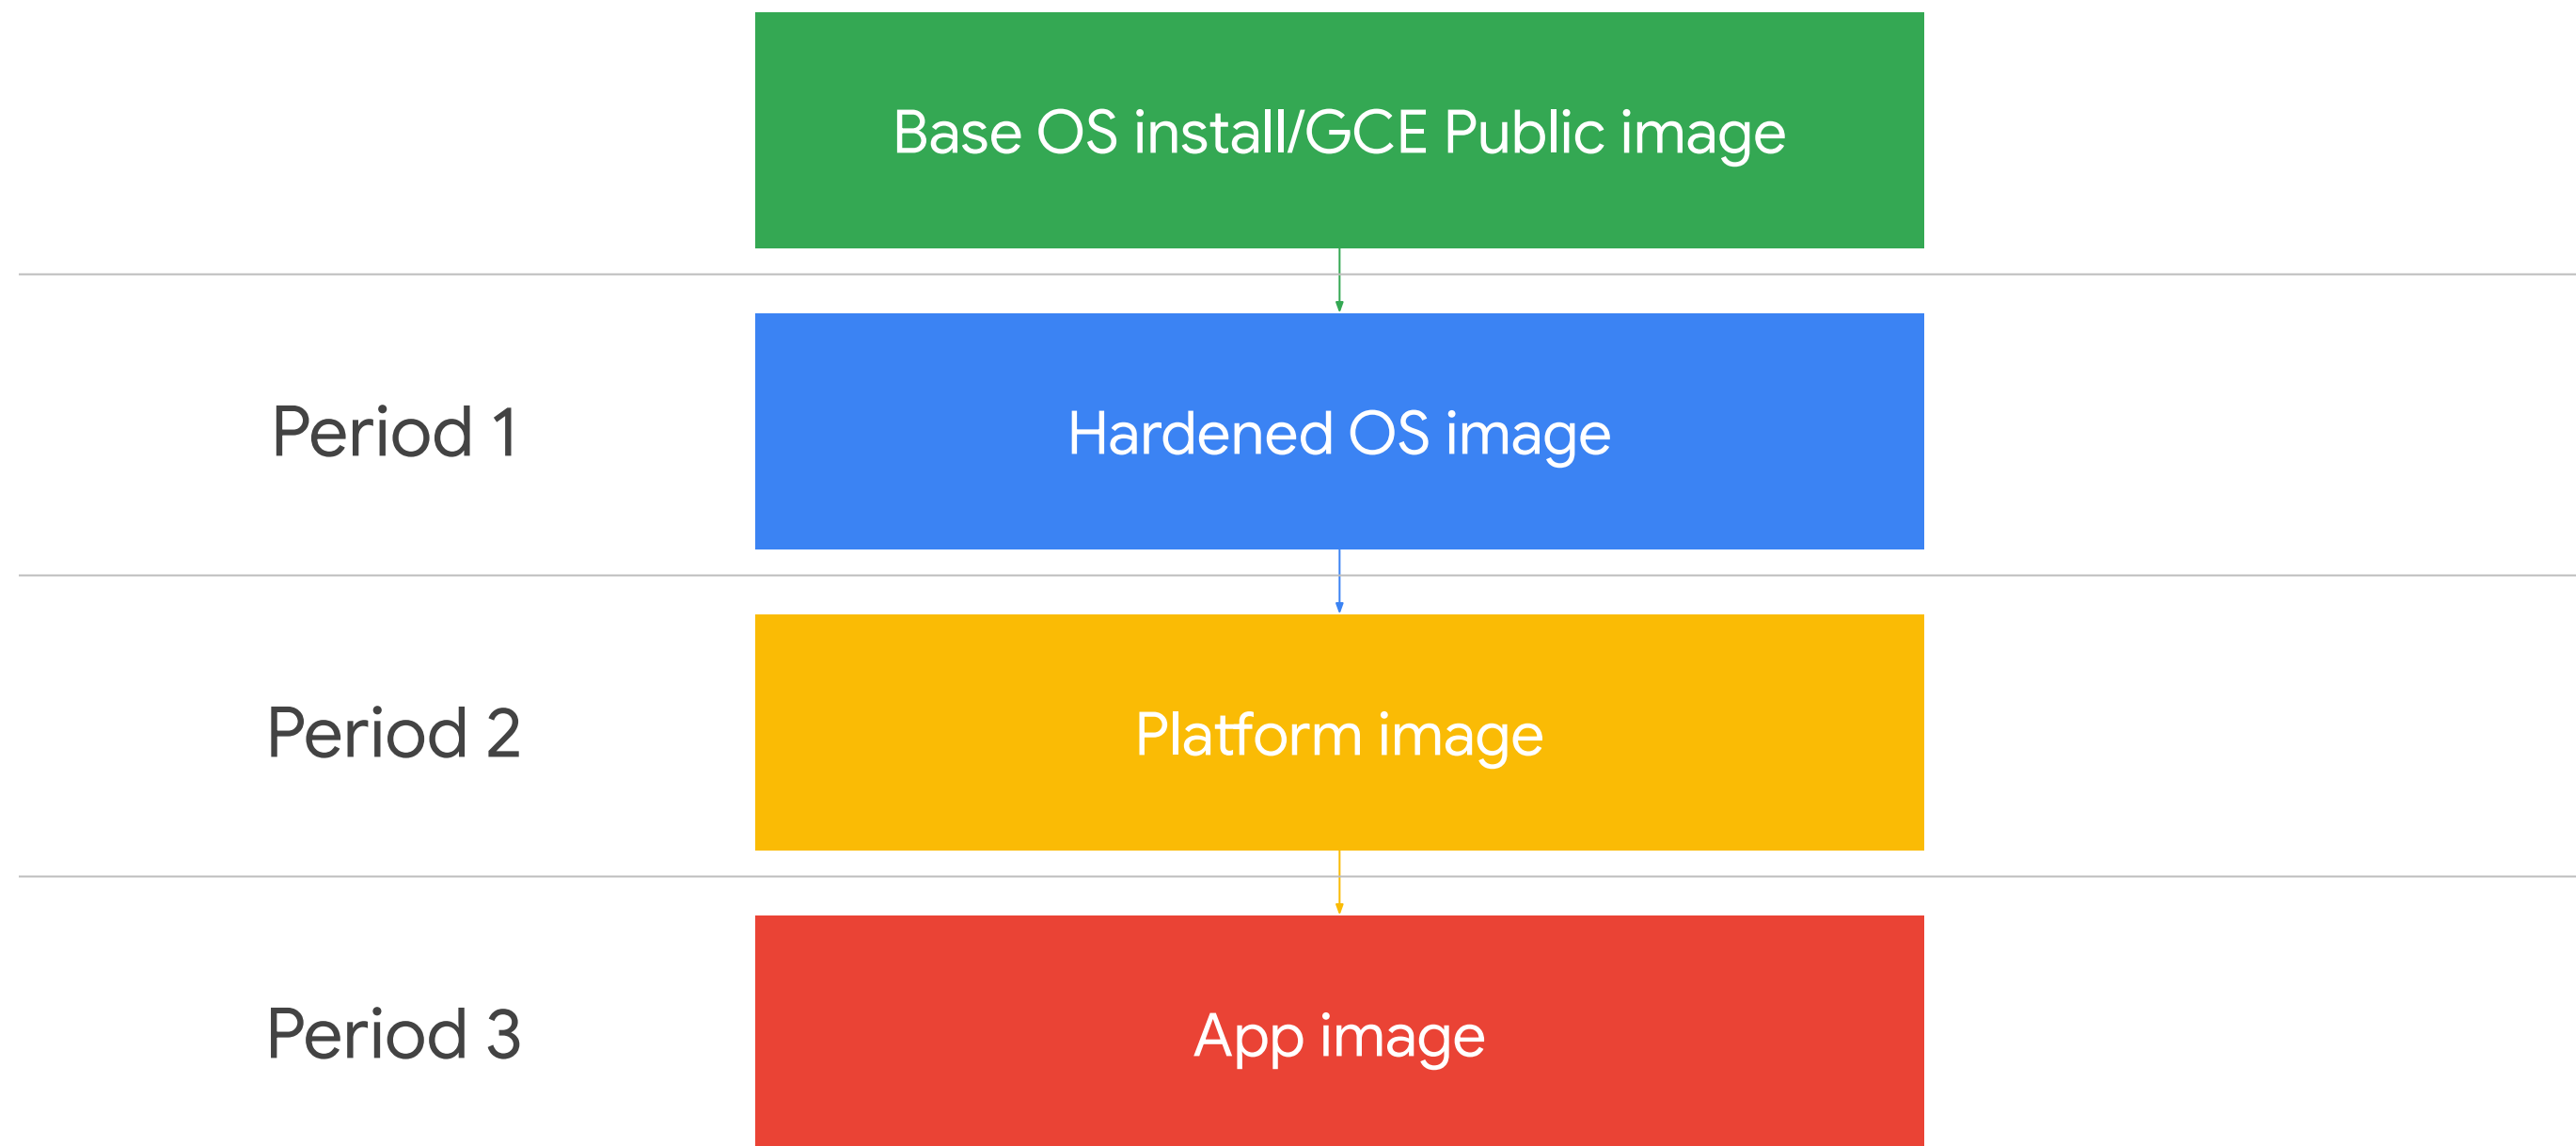
# Basic image management scheme



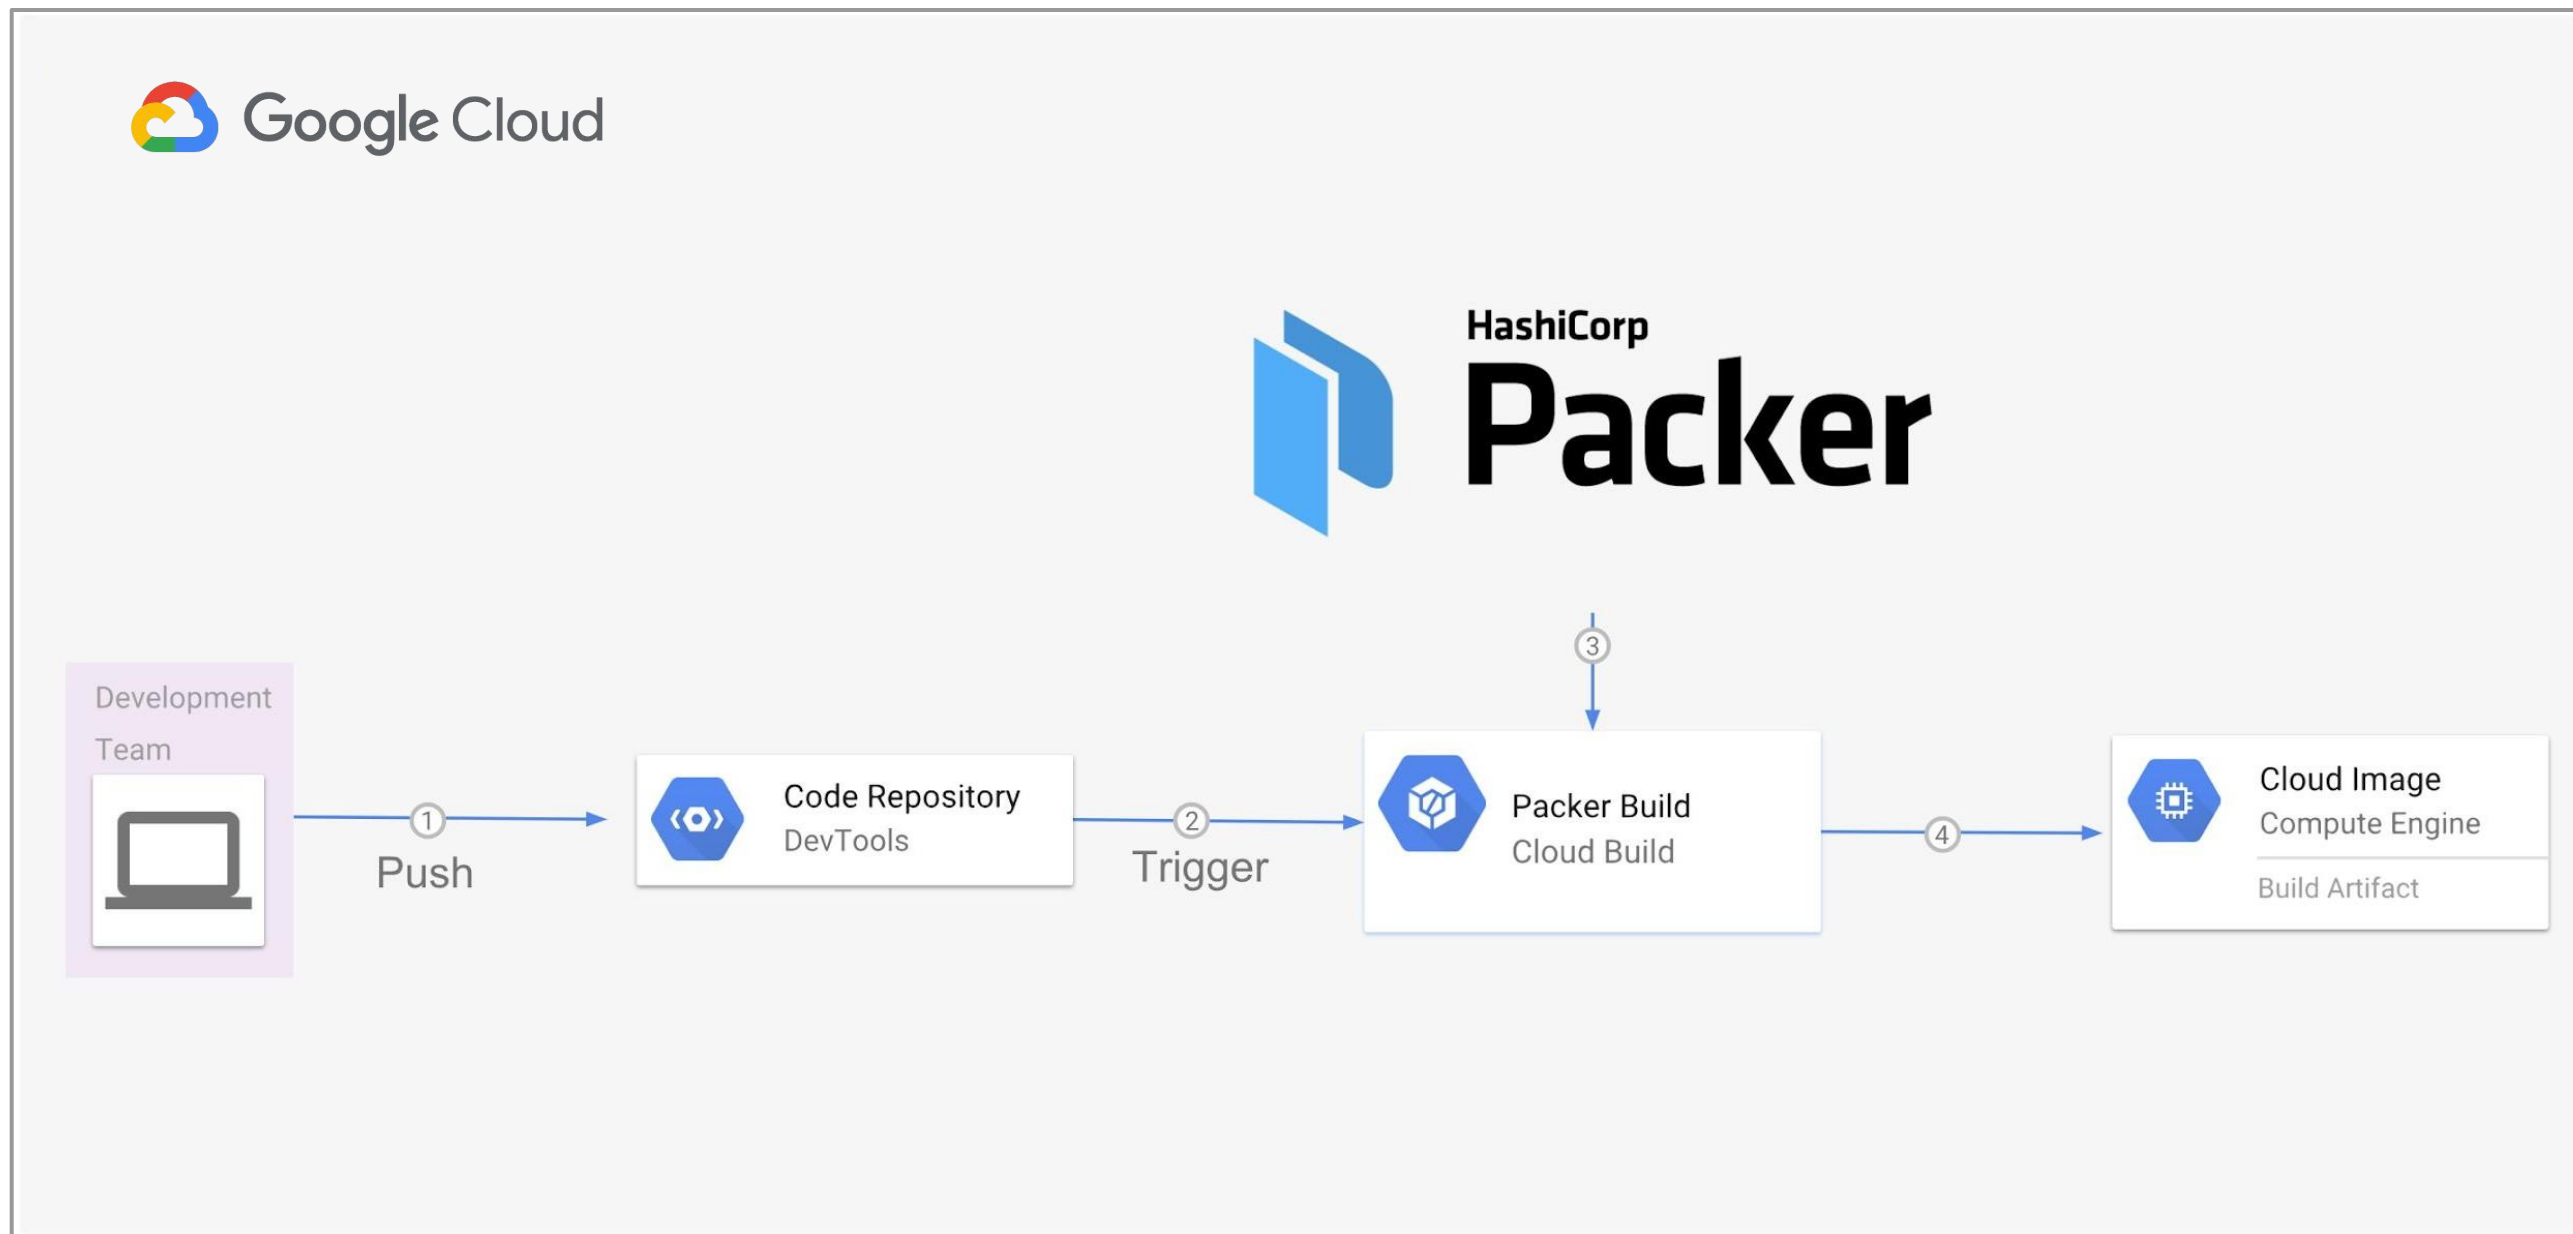
# Basic image management scheme



# Basic image management scheme



# Packer can automate image builds



---

# Leverage Agent Policies to aid with agent automation

- Automate installation and maintenance of Monitoring and Logging agents
- May apply to a fleet of VMs matching user-specified criteria
  - Current support for Linux VMs which support the agents

Here's an example policy that targets all CentOS 7 VMs with the labels *env=test* and *app=myproduct*

```
gcloud alpha compute instances ops-agents policies create  
ops-agents-policy-safe-rollout \  
--agent-rules=  
"type=logging,version=current-major,package-state=installed,enable-autoupgrade=  
true; \  type=metrics,version=current-major,  
package-state=installed,enable-autoupgrade=true" \  
--os-types=short-name=centos,version=7 --group-labels=env=test,app=myproduct
```



---

# Agenda

Working with Agents  
Monitoring

Logging

Images and Agent Policies

Non-VM Resources

Exposing Custom Metrics



---


# App Engine

- Standard and Flex support Monitoring
  - [Check documentation](#) for metric details
- Standard and Flex support Logging
  - Write to `stdout` or `stderr` from code
  - May also use Logging APIs (like Winston on Node.js)
- Logs viewable under [GAE Application](#) resource



# Google Kubernetes Engine

**Operations**

☒ Enable Cloud Operations for GKE 

System and workload logging and monitoring

System logging and monitoring only (beta)

System and workload logging only (Monitoring disabled)


System monitoring only (Logging disabled)

Legacy Logging and Monitoring

**Other**

☐ Enable...

☐ Enable...

☐ Enable GKE usage metering 



# GKE monitoring with the default dashboard

INFRASTRUCTURE		WORKLOADS		SERVICES	
Name	Type ?	Ready ?	Incidents	CPU Utilization ?	Memory Utilization ?
▼ ● monitor-me	Cluster	11 ✓	0 ✓	3.00 <div><div></div></div> 7.50%	11GiB <div><div></div></div> 21.42%
▼ ● gke-monitor-me-default-pool-9906baeb-682q	Node	4 ✓	0 ✓	1.00 <div><div></div></div> 9.30%	3.6GiB <div><div></div></div> 19.47%
▶ ● application-controller-manager-0	Pod	✓	0 ✓	0.10 <div><div></div></div> 9.39%	30MiB <div><div></div></div> 29.01%
▶ ● fluentd-gcp-v3.1.1-kvg2j	Pod	✓	0 ✓	0.10 <div><div></div></div> 10.69%	500MiB <div><div></div></div> 36.57%
▶ ● heapster-gke-7b9b95d8cd-zzjcr	Pod		0 ✓	0.06 <div><div></div></div> 2.35%	211MiB <div><div></div></div> 5.21%
▶ ● kube-proxy-gke-monitor-me-default-pool-9906baeb-682q	Pod		0 ✓	0.10 <div><div></div></div> 1.94%	<div><div></div></div> 22MiB
▶ ● metrics-server-v0.3.1-5c6fbf777-vlkrl	Pod	✓	0 ✓	0.05 <div><div></div></div> 2.72%	355MiB <div><div></div></div> 28.36%
▶ ● prometheus-to-sd-h7qv5	Pod	✓	0 ✓	1.0e-3 <div><div></div></div> 21.14%	20MiB <div><div></div></div> 20.47%
▶ ● wordpress-1-mysql-0	Pod		0 ✓	0.20 <div><div></div></div> 19.86%	<div><div></div></div> 110MiB
▶ ● gke-monitor-me-default-pool-9906baeb-kqrp	Node	2 ✓	0 ✓	1.00 <div><div></div></div> 7.46%	3.6GiB <div><div></div></div> 16.88%
▶ ● gke-monitor-me-default-pool-9906baeb-nnxr	Node	5 ✓	0 ✓	1.00 <div><div></div></div> 5.53%	3.6GiB <div><div></div></div> 26.41%





# View the cluster from three perspectives

INFRASTRUCTURE   WORKLOADS   SERVICES									
Name	Type ?	Ready ?	Incidents	CPU Utilization ?		Memory Utilization ?			
monitor-me	Cluster	11 ✓	0 ✓	3.00	7.50%	11GiB	21.42%		
gke-monitor-me-default-pool-9906baeb-682q	Node	4 ✓	0 ✓	1.00	9.30%	3.6GiB	19.47%		
application-controller-manager-0	Pod	✓	0 ✓	0.10	9.39%	30MiB	29.01%		
fluentd-gcp-v3.1.1-kvg2j	Pod	✓	0 ✓	0.10	10.69%	500MiB	36.57%		
heapster-gke-7b9b95d8cd-zzjcr	Pod		0 ✓	0.06	2.35%	211MiB	5.21%		
kube-proxy-gke-monitor-me-default-pool-9906baeb-682q	Pod		0 ✓	0.10	1.94%	22MiB			
metrics-server-v0.3.1-5c6fbf777-vlkrl	Pod	✓	0 ✓	0.05	2.72%	355MiB	28.36%		
prometheus-to-sd-h7qv5	Pod	✓	0 ✓	1.0e-3	21.14%	20MiB	20.47%		
wordpress-1-mysql-0	Pod		0 ✓	0.20	19.86%	110MiB			
gke-monitor-me-default-pool-9906baeb-kqrp	Node	2 ✓	0 ✓	1.00	7.46%	3.6GiB	16.88%		
gke-monitor-me-default-pool-9906baeb-nnxr	Node	5 ✓	0 ✓	1.00	5.53%	3.6GiB	26.41%		



# Resources preceded by a status indicator

INFRASTRUCTURE   WORKLOADS   SERVICES									
Name	Type ?	Ready ?	Incidents	CPU Utilization ?		Memory Utilization ?			
monitor-me	Cluster	11 ✓	0 ✓	3.00	7.50%	11GiB	21.42%		
gke-monitor-me-default-pool-9906baeb-682q	Node	4 ✓	0 ✓	1.00	9.30%	3.6GiB	19.47%		
application-controller-manager-0	Pod	✓	0 ✓	0.10	9.39%	30MiB	29.01%		
fluentd-gcp-v3.1.1-kvg2j	Pod	✓	0 ✓	0.10	10.69%	500MiB	36.57%		
heapster-gke-7b9b95d8cd-zzjcr	Pod		0 ✓	0.06	2.35%	211MiB	5.21%		
kube-proxy-gke-monitor-me-default-pool-9906baeb-682q	Pod		0 ✓	0.10	1.94%	22MiB			
metrics-server-v0.3.1-5c6fbf777-vlkrl	Pod	✓	0 ✓	0.05	2.72%	355MiB	28.36%		
prometheus-to-sd-h7qv5	Pod	✓	0 ✓	1.0e-3	21.14%	20MiB	20.47%		
wordpress-1-mysql-0	Pod		0 ✓	0.20	19.86%	110MiB			
gke-monitor-me-default-pool-9906baeb-kqrp	Node	2 ✓	0 ✓	1.00	7.46%	3.6GiB	16.88%		
gke-monitor-me-default-pool-9906baeb-nnxr	Node	5 ✓	0 ✓	1.00	5.53%	3.6GiB	26.41%		



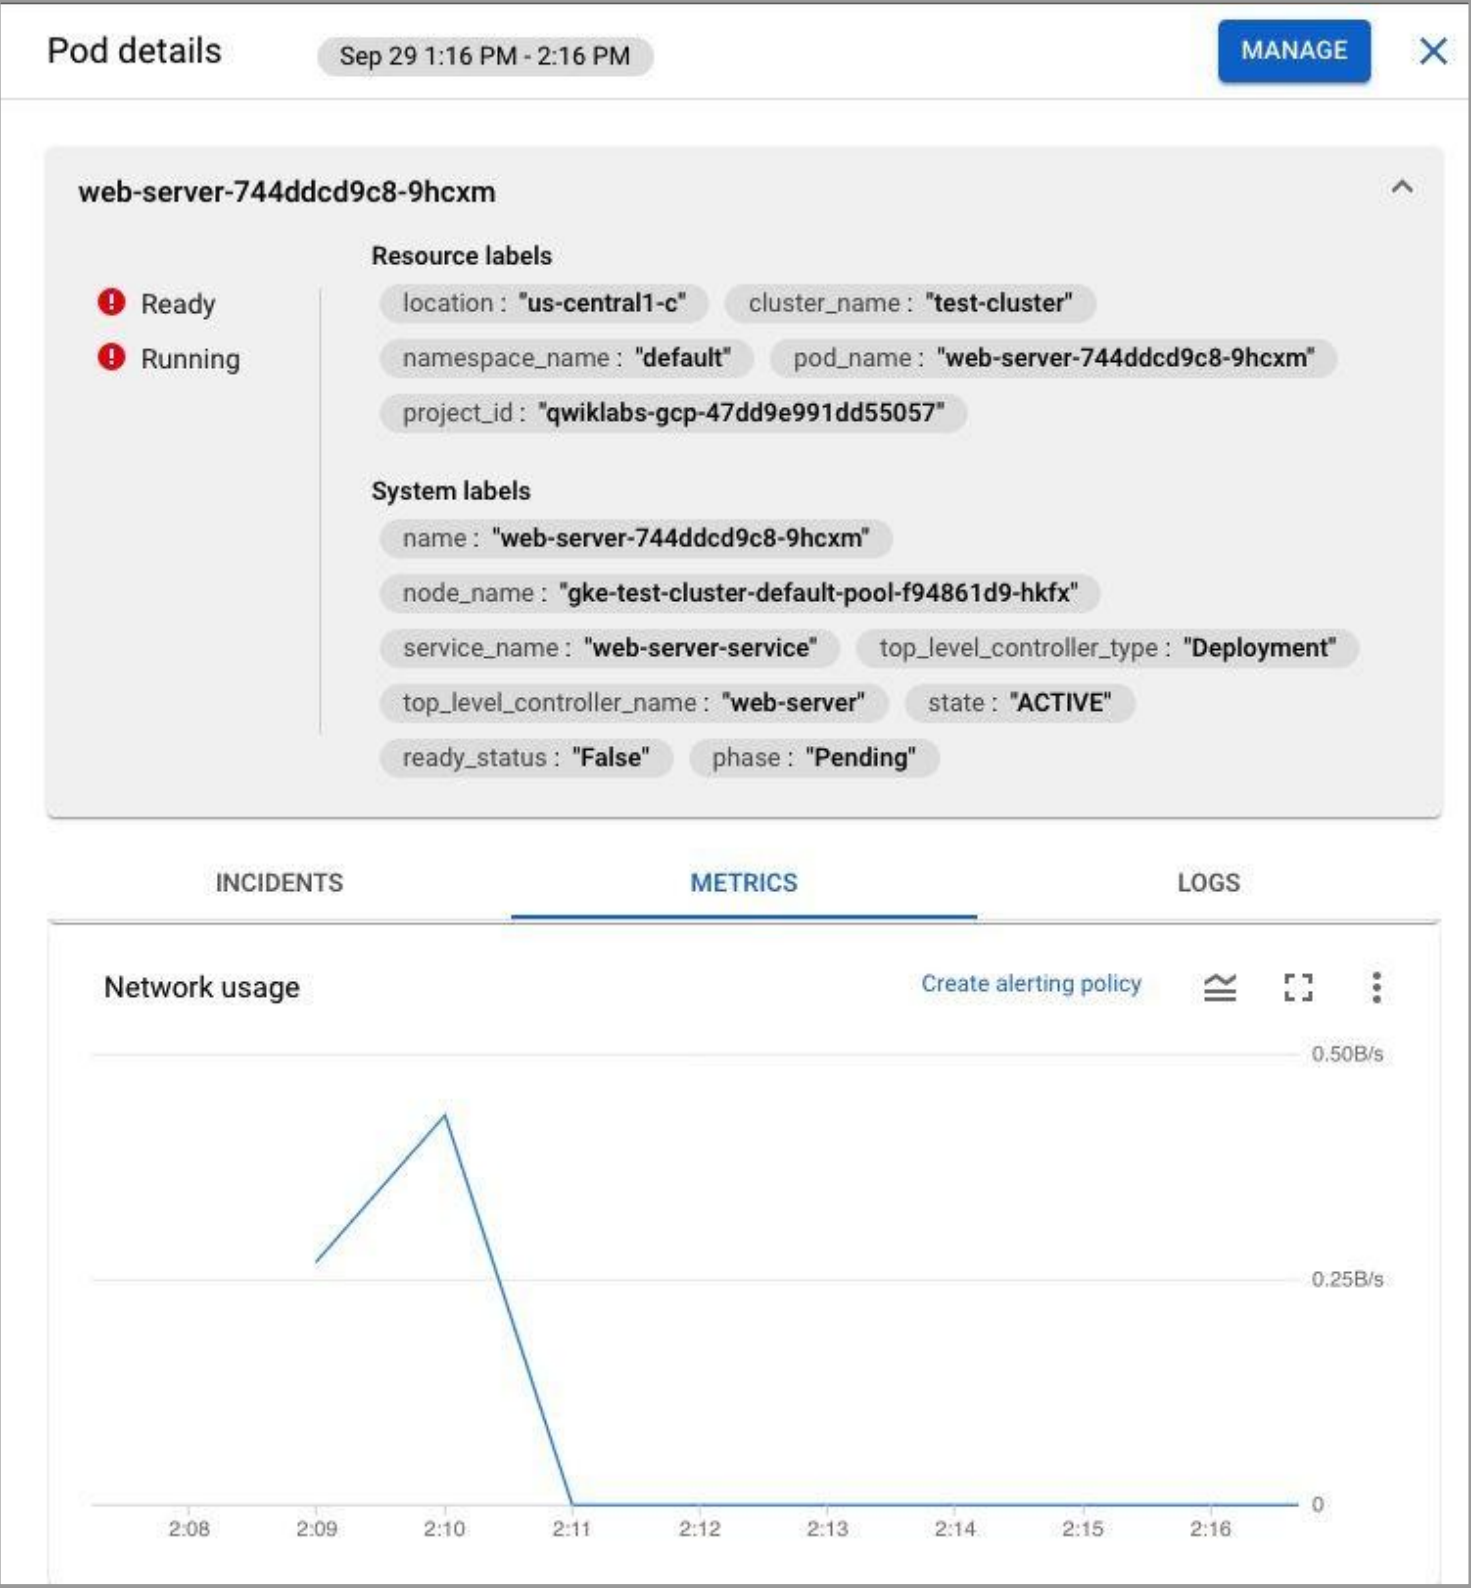
# Status information provided for each GKE object



INFRASTRUCTURE			WORKLOADS			SERVICES		
Name			Type ?	Ready ?	Incidents	CPU Utilization ?		Memory Utilization ?
▼	●	monitor-me	Cluster	11 ✓	0 ✓	3.00	7.50%	11GiB 21.42%
▼	●	gke-monitor-me-default-pool-9906baeb-682q	Node	4 ✓	0 ✓	1.00	9.30%	3.6GiB 19.47%
▶	●	application-controller-manager-0	Pod	✓	0 ✓	0.10	9.39%	30MiB 29.01%
▶	●	fluentd-gcp-v3.1.1-kvg2j	Pod	✓	0 ✓	0.10	10.69%	500MiB 36.57%
▶	●	heapster-gke-7b9b95d8cd-zzjcr	Pod		0 ✓	0.06	2.35%	211MiB 5.21%
▶	●	kube-proxy-gke-monitor-me-default-pool-9906baeb-682q	Pod		0 ✓	0.10	1.94%	22MiB
▶	●	metrics-server-v0.3.1-5c6fbf777-vlkrl	Pod	✓	0 ✓	0.05	2.72%	355MiB 28.36%
▶	●	prometheus-to-sd-h7qv5	Pod	✓	0 ✓	1.0e-3	21.14%	20MiB 20.47%
▶	●	wordpress-1-mysql-0	Pod		0 ✓	0.20	19.86%	110MiB
▶	●	gke-monitor-me-default-pool-9906baeb-kqrp	Node	2 ✓	0 ✓	1.00	7.46%	3.6GiB 16.88%
▶	●	gke-monitor-me-default-pool-9906baeb-nnxr	Node	5 ✓	0 ✓	1.00	5.53%	3.6GiB 26.41%



# Select a pod to view details including metrics







# Logs tab displays the latest entries

Pod details

Sep 29 1:16 PM - 2:16 PM

MANAGE

×

web-server-744ddcd9c8-9hcxm

! Ready

! Running

Resource labels

location: "us-central1-c" cluster\_name: "test-cluster"

namespace\_name: "default" pod\_name: "web-server-744ddcd9c8-9hcxm"

project\_id: "qwiklabs-gcp-47dd9e991dd55057"

System labels

name: "web-server-744ddcd9c8-9hcxm"

node\_name: "gke-test-cluster-default-pool-f94861d9-hkfx"

service\_name: "web-server-service" top\_level\_controller\_type: "Deployment"

top\_level\_controller\_name: "web-server" state: "ACTIVE"

ready\_status: "False" phase: "Pending"

INCIDENTS

METRICS

LOGS

Logs

Default

Filter

↺

↻

▶

2020-09-29 14:08:51.000 CDT

Successfully assigned default/web-server-744ddcd9c8-9hcxm to gke-t...

▶

2020-09-29 14:08:52.000 CDT

Pulling image "nginx:latest"

▶

2020-09-29 14:08:55.000 CDT

Successfully pulled image "nginx:latest"

▶

2020-09-29 14:08:56.000 CDT

Created container nginx-1

▶

2020-09-29 14:08:56.000 CDT

Started container nginx-1

▶

2020-09-29 14:08:56.328 CDT

/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will at...



---

# What is Prometheus?

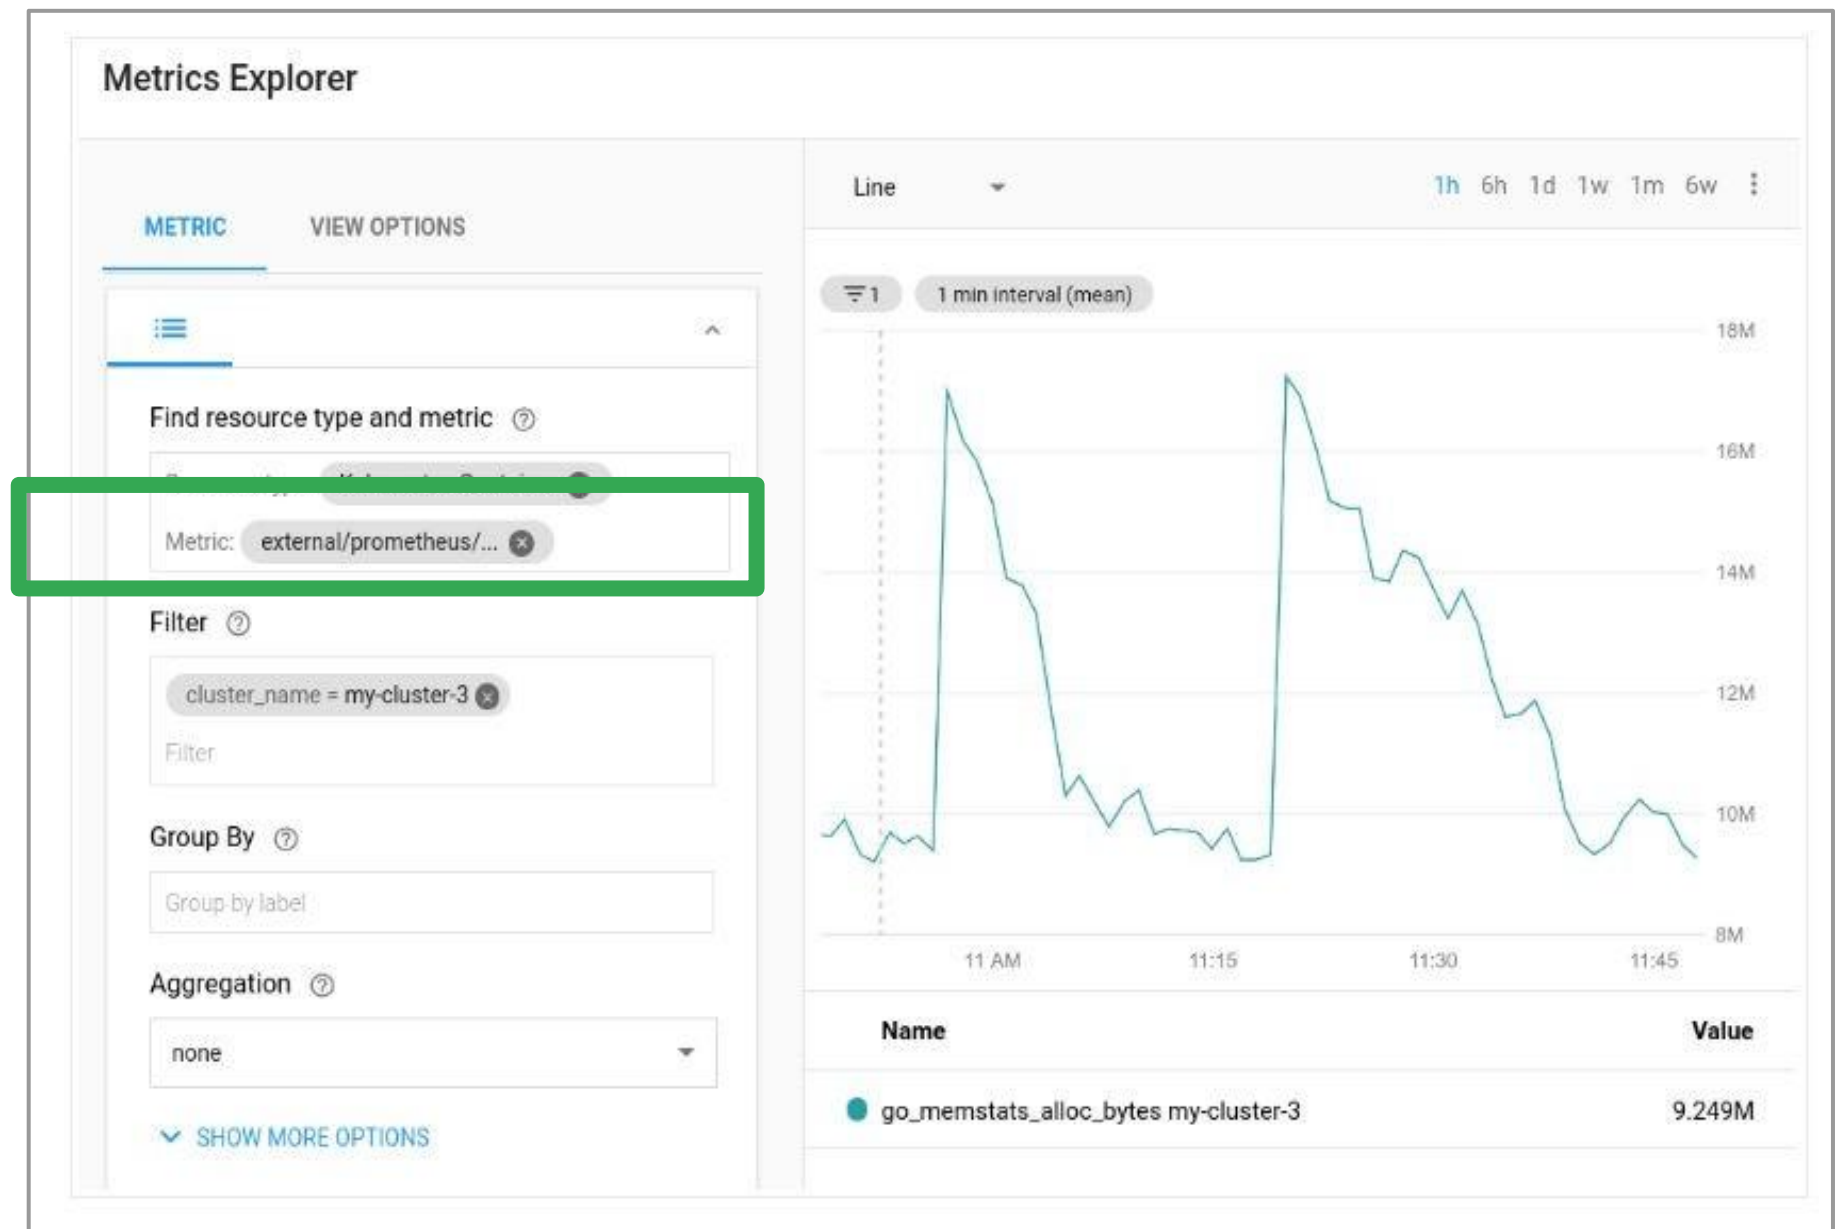
- Prometheus is an optional monitoring tool for Kubernetes
  - Supported with GKE Monitoring
- Service metrics using Prometheus exposition format can be exported and made visible as external metrics



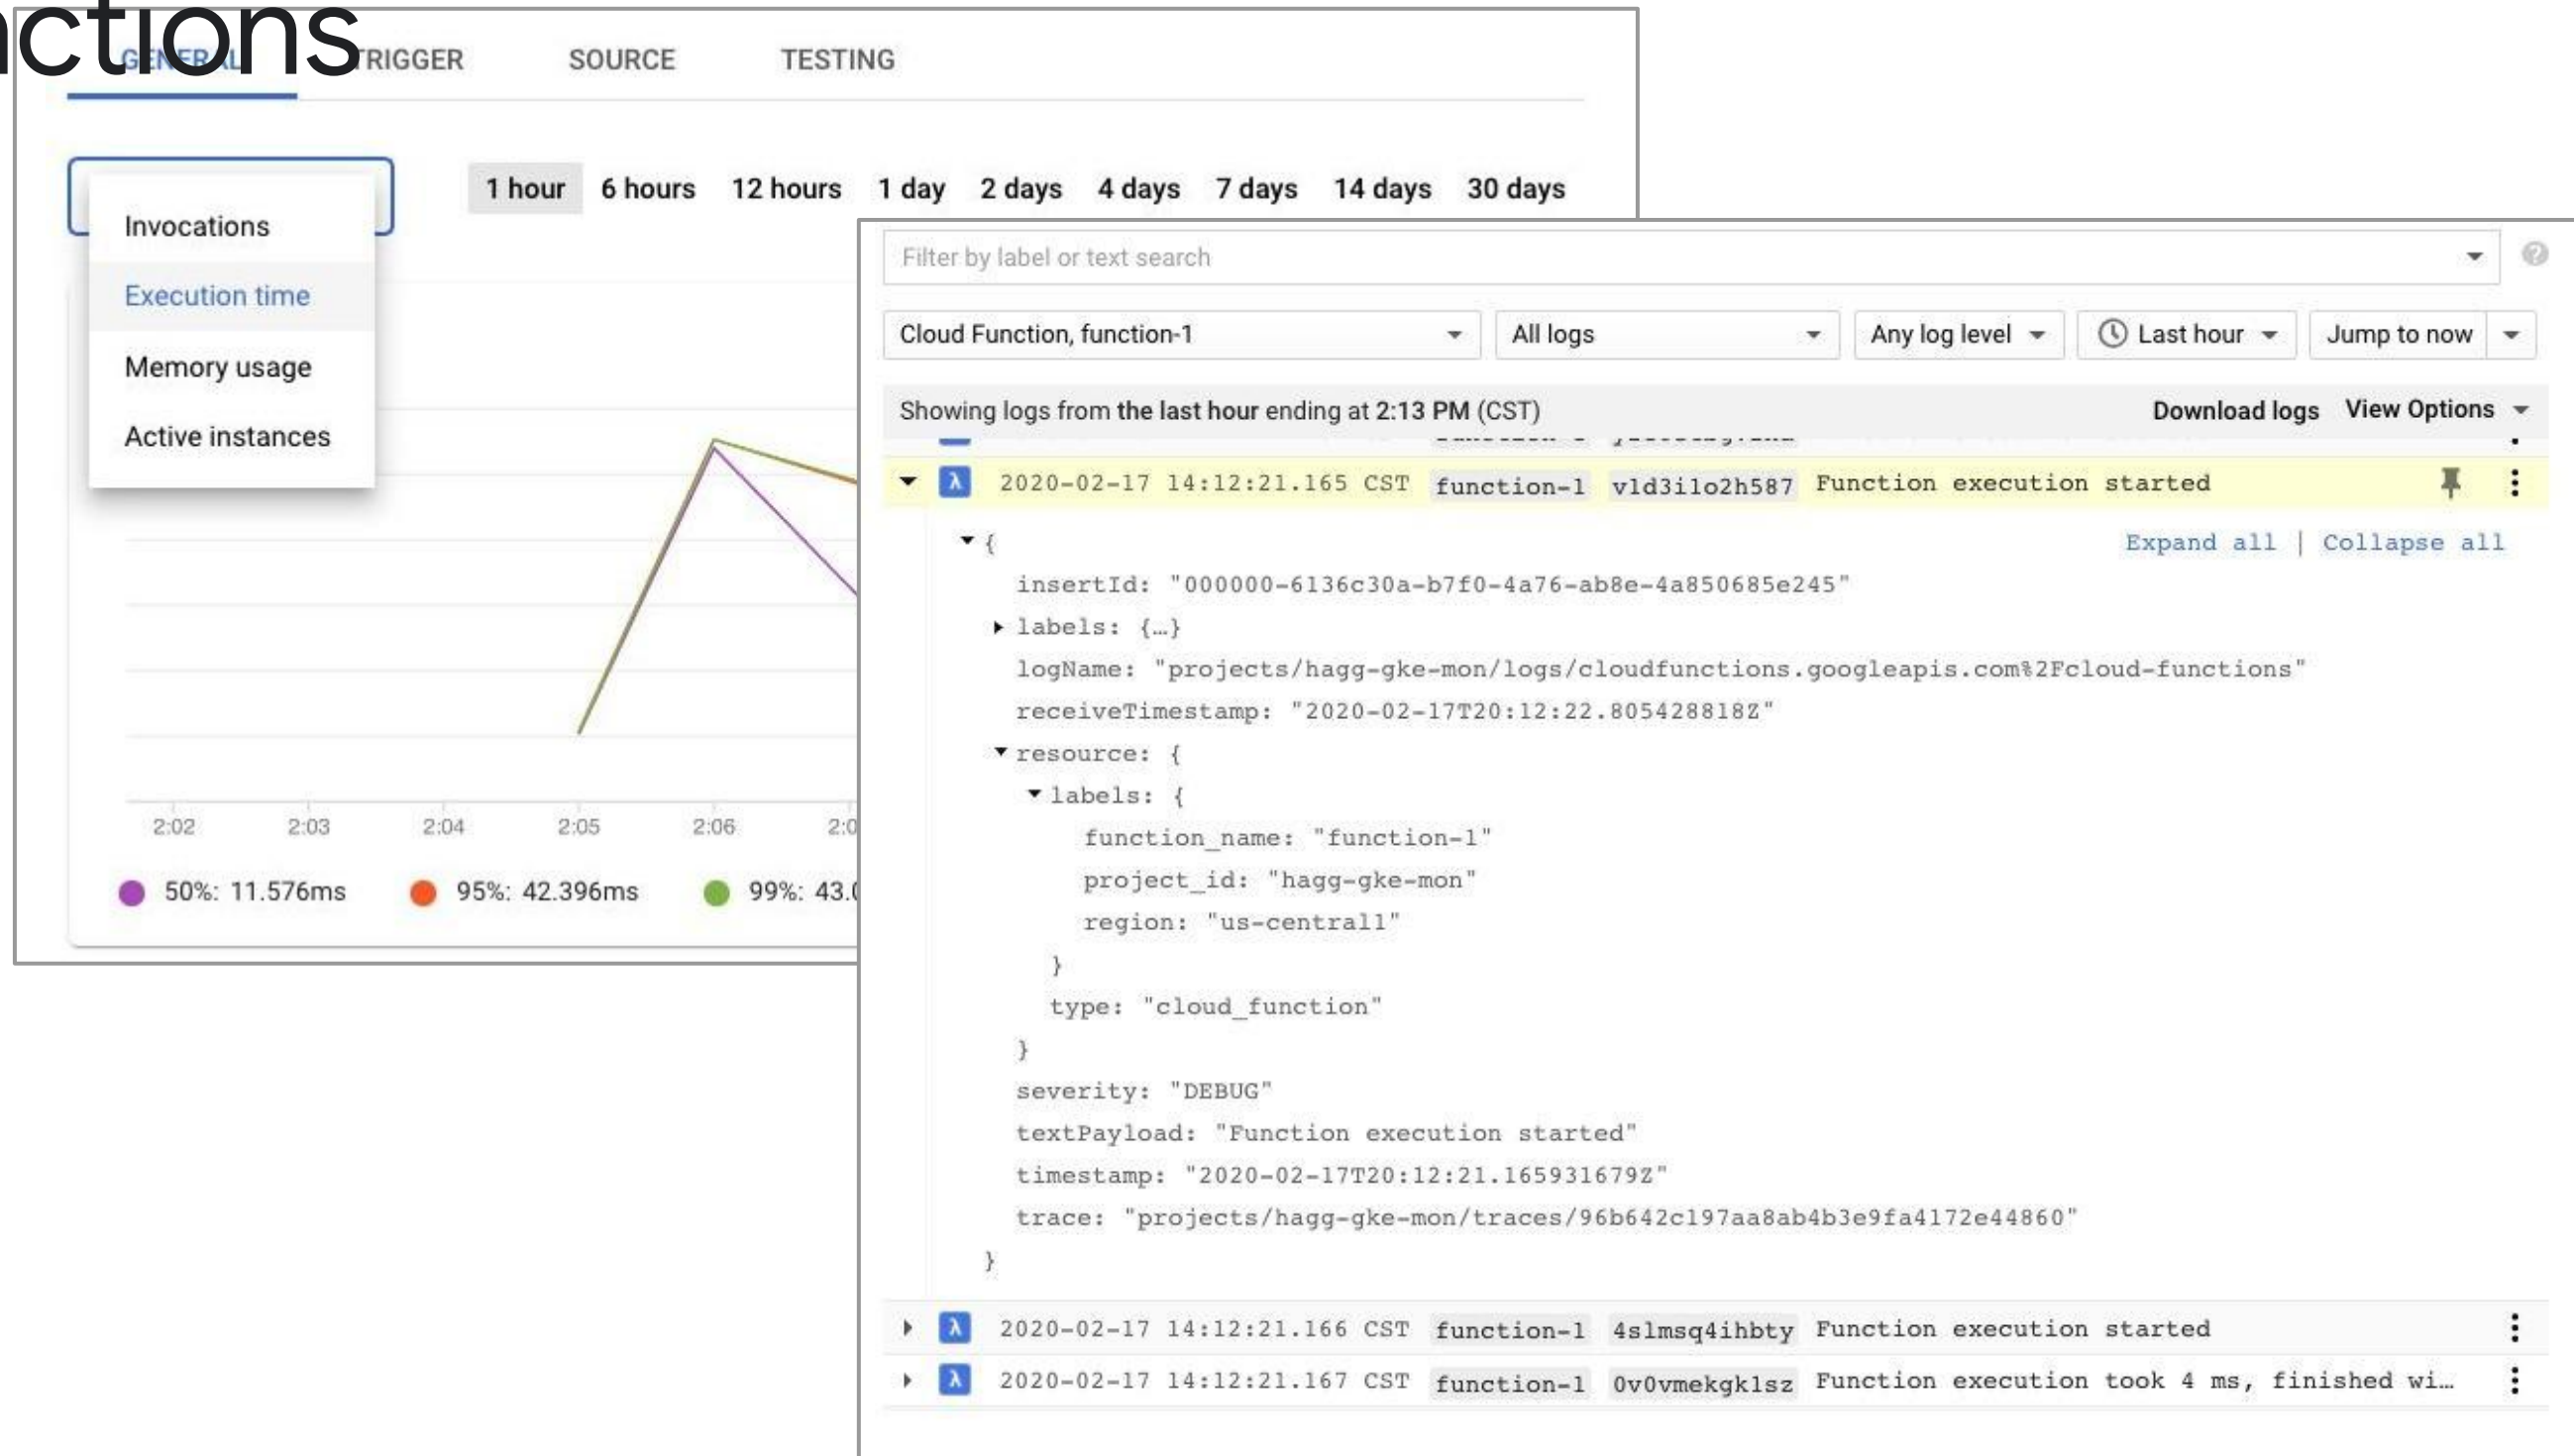


# Configure Prometheus for GKE

- Install Prometheus and the Collector
- Metrics can be viewed as external metrics
  - *external/prometheus/\**

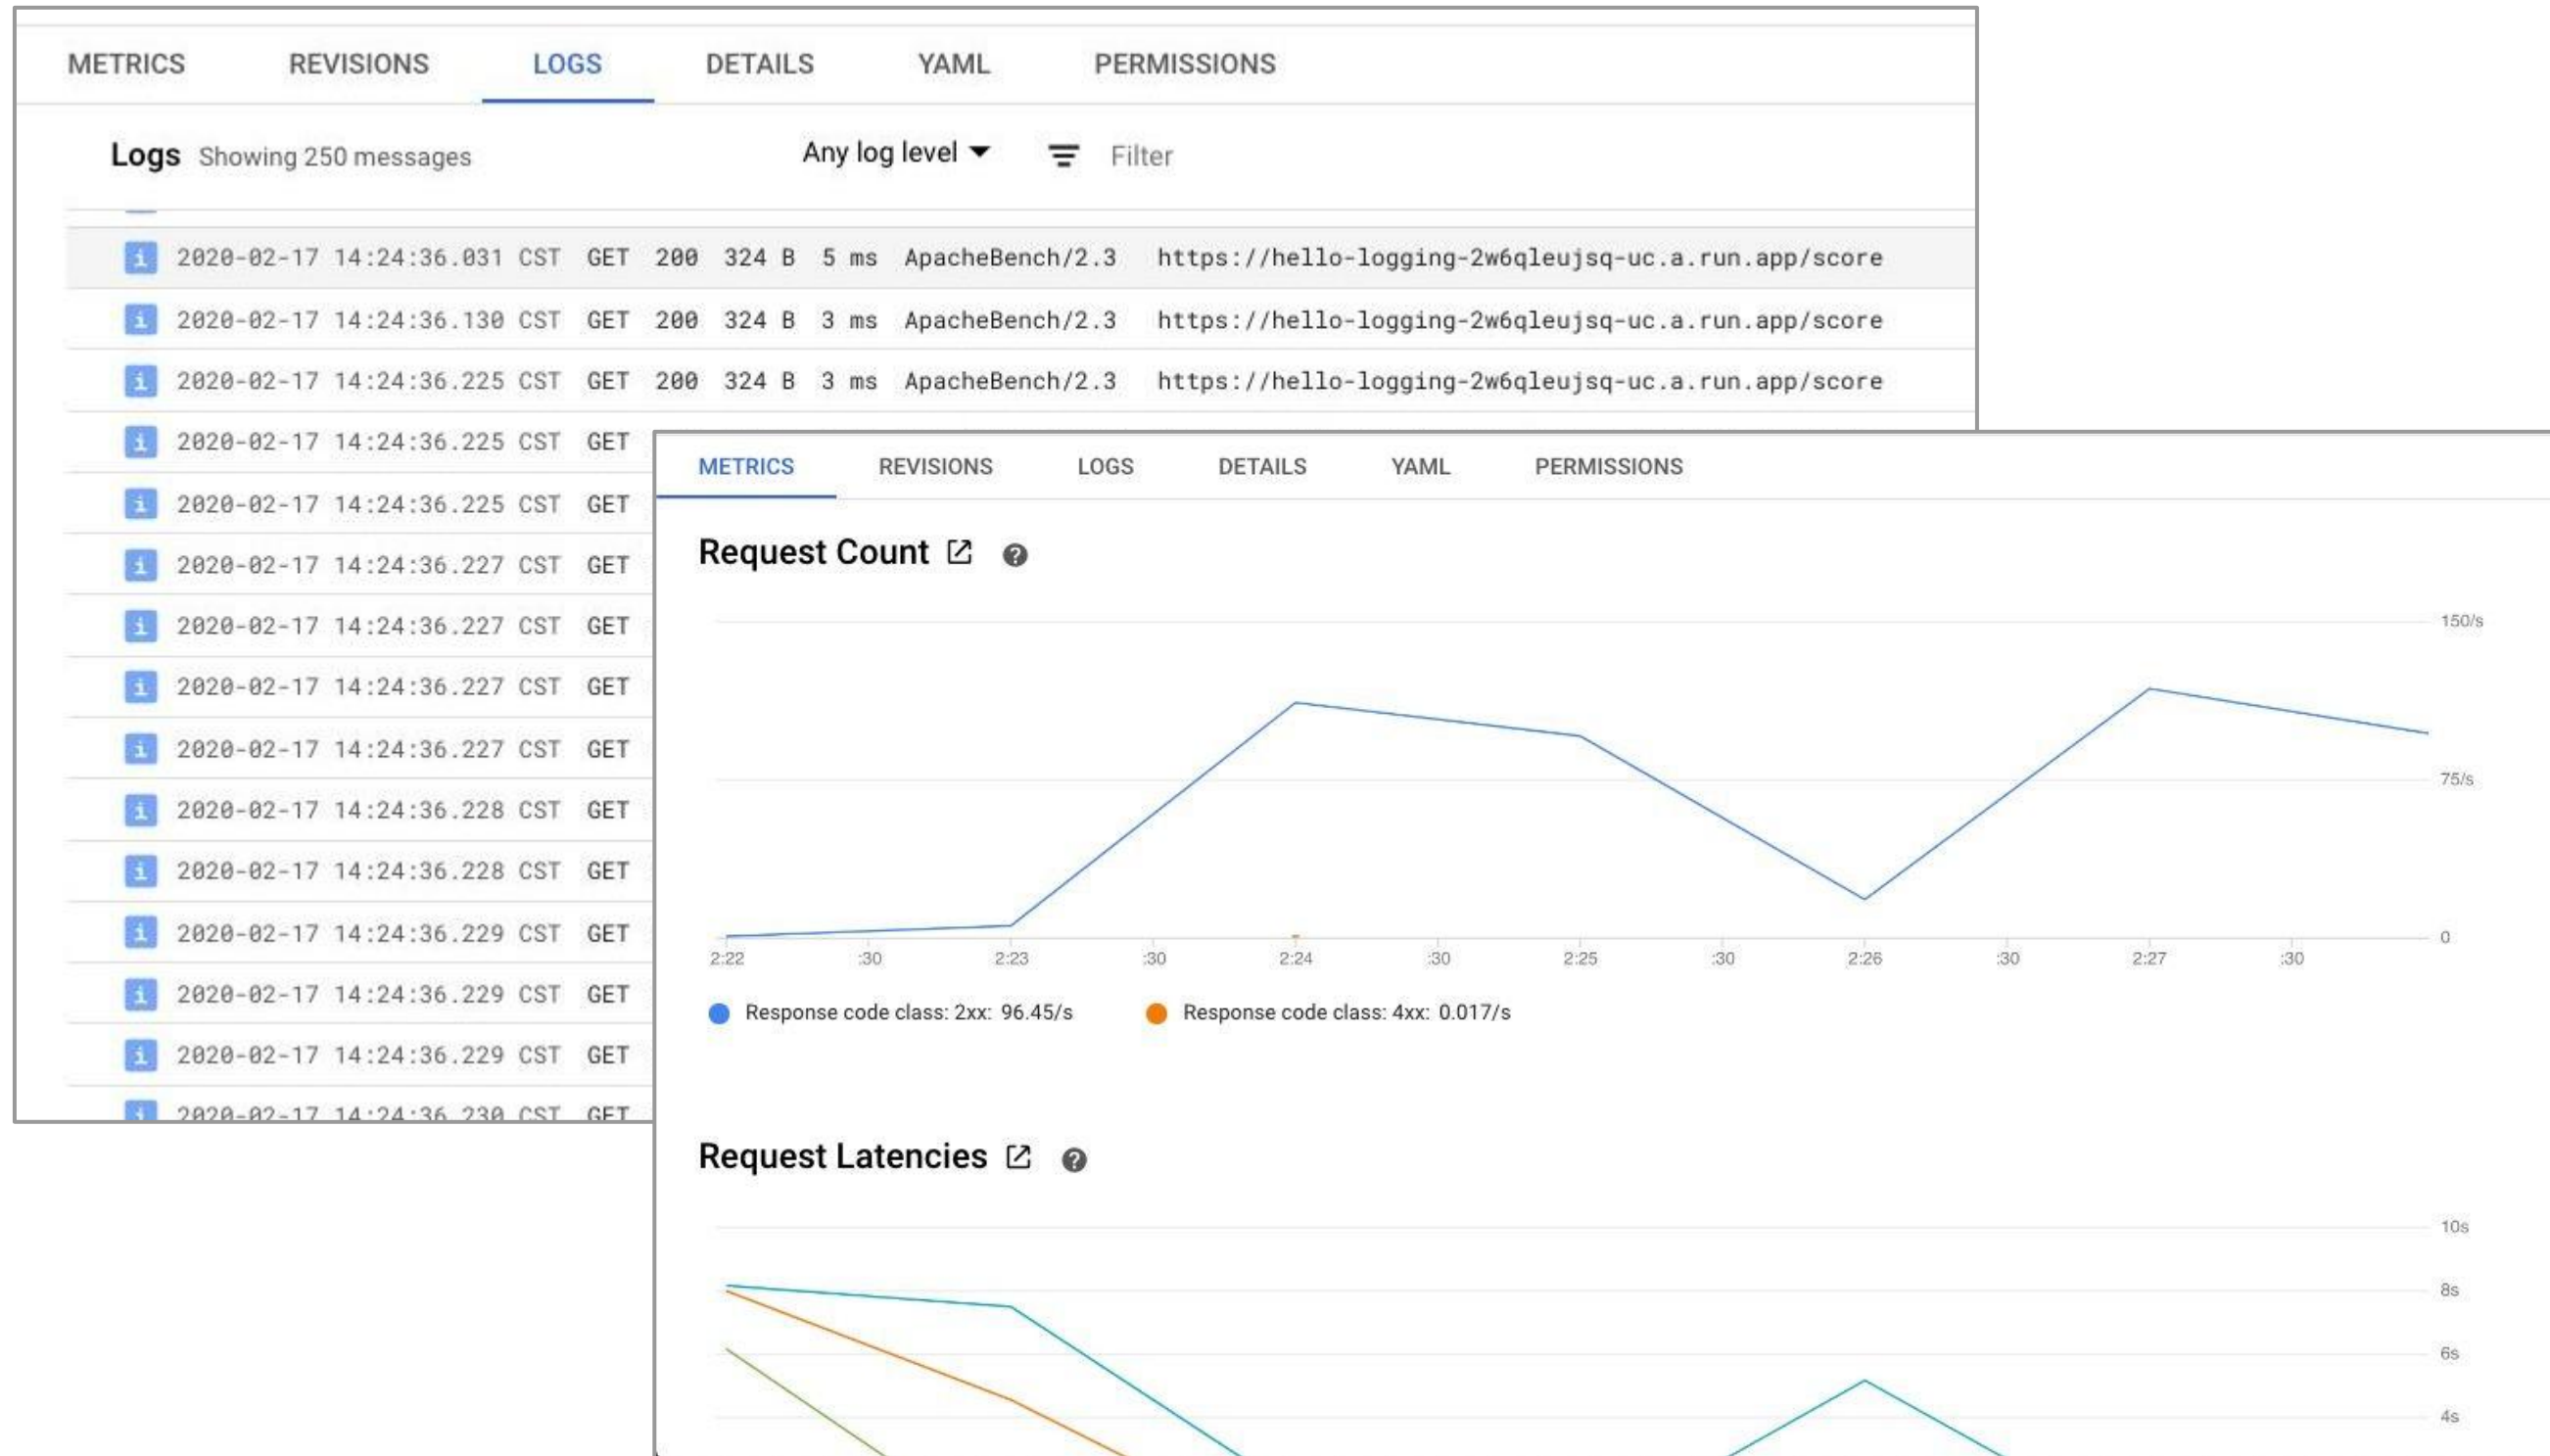


# Cloud Functions





# Cloud Run



---

# Agenda

Working with Agents  
Monitoring

Logging

Images and Agent Policies

Non-VM Resources

Exposing Custom Metrics

---

# Exposing custom metrics

Two fundamental approaches:

- Use the Cloud Monitoring API
- Use OpenCensus

# Custom metrics

Custom metric descriptor example in Python:

```
client = monitoring_v3.MetricServiceClient()
project_name =
client.project_path(project_id)
descriptor = monitoring_v3.types.MetricDescriptor()

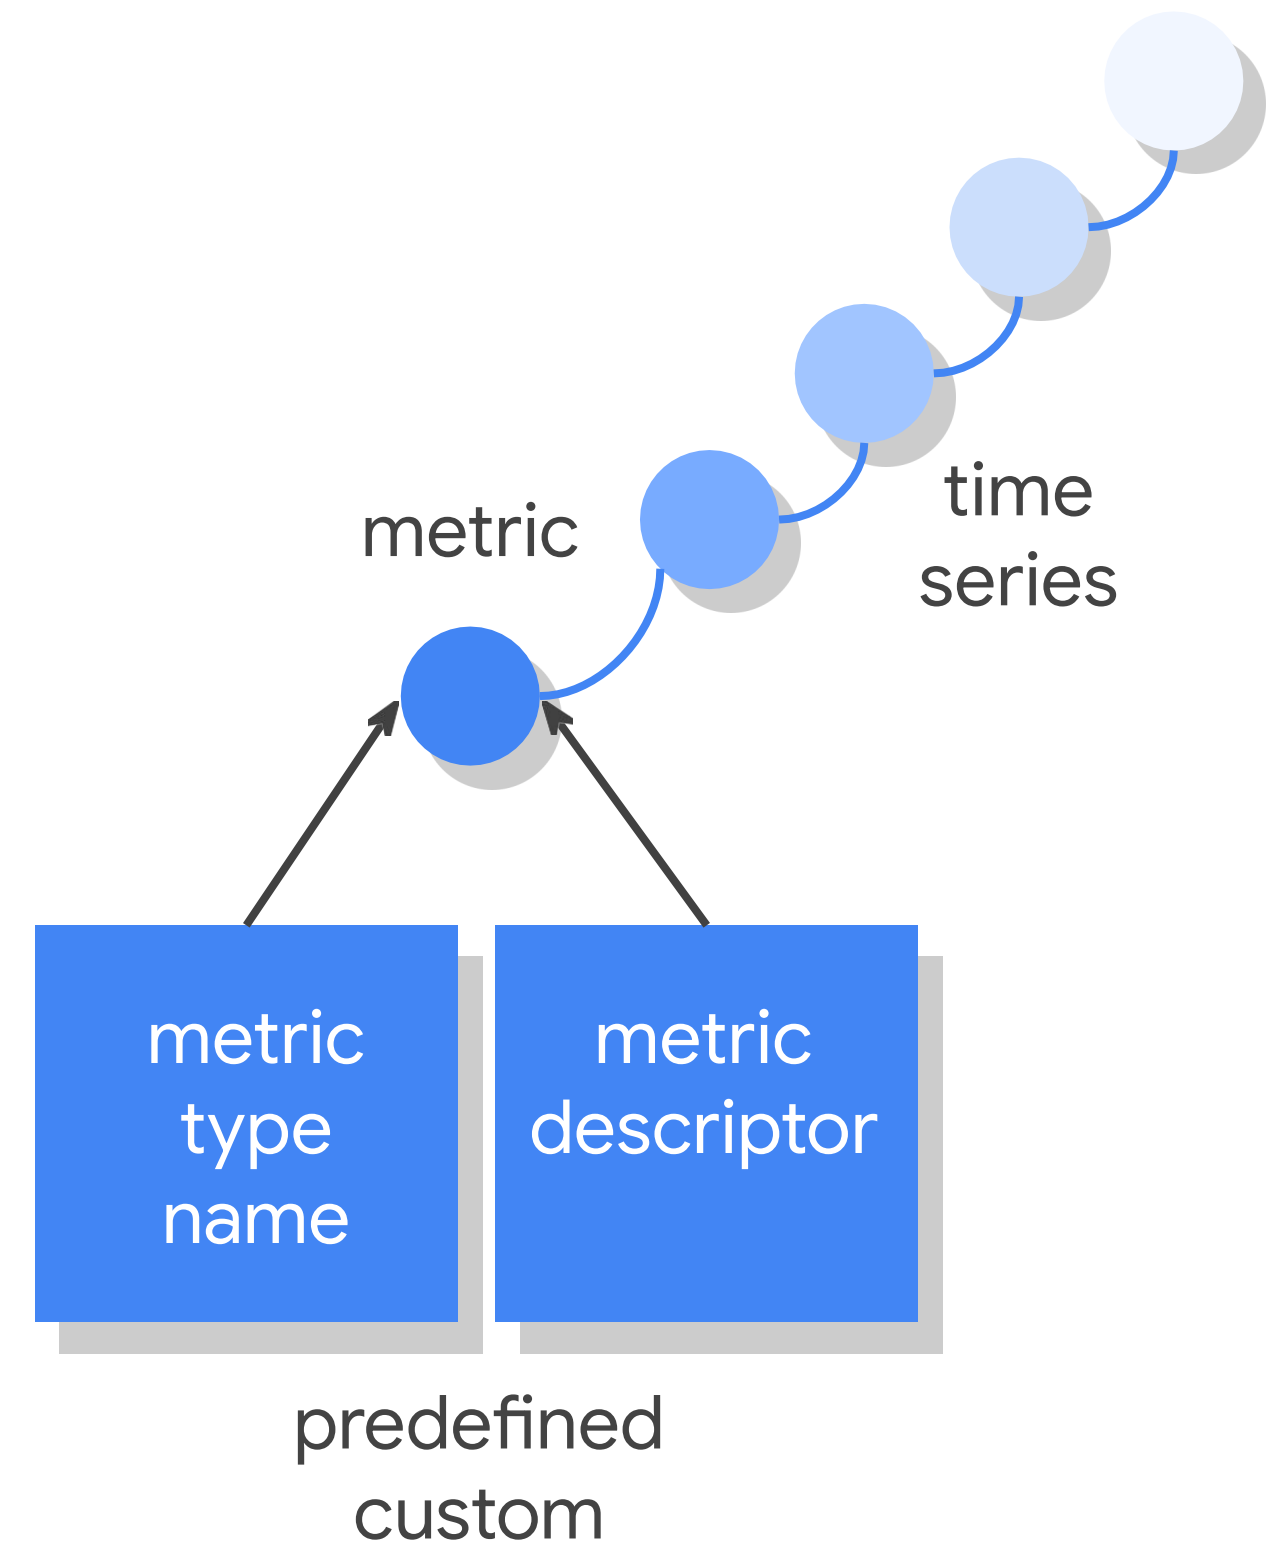
descriptor.type =
('custom.googleapis.com/my_metric')

descriptor.metric_kind = (
    monitoring_v3.enums.MetricDescriptor.MetricKind.GAUGE)

descriptor.value_type = (
    monitoring_v3.enums.MetricDescriptor.ValueType.DOUBLE)

descriptor.description = 'Custom metric example.'

client.create_metric_descriptor(project_name, descriptor)
```





# Custom metrics

Custom metric descriptor example in Python:

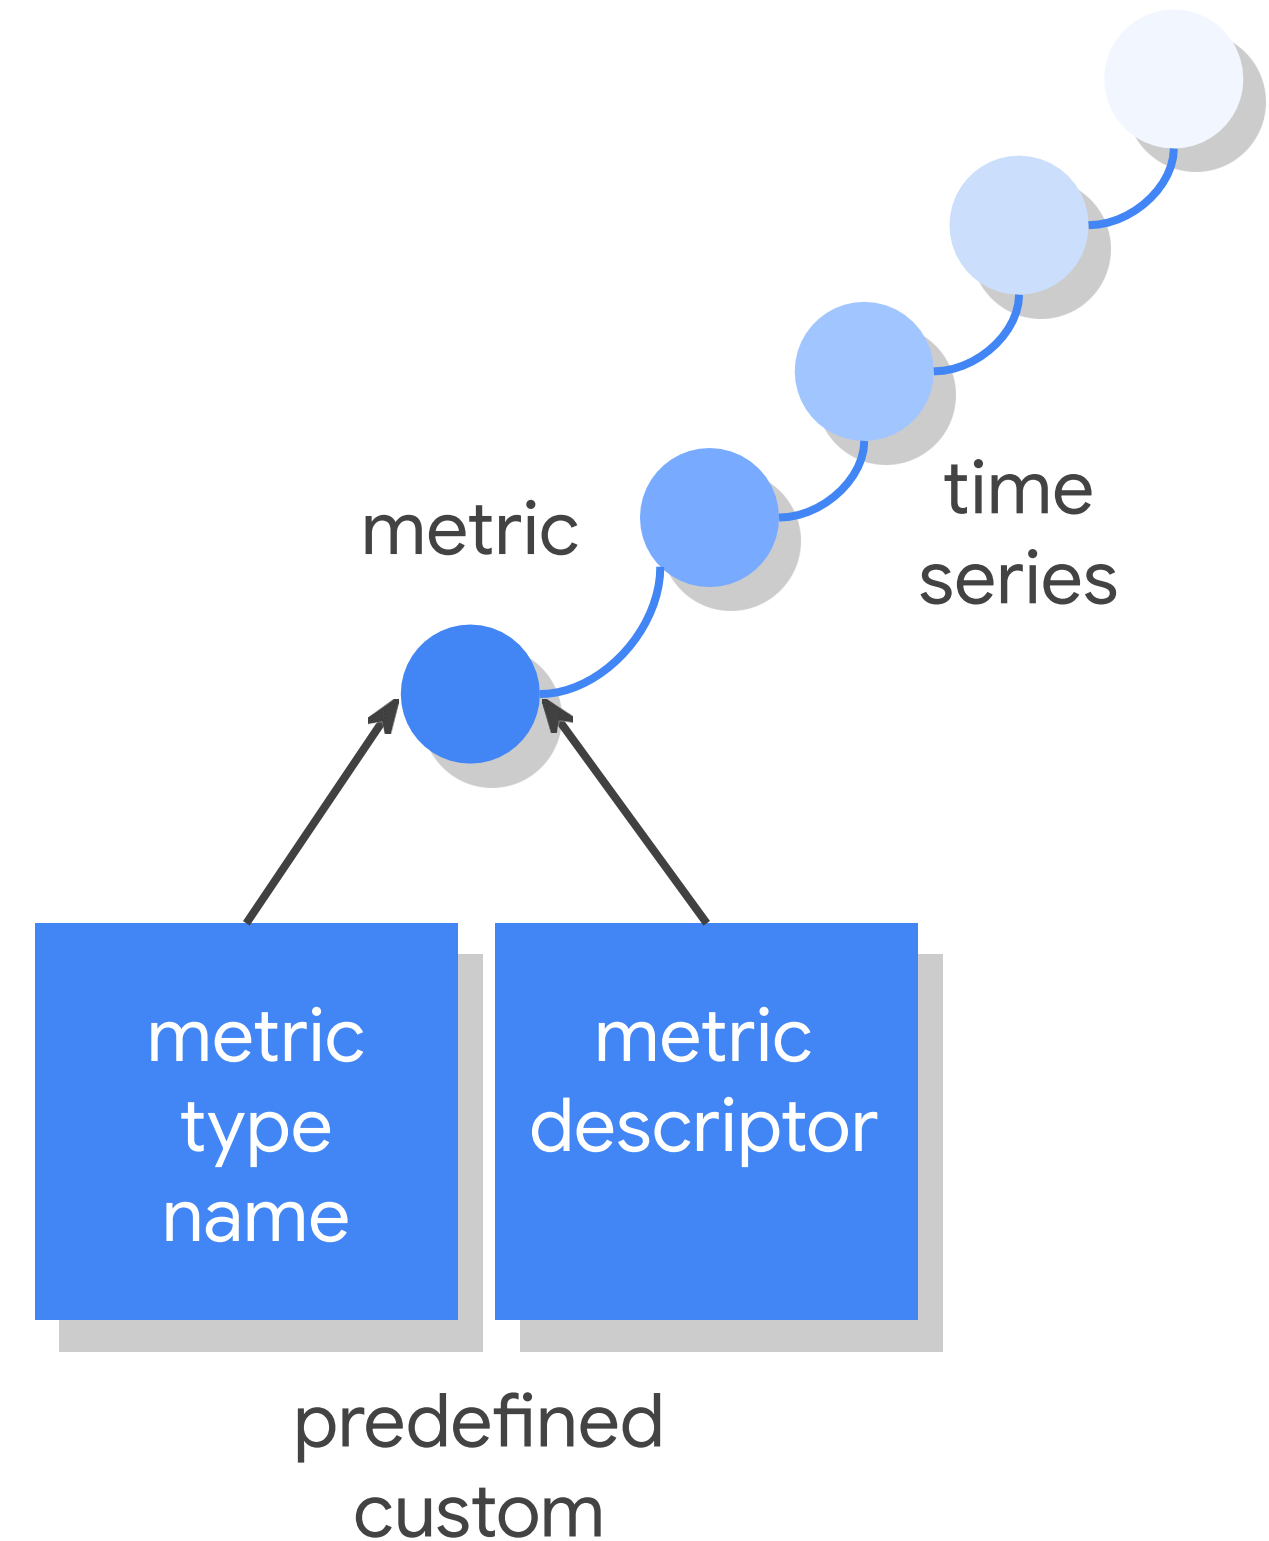
```
client = monitoring_v3.MetricServiceClient()
project_name =
client.project_path(project_id)
descriptor = monitoring_v3.types.MetricDescriptor()

descriptor.type =
('custom.googleapis.com/my_metric')

descriptor.metric_kind = (
    monitoring_v3.enums.MetricDescriptor.MetricKind.GAUGE)

descriptor.value_type = (
    monitoring_v3.enums.MetricDescriptor.ValueType.DOUBLE)

descriptor.description = 'Custom metric example.'
client.create_metric_descriptor(project_name, descriptor)
```



# Custom metrics

Custom metric descriptor example in Python:

```
client = monitoring_v3.MetricServiceClient()
project_name =
client.project_path(project_id)

descriptor = monitoring_v3.types.MetricDescriptor()

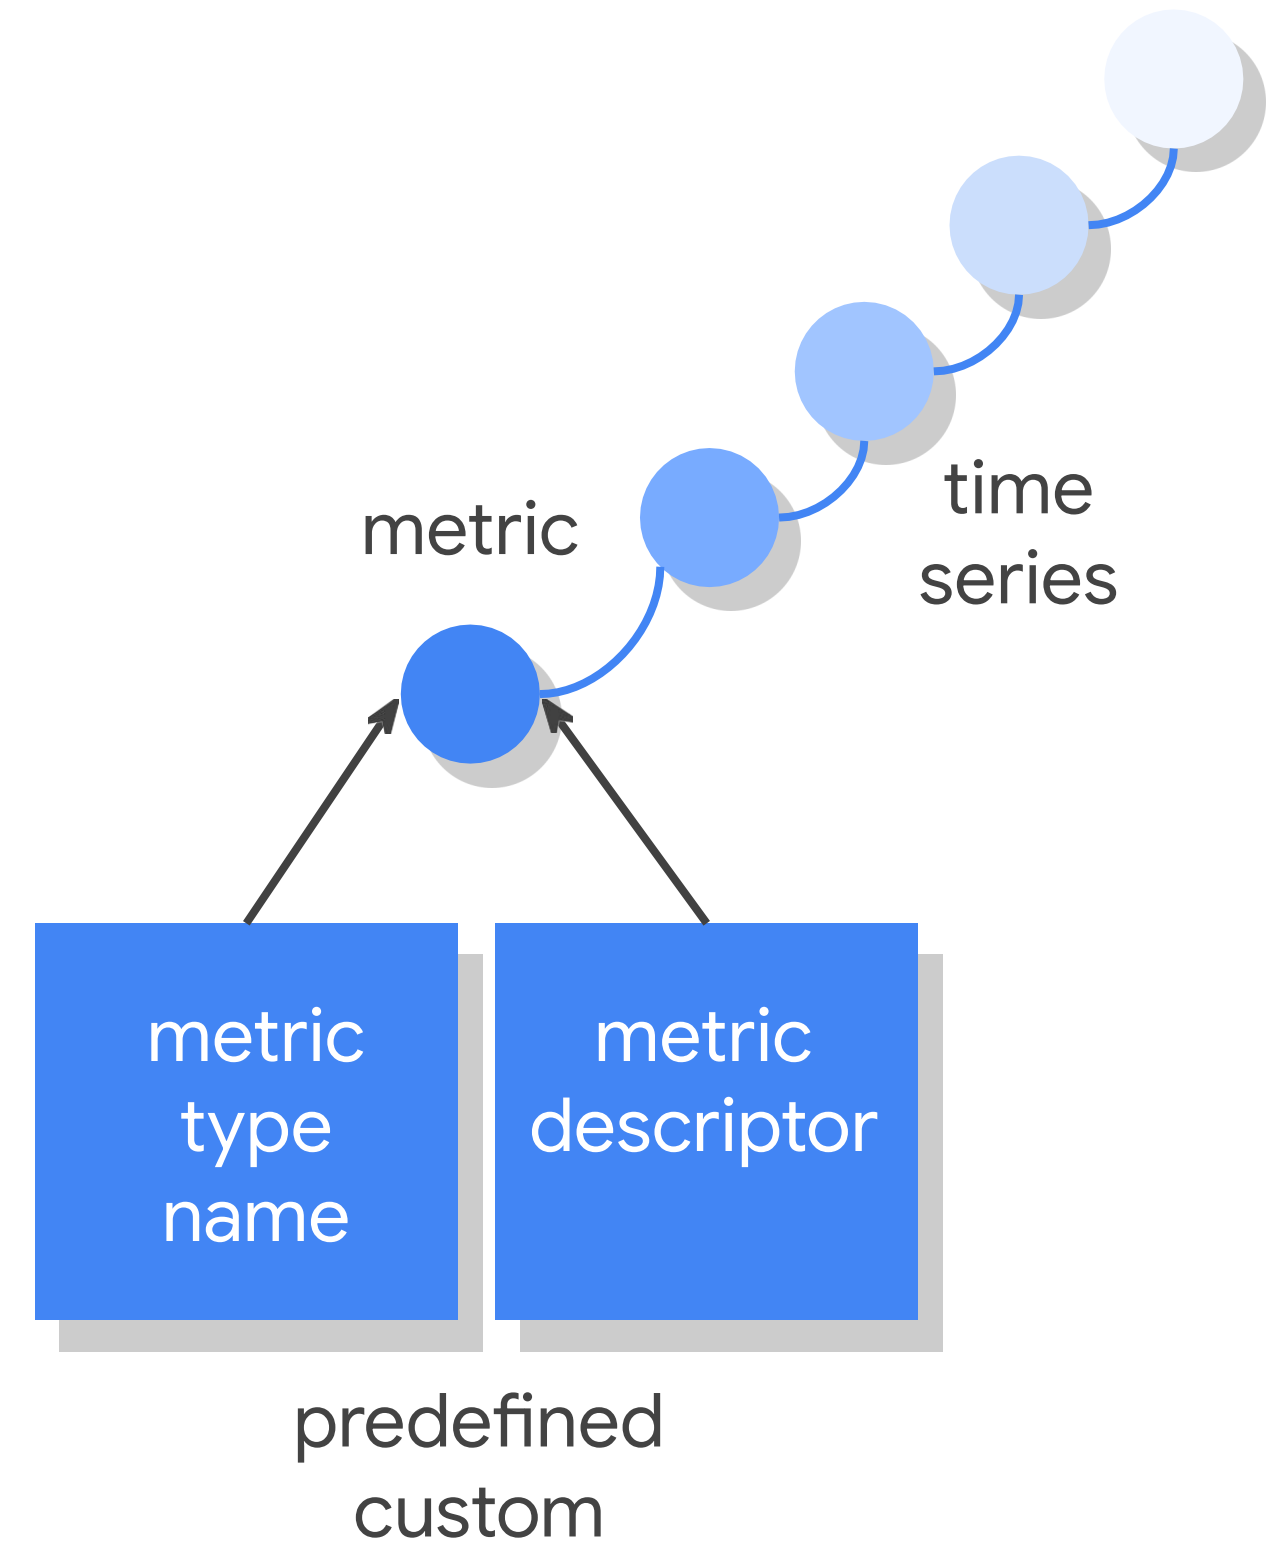
descriptor.type = ('custom.googleapis.com/my_metric')

descriptor.metric_kind = (
    monitoring_v3.enums.MetricDescriptor.MetricKind.GAUGE
)

descriptor.value_type = (
    monitoring_v3.enums.MetricDescriptor.ValueType.DOUBLE
)

descriptor.description = 'Custom metric example.'

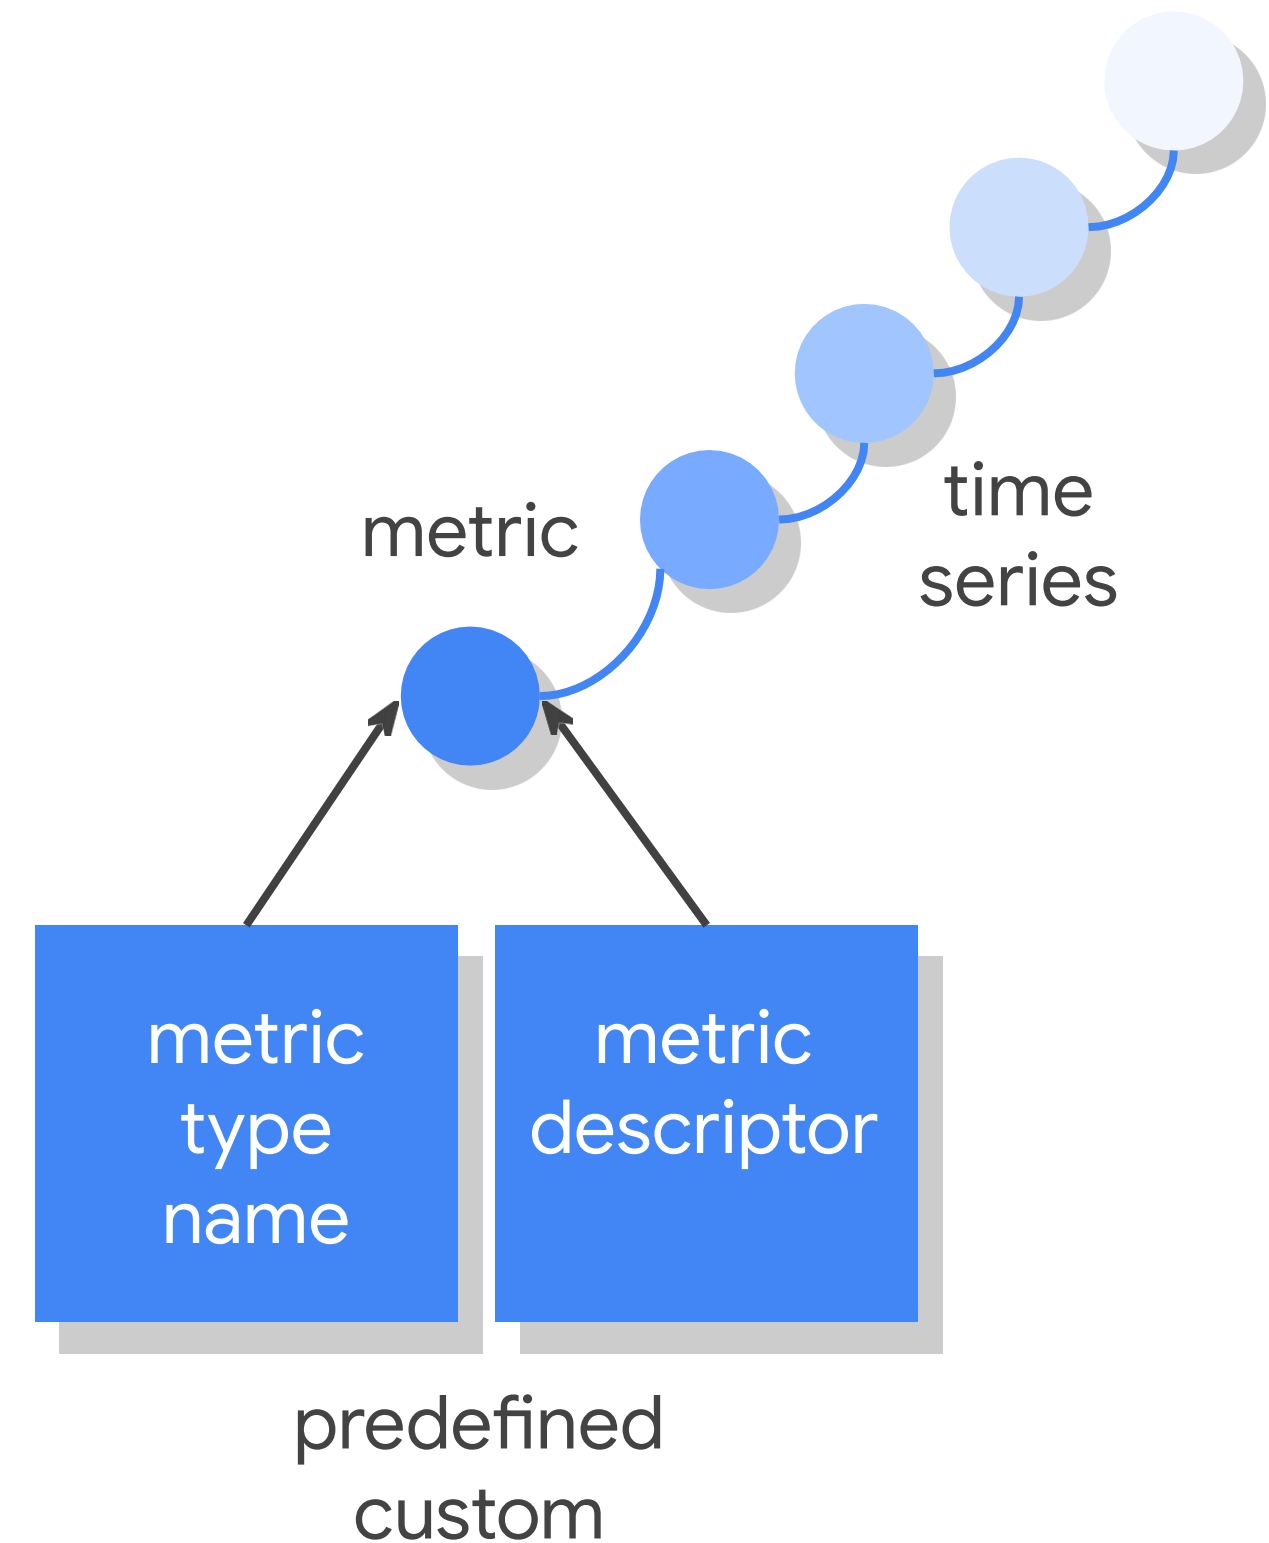
client.create_metric_descriptor(project_name, descriptor)
```



# Writing metrics

Using our custom metric, again in Python:

```
client =
monitoring_v3.MetricServiceClient()
project_name =
series = monitoring_v3.types.TimeSeries()
series.metric.type =
('custom.googleapis.com/my_metric')
series.resource.type = 'gce_instance'
series.resource.labels['instance_id'] = '126789012345
6789' series.resource.labels['zone'] =
point = series.points.add()
point.value.double_value =
3.14 now = time.time()
point.interval.end_time.seconds =
int(now)
client.create_time_series(project_name, [series])
```



# Writing metrics

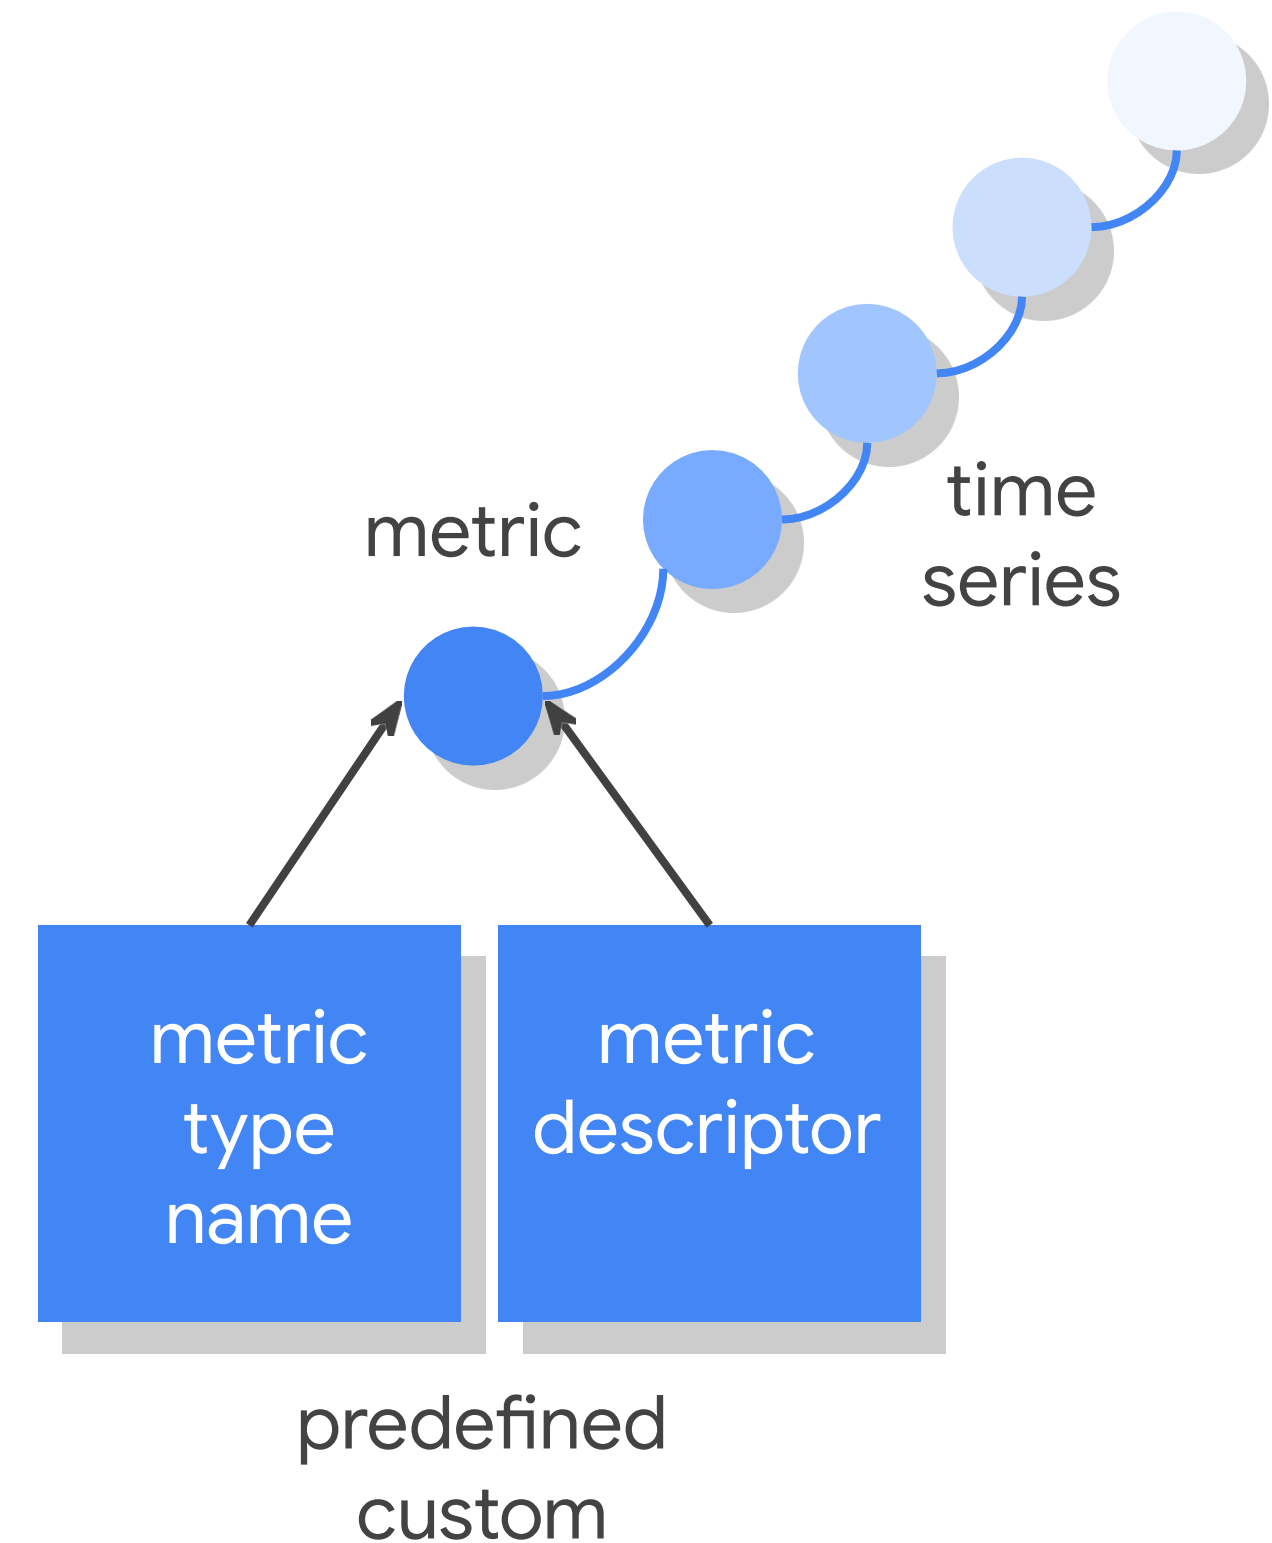
Using our custom metric, again in Python:

```
client =
monitoring_v3.MetricServiceClient()
project_name =
client.project_path(project_id)

series = monitoring_v3.types.TimeSeries()
series.metric.type =
('custom.googleapis.com/my_metric')
series.resource.type = 'gce_instance'
series.resource.labels['instance_id'] = '126789012345
6789' series.resource.labels['zone'] =
'us-central1-f'

point = series.points.add()
point.value.double_value =
3.14 now = time.time()
point.interval.end_time.seconds = int(now)

client.create_time_series(project_name, [series])
```



# Writing metrics

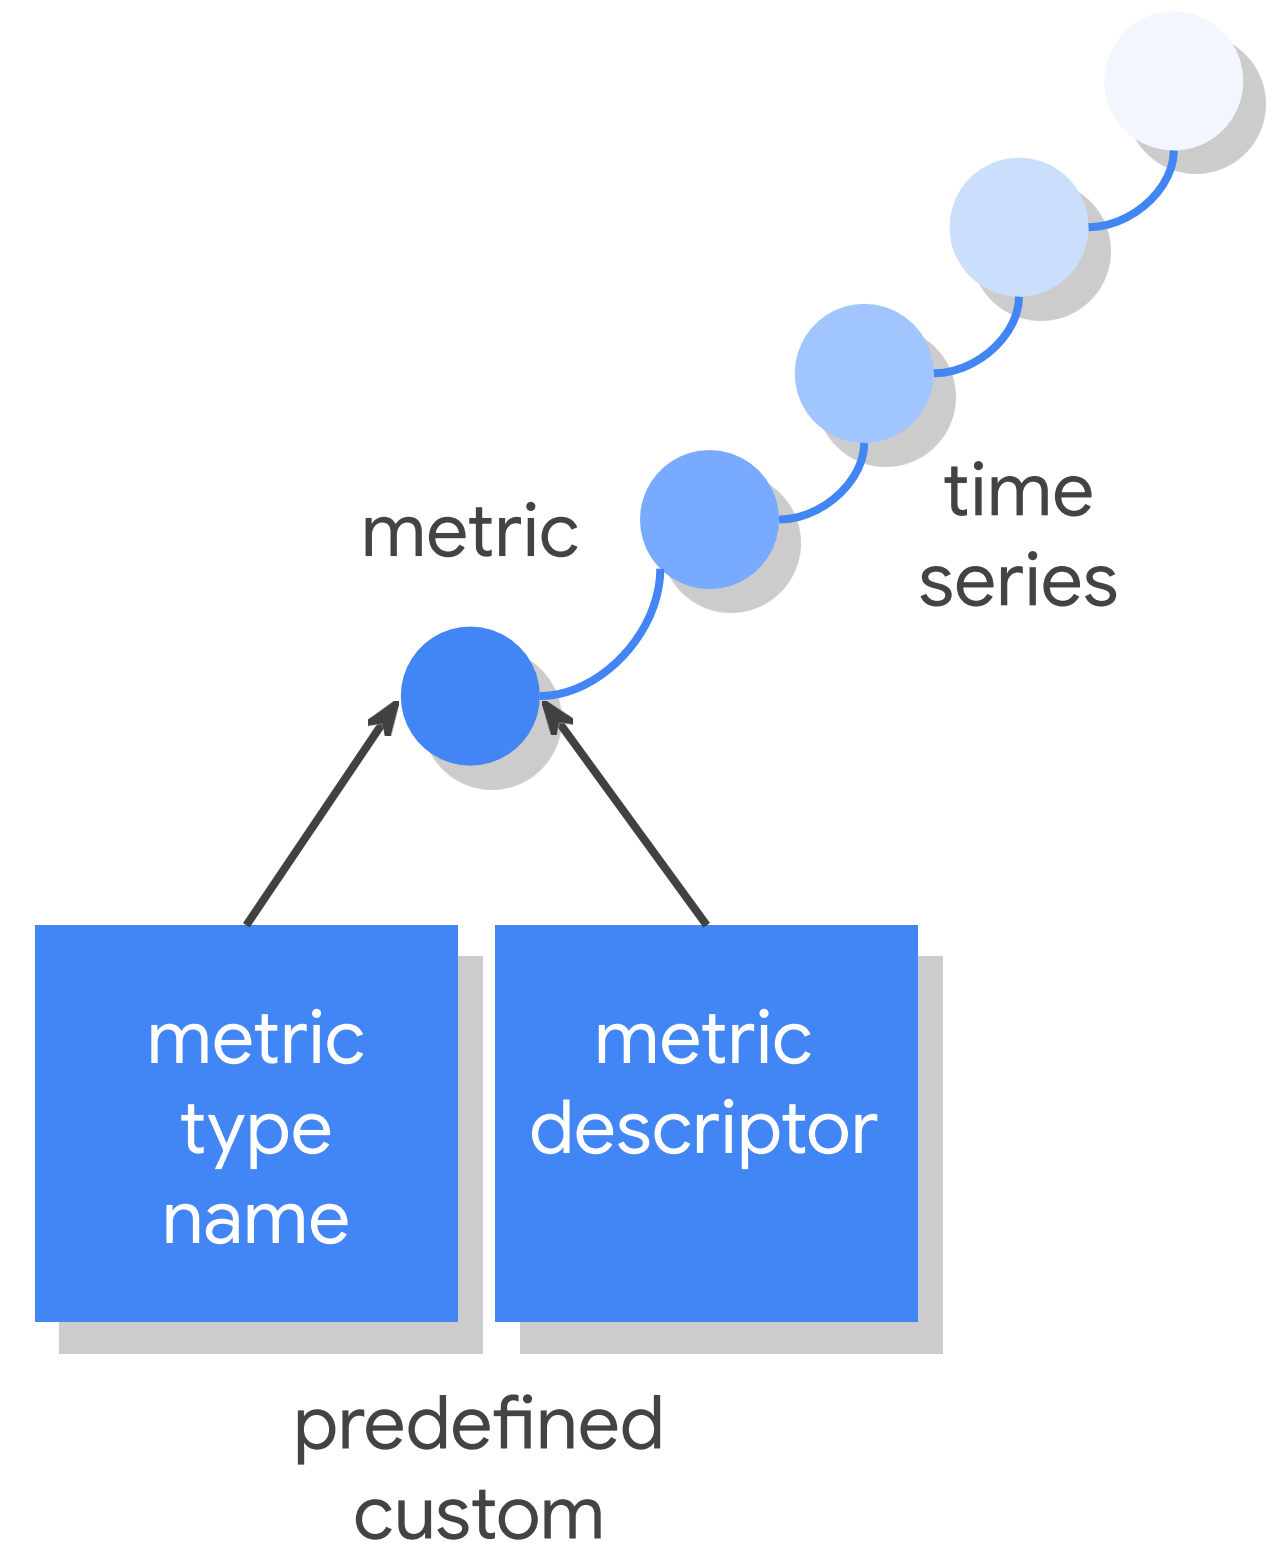
Using our custom metric, again in Python:

```
client =
monitoring_v3.MetricServiceClient()
project_name =
client.project_path(project_id)

series = monitoring_v3.types.TimeSeries()
series.metric.type =
('custom.googleapis.com/my_metric')
series.resource.type = 'gce_instance'
series.resource.labels['instance_id']='126789012345
6789' series.resource.labels['zone'] =
'us-central1-f'

point = series.points.add()
point.value.double_value =
3.14 now = time.time()
point.interval.end_time.seconds = int(now)

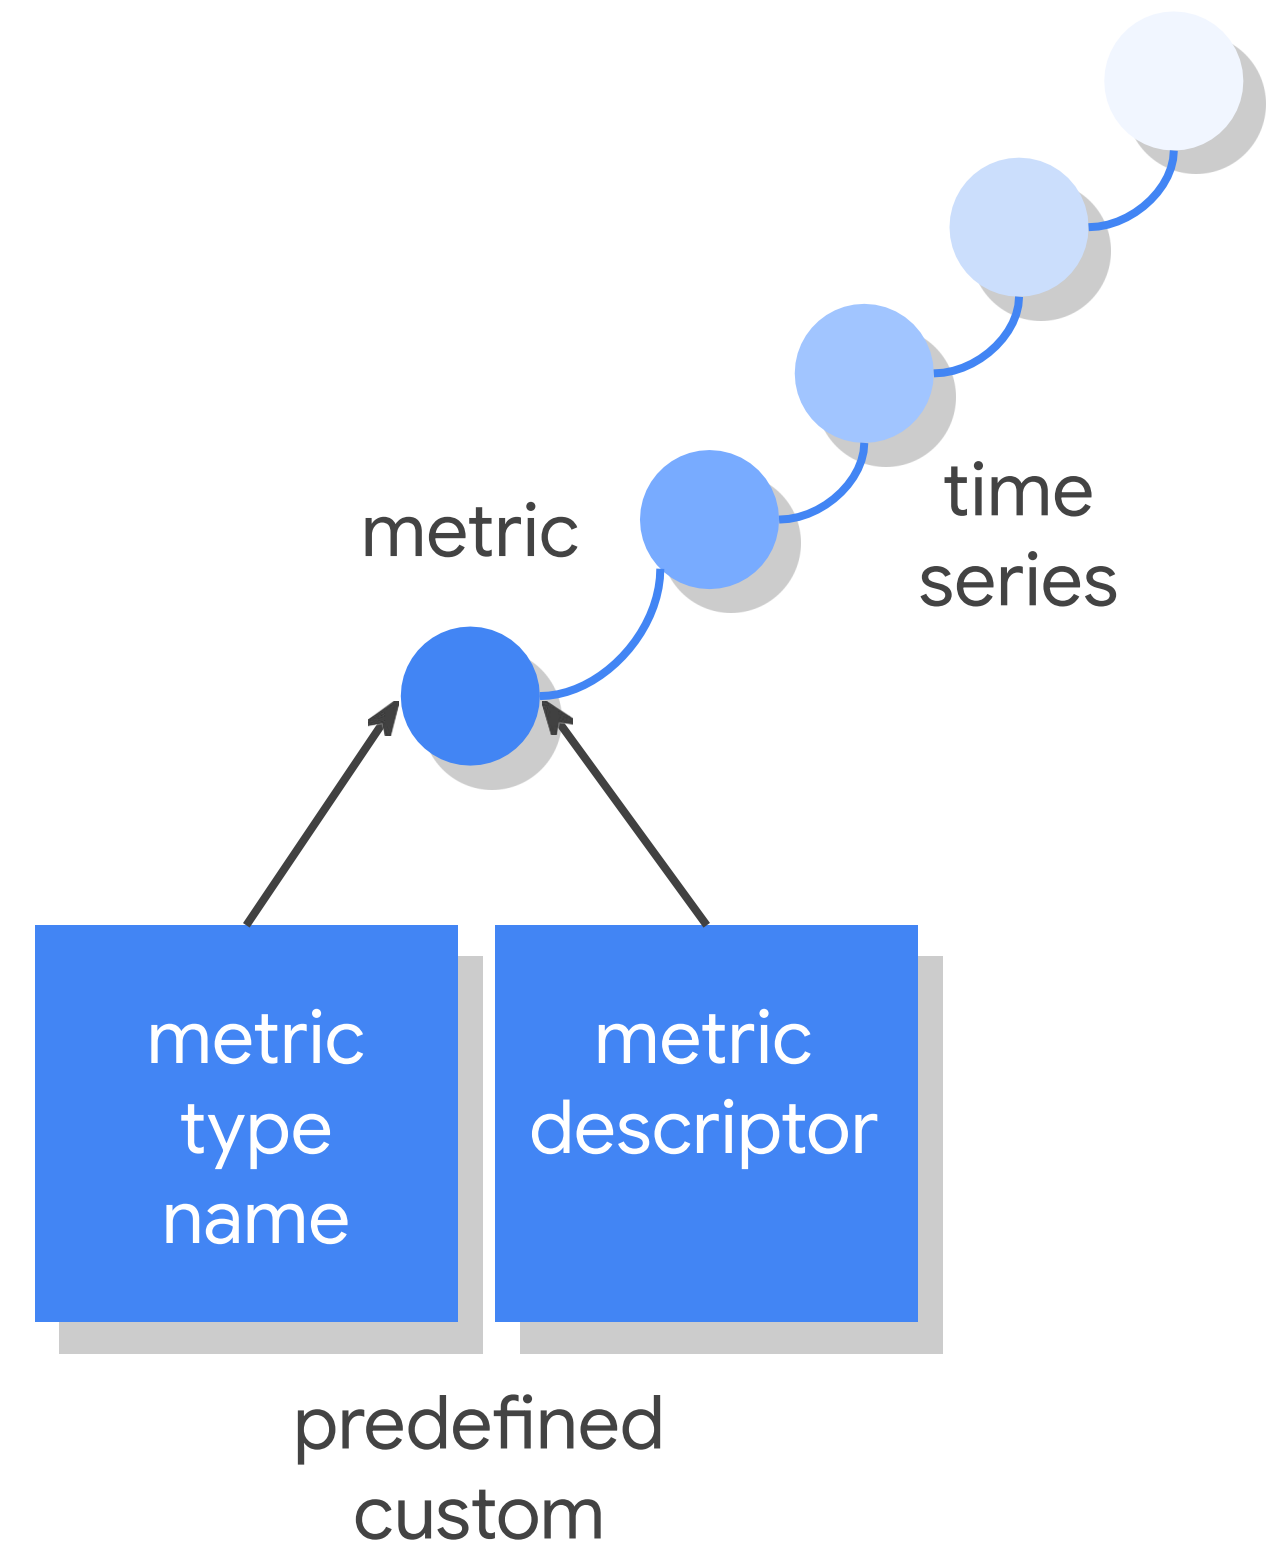
client.create_time_series(project_name, [series])
```



# Writing metrics

Using our custom metric, again in Python:

```
client =  
monitoring_v3.MetricServiceClient()  
project_name =  
series = monitoring_v3.types.TimeSeries()  
series.metric.type =  
( 'custom.googleapis.com/my_metric' )  
series.resource.type = 'gce_instance'  
series.resource.labels['instance_id'] = '126789012345'  
series.resource.labels['zone'] = 'us-east1-b'  
point = series.points.add()  
point.value.double_value = 3.14  
now = time.time()  
point.interval.end_time.seconds = int(now)  
  
client.create_time_series(project_name, [series])
```



---

# What is OpenCensus?

- Open-source library to help capture, manipulate, and export traces and metrics
  - Works with microservices and monoliths
- Supports many mainstream languages
  - Java, Python, Node.js, Go, C#, Erlang, and C++
- Low overhead and broadly supported
- OpenCensus is merging with OpenTracing to become OpenTelemetry
  - APIs planned to be backwards compatible



---

# Metrics expressed as measures and measurements

- A **Measure** represents a metric being recorded
  - **Name**: unique identifier
  - **Description**: purpose of the measure
  - **Unit**: string unit specifier, like “By”, “1”, or “ms”
    - [Unit codes](#)
  - Two measure value types: Int64 or a Float64
- A **Measurement** is a data point recorded as a **Measure**



---

# Views describe how measurements are collected

- A **View** represents the coupling of an **Aggregation** applied to a **Measure** and optionally **Tags**
- They contain:
  - Name: unique view name
  - Description
  - **Measure**: Measurement type
  - **TagKeys**: tagkeys used to group and filter metrics
  - **Aggregation**: How is the data gathered
    - Count, Distribution, Sum, or LastValue

# Exposing metrics from GCE using OpenCensus

